


Suffix trees

A *suffix tree* for a string **S** can be derived from a suffix trie by

- collapsing each path consisting of *unary* branch nodes into a single edge
- labelling each new edge with the concatenation of the labels from the corresponding trie edges





The suffix tree of a string S

First construct S^* from S by appending a character (\$) that does not appear in S .

The **suffix tree** of S^* (loosely, of S) is a rooted tree in which


- each edge is labelled by a string (its **edge label**)
- all branch nodes have ≥ 2 children
- no two children of a node have edge labels beginning with the same character
- there is a one-to-one correspondence between suffixes of S^* and leaf nodes v , in the sense that the path label of v is precisely that suffix
 - the **path label** of a node is the concatenation of edge labels on the path from the root to that node




Properties of suffix trees

- The number of leaf nodes is $n+1$ ($n = |S|$)
- The number of branch nodes is $\leq n$
 - proof is an exercise
- So the total number of nodes is $O(n)$
- The sum of the lengths of the edge labels is the same as in the suffix trie, so no better than $O(n^2)$ in the worst case
- To avoid this space overhead, use *short edge labels*
 - replace a string X by a pair of integers (i, j) such that $X = S^*(i \dots j)$
 - each short edge label has length $O(1)$
 - so the sum of the lengths of the edge labels is $O(n)$

suffix tree with full edge labels



mississippi\$
123456789012



suffix tree with short edge labels

Suffix tree construction 1

Let **S** be a string of length **n**. Append a **\$** character to form **S***

Naïve method for creating suffix tree **T**:

1. Create suffix trie for **S***

```
SuffTrie suffTrie = new SuffTrie();  
for (int j=1; j<=n+1; j++)  
    insert(suffTrie, S* [j..n+1]);
```

2. Amalgamate edges of paths containing unary branch nodes into single edges
3. Create short edge labels

Worst-case time complexity:

1. **O(n²)**
2. **O(n²)**
3. **O(n²)**

Overall: **O(n²)**. Space complexity: **O(n²)**

Suffix tree construction 2

Better method for creating suffix tree T :

Create suffix tree (with short labels) for S^* directly


```
SuffTree suffTree = new SuffTree();  
for (int j=1; j<=n+1; j++)  
    insert(suffTree, S*[j..n+1]);
```

To insert a suffix, follow an existing path and split from a certain point


Worst-case time complexity still $O(n^2)$

But space complexity $O(n)$

And average time complexity $O(n \log n)$



First 7 suffixes inserted; insert suffix 8



Suffix tree construction 3

Even better algorithms

For a string of length n over an alphabet Σ :

Weiner (1973), McCreight (1976), Ukkonen (1995)

All $O(n \log |\Sigma|)$ time and $O(n)$ space for suffix tree construction

More recently:

Farach (1997)


$O(n)$ time and space for suffix tree construction

These algorithms are beyond the scope of the course

Problem solving with a suffix tree 1


- **Searching a text T of length n , for a string S of length m**
 - build the suffix tree for T
 - $O(n)$ time
 - follow a path from the root matching characters
 - $O(m)$ time
 - all occurrences are represented by the leaf node descendants from the point where the search ends
 - can search a text of length n for k short strings of total length r in $O(n + r)$ time
 - e.g. find all occurrences of **is** in **mississippi**

suffix tree with full edge labels



each descendant leaf is an occurrence

search ends here




suffix tree with short edge labels


Problem solving with a suffix tree 2

- Finding a *longest repeated substring* in a text T , length n
 - build the suffix tree for T
 - $O(n)$ time
 - traverse the tree
 - a longest repeat is represented by a branch node having greatest string depth
 - the *string depth* of a node is the sum of the lengths of the edge labels on the path from the root to that node
 - i.e., it is the length of the path label of that node
 - traversal is also $O(n)$, so the entire algorithm is $O(n)$

suffix tree with full edge labels



Here it is –
representing
string issi




Problem solving with a suffix tree 3

- Finding a *longest common substring (LCSt)* of strings S_1, S_2 of lengths m, n
 - example - strings $S_1 = \text{dadbcdb}$ and $S_2 = \text{abcdacda}$
 - build the *generalised suffix tree T* for S_1 and S_2
 - this is the suffix tree for $S = S_1 \# S_2$, i.e., T stores the suffixes of the concatenation of $S_1, \#, S_2, \$$ in that order (where neither $\#$ nor $\$$ occurs in S_1 or S_2)
 - traverse the tree — a longest common substring is represented by a *common branch* node with greatest string depth
 - a branch node is *common* if it has two descendant leaf nodes representing positions starting in S_1 and S_2


Example - strings $S_1 = \text{dadbcd}\mathbf{b}$ and $S_2 = \text{abcdacda}$.

12345678901234567

Form $S^*=\text{dadbcd}\mathbf{b}\#\text{abcdacda}\$$ and the generalised suffix tree T :



● = common branch node, string depth 2



● = common branch node, string depth 3

Longest common substring is **bcd**

Identifying common branch nodes

Maintain booleans at each node v of T as follows:

Define $b_i(v)$ to be **true** if and only if the path label of a descendant leaf node of v starts in string S_i ($i \in \{1, 2\}$)

A branch node v of T is **common** if and only if $b_1(v) \wedge b_2(v)$

Computing the b_i values

Let v be a leaf node corresponding to suffix no. j of S^* , so path label of v is $S^*(j \dots m+n+2)$

$b_1(v) = \text{true}$ if and only if $1 \leq j \leq m$

$b_2(v) = \text{true}$ if and only if $m+2 \leq j \leq m+n+1$

For a branch node v , $b_i(v) = b_i(w_1) \vee b_i(w_2) \vee \dots \vee b_i(w_r)$

where w_1, \dots, w_r are the children of v ($i \in \{1, 2\}$)

LCSt - putting it all together

Given strings S_1, S_2 of lengths m, n .

1. Build the generalised suffix tree T — $O(m+n)$ time
2. Calculate the $b_i(v)$ values — $O(m+n)$ time using a traversal of T
3. Output the path label of a common branch node w of T with maximum string depth — $O(m+n)$ time

(In practice, w may be computed during the traversal in 2)

Overall $O(m+n)$ time and space to find LCSt of S_1, S_2

Finding an LCSt of S_1 and S_2 :


Why did we need to append # to S_1 ?

Why not just consider $S_1S_2\$$?

Example - suppose $S_1=a$ and $S_2=bab$

Clearly the unique LCSt of S_1 and S_2 is a

Consider the suffix tree for $S_1S_2\$ = abab\$$ rather than $S_1\#S_2\$$ as before



Algorithm incorrectly reports that ab is an LCSt of S_1 and S_2