

Optimisation problems

- Many decision problems are decision versions of *optimisation problems*
 - problems that involve maximising or minimising a value over a set of feasible solutions
- Formally, an *optimisation problem* has:
 - A set of instances I
 - A function SOL that associates with any instance the set of feasible solutions
 - A measure function m that assigns a non-negative integer $m(x,y)$ to any feasible solution y for a given instance x
 - a **GOAL**, that is, either max or min
- Given $x \in I$, let $m^*(x) = \text{GOAL}\{m(x,y) : y \in SOL(x)\}$ denote the *optimal measure*
- Given $x \in I$, the objective is to find, an *optimal solution*, i.e. a feasible solution $y^* \in SOL(x)$ such that $m(x,y^*) = m^*(x)$

Example optimisation problem

Maximum Satisfiability (MAX-SAT):

- **Instance:** Boolean formula B in CNF (as for SAT)
- **Feasible solutions:** All truth assignments for B
- **Measure:** Number of clauses of B that are satisfied
- **Goal:** max

Objective is to find a truth assignment for B which satisfies the largest number of clauses of B simultaneously

Example:

$$(x_1 \vee \bar{x}_2 \vee \bar{x}_3) \wedge (\bar{x}_1 \vee x_2 \vee \bar{x}_3) \wedge (\bar{x}_1 \vee \bar{x}_2 \vee x_3)$$

Setting $f(x_1) = F$ and $f(x_2) = f(x_3) = T$ satisfies two clauses

Setting $g(x_1) = g(x_2) = g(x_3) = T$ satisfies three clauses

Decision version of an optimisation problem

- Given an optimisation problem Π , the *decision version* of Π , denoted by Π_d , is defined as follows:
 - Instance: any instance $x \in I$; integer k (the target)
 - Question: is there a feasible solution $y \in SOL(x)$ such that $m(x,y) \leq k$ (if $GOAL=\min$)
or $m(x,y) \geq k$ (if $GOAL=\max$)?
- Usually Π is solvable in polynomial time if and only if Π_d is
- Example decision version – MAX-SAT-D:
 - Instance: Boolean formula in CNF with m clauses and target integer k
 - Question: Is there a truth assignment that satisfies k or more clauses simultaneously?
- Restriction: MAX-SAT-D with $k=m$ is SAT so MAX-SAT-D is NP-complete

The Classes NPO and PO

- **NPO:** Optimisation problems whose decision versions are in NP
 - Examples: MAX-SAT, Minimum Vertex Cover, Maximum Clique, Minimum Graph Colouring, Maximum Cut, . . .
- **PO:** those NPO problems solvable in polynomial time
 - Example: Maximum Matching
- We say that a problem Π in NPO is **NP-hard** if its decision version Π_d is NP-complete
 - Example: MAX-SAT
- **Theorem 1:** P = NP if and only if PO = NPO
- **Theorem 2:** If P \neq NP and Π is NP-hard then $\Pi \notin$ PO

Approximation algorithms

Let Π be a problem in NPO

- An **exact** or *optimising algorithm* finds an optimal solution, given any instance x of Π

If Π is NP-hard, we consider *approximation algorithms* for Π

- An *approximation algorithm* A for Π returns, in polynomial time, $A(x)$, where $A(x) \in SOL(x)$, for any instance x of Π
- A has a *performance guarantee* c ($c \geq 1$) if
 - GOAL=min and $m(x, A(x)) \leq c \times m^*(x)$ for all instances x of Π , or
 - GOAL=max and $m(x, A(x)) \geq 1/c \times m^*(x)$ for all instances x of Π
- Refer to A as a *c-approximation algorithm*
 - The closer c is to 1, the better the feasible solution A guarantees to deliver

Examples of approximation algorithms

Example 1

- Approximation algorithm for **Minimum Vertex Cover** with performance guarantee **2**
 - see tutorial exercises

Example 2

- Approximation algorithm for **MAX-SAT** with performance guarantee **2**
- Returns a truth assignment satisfying at least half the number of clauses satisfied by an optimal truth assignment
- In fact, returns a truth assignment that satisfies at least half the number of clauses in the given formula

Approximation algorithm for MAX-SAT

```
/** Input: Boolean formula b in CNF
 * Output: Truth assignment f */
for ( variable x : b )
    f(x) = true;
while (b has at least one clause)
{   let u be a variable in b;
    p = number of clauses in b which contain u;
    q = number of clauses in b which contain  $\bar{u}$ ;
    if ( p  $\geq$  q ) {
        f(u) = true;
        remove clauses containing u from b;
        delete occurrences of  $\bar{u}$  from b;
    }
    else
    {   f(u) = false;
        remove clauses containing  $\bar{u}$  from b;
        delete occurrences of u from b;
    }
    delete empty clauses from b;
}
return f;
```

Claim: this algorithm returns a truth assignment which satisfies at least half of the m clauses

Proof: By induction on the number of variables n in any CNF formula B

Base step $n=1$: B is u_1 or \bar{u}_1 or $(u_1 \vee \bar{u}_1)$ or $u_1 \wedge \bar{u}_1$ or $u_1 \wedge (u_1 \vee \bar{u}_1)$ or $\bar{u}_1 \wedge (u_1 \vee \bar{u}_1)$
– in all cases, algorithm satisfies at least one out of one or two clauses

Inductive step: assume true for all CNF formulae with $n-1$ variables.

Let u be the first variable to which a value has been assigned. Assume $p \geq q$ (argument similar for $p < q$).

Let B' be CNF formula resulting from deletions of clauses and literals.
Assume B' contains r clauses.

By the induction hypothesis, algorithm returns a truth assignment f satisfying $\geq r/2$ clauses of B' .

Let $f(u)=\text{true}$; then f satisfies at least

$$p + r/2 \geq p + (m-p-q)/2 \geq m/2$$

clauses of B . \square

- **Theorem:** MAX-SAT has a **2**-approximation algorithm
- A more sophisticated algorithm achieves a performance guarantee of **1.2551** (Avidor et al, 2005)

The Class APX


- NPO problems Π that admit a **c**-approximation algorithm, for some constant $c \geq 1$
- Problem Π is said to be **c-approximable** or **approximable within (a factor of) c**
- Examples:
 - Minimum Vertex Cover (**c=2**)
 - TSP under triangle inequality (**c=3/2**) – see below
 - MAX-SAT (**c=1.2551**)
- We have **PO \subseteq APX \subseteq NPO**

Example 3: the *Travelling Salesman Problem (TSP)*

- Approximation algorithms based on *spanning trees*
- Assume *triangle inequality* is satisfied
 - $d(x, z) \leq d(x, y) + d(y, z) \quad \forall x, y, z$
 - true in most applications
- Claim: For a TSP instance I let $\text{OPT}(I)$ be the length of the shortest tour, and $T(I)$ the cost of a minimum weight spanning tree; then $T(I) < \text{OPT}(I)$
 - because, discarding one edge from shortest tour gives a spanning tree
- Algorithm M (*“twice-around the tree”*):


find a minimum weight spanning tree T ;
carry out a depth-first traversal of T , following each
edge once in each direction;
take shortcuts to avoid visiting cities more than once;
return the tour so constructed;

Illustration of Algorithm M




Here, edge weight = Euclidean distance between the vertices

Illustration of Algorithm M




Here, edge weight = Euclidean distance between the vertices

Illustration of Algorithm M




distance travelled = $2 \times T(I)$

Illustration of Algorithm M



twice around the
tree with shortcuts


Illustration of Algorithm M



The final tour given
by Algorithm M

- distance travelled = $M(I) \leq 2 \times T(I) < 2 \times OPT(I)$ (triangle inequality)
(from earlier)
- Hence performance guarantee of 2

Illustration of Algorithm M



- More complex algorithm: Algorithm C (*Christofides' algorithm*)
- Involves computing minimum weight spanning tree, minimum weight matching and Eulerian Cycle
- Gives a performance guarantee of $3/2$