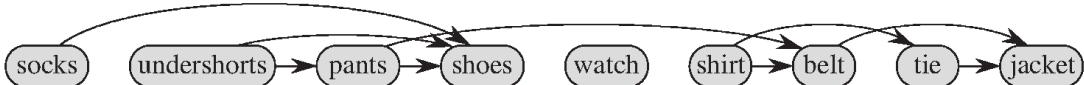


Exercises, chapter 6, solutions

1. Many different solutions are possible. Here is one, from:

- T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein: *Introduction to Algorithms (Third Edition)*, The MIT Press, 2009. [Figure 22.7]



2. (a) The difference between DFS in Chapter 2 and DFSTopologicalSort in Chapter 6 is that the latter uses some additional information in the form of a list named *sorted* for storing vertices of V . More precisely, whenever DFSTopologicalSort reaches a dead end at some vertex u during the depth-first search (i.e., when all neighbors of u have been taken care of in the **foreach**-loop on lines 2–4 and it's time to backtrack), it will insert u at the current head of *sorted* (line 5 of DFSTopologicalSort) before returning from the recursive call.
- (b) It is useful because it lets a critical path in G be computed efficiently with the “relaxation” technique. For relaxation to work, after we have used the estimated length ℓ of a longest path to some vertex u to update the estimated lengths of longest paths to the vertices in u 's adjacency list, the value of ℓ is not allowed to change. Considering the vertices in the order obtained by a topological sort of G guarantees that all vertices that have paths leading to u will have been handled before u , so the final value of ℓ is available once it is time to use ℓ .

3. Algorithm for checking if an undirected graph $G = (V, E)$ contains a cycle:

- If $|E| \geq |V|$ then return TRUE.
- Otherwise ($|E| \leq |V| - 1$), run HasCycle(G) below, which does a depth-first search of G using a procedure Modified-DFS. For any unvisited vertex i , Modified-DFS returns TRUE if and only if the connected component in G that i belongs to contains a cycle. The parameter *parent* specifies the vertex from which the search came; it is excluded when handling the neighbors of u so that $\langle \text{parent}, u, \text{parent} \rangle$ will not be considered to be a cycle.

```

HasCycle( $G$ )
1 Create a Boolean array visited of length  $|V|$  and initialize all of its entries to FALSE.
2 has_cycle  $\leftarrow$  FALSE
3 for  $i \leftarrow 1$  to  $|V|$  do
4   if not visited[ $i$ ] then
5     has_cycle  $\leftarrow$  has_cycle or Modified-DFS( $i, null$ )
6 return has_cycle

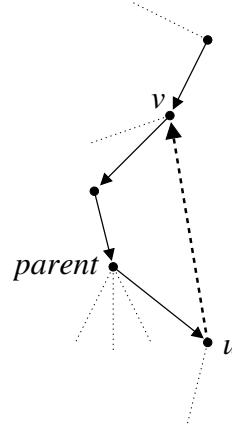
Modified-DFS( $u, parent$ )
1 visited[ $u$ ]  $\leftarrow$  TRUE
2 tmp  $\leftarrow$  FALSE
3 foreach  $v$  in AdjacencyList( $G, u$ ) do
4   if  $v \neq parent$  then
5     if visited[ $v$ ] then      /* A cycle has been found. */
6       tmp  $\leftarrow$  TRUE
7     else
8       tmp  $\leftarrow$  tmp or Modified-DFS( $v, u$ )
9 return tmp
  
```

Explanation:

First observe that the number of edges in any undirected graph $G = (V, E)$ without cycles is at most $|V| - 1$. Consequently, if the input graph satisfies $|E| \geq |V|$ then it must have a cycle, in which case the algorithm can return TRUE immediately in step (i).

On the other hand, if $|E| \leq |V| - 1$ then G may or may not have a cycle. To find the answer, the algorithm will perform a depth-first search of G in step (ii). There are two possibilities:

- If G contains a cycle, there is some edge that connects a vertex u to another vertex v such that: (i) v was visited by the depth-first search before visiting u ; and (ii) v is not the parent of u in the depth-first search. This will be discovered by the algorithm on line 5 of **Modified-DFS** during the visit to u because then v is in $\text{AdjacencyList}(G, u)$ and: (i) $\text{visited}[v]$ will already have been set to TRUE; and (ii) v is different from parent . See the following figure for an illustration:



When this situation occurs, line 6 of **Modified-DFS** sets tmp equal to TRUE. Next, since the logical **or**-operation between TRUE and anything else gives TRUE, the value TRUE will be returned to the calling function and all the way back up to **HasCycle**, where it is assigned to the variable has_cycle . The value of has_cycle can never be changed back to FALSE, so the value TRUE is returned at the end.

- If G doesn't contain any cycles, $\text{visited}[v]$ will never be TRUE for any neighbor v of the currently visited vertex during the depth-first search (except for the parent of u in the depth-first search). Therefore, line 6 of **Modified-DFS** will never be executed. Since the logical **or**-operation between FALSE and FALSE is FALSE, the value of has_cycle will not change from its initial value FALSE, and the algorithm will correctly return FALSE at the end.

Time complexity analysis:

Checking if $|E| \geq |V|$ holds in step (i) can be done in $O(|V|)$ time by counting the edges one by one until all edges have been counted or we reach $|V|$. Step (ii) takes $O(|V| + |E|)$ time according to the analysis of the depth-first search algorithm, but here we also have $O(|V| + |E|) = O(|V|)$ because this step is only executed when $|E| \leq |V| - 1$. In total, the algorithm's time complexity is $O(|V|)$.