

Chapter 8: Routing, Arbitrage

Panos Louridas

Athens University of Economics and Business
Real World Algorithms
A Beginners Guide
The MIT Press


Outline

1 Internet Routing

2 The Bellman-Ford Algorithm

3 Arbitrage


The TCP/IP Protocol Stack with HTTP




Routing

- The Internet consists of a set of *autonomous systems*.
- Routing between autonomous systems is done with the Border Gateway Protocol (BGP).
- Routing inside an autonomous system is done with the Routing Information Protocol (RIP).

An Autonomous System with Five Routers




RIP Example (1)



	R_1	R_2	R_3	R_4	R_5
R_1	0	1/D	-	5/D	-
R_2	1/D	0	1/D	-	-
R_3	-	1/D	0	2/D	3/D
R_4	5/D	-	2/D	0	-
R_5	-	-	3/D	-	0

Initial state.


RIP Example (2)



	R_1	R_2	R_3	R_4	R_5
R_1	0	1/D	$2/R_2$	5/D	-
R_2	1/D	0	1/D	-	-
R_3	-	1/D	0	2/D	3/D
R_4	5/D	-	2/D	0	-
R_5	-	-	3/D	-	0

$R_2 \rightarrow R_1.$


RIP Example (3)



	R_1	R_2	R_3	R_4	R_5
R_1	0	$1/D$	$2/R_2$	$5/D$	-
R_2	$1/D$	0	$1/D$	$3/R_3$	$4/R_3$
R_3	-	$1/D$	0	$2/D$	$3/D$
R_4	$5/D$	-	$2/D$	0	-
R_5	-	-	$3/D$	-	0

$R_3 \rightarrow R_2.$


RIP Example (4)



	R_1	R_2	R_3	R_4	R_5
R_1	0	$1/D$	$2/R_2$	$4/R_2$	$5/R_2$
R_2	$1/D$	0	$1/D$	$3/R_3$	$4/R_3$
R_3	-	$1/D$	0	$2/D$	$3/D$
R_4	$5/D$	-	$2/D$	0	-
R_5	-	-	$3/D$	-	0

$R_2 \rightarrow R_1.$


RIP Example (5)



	R_1	R_2	R_3	R_4	R_5
R_1	0	$1/D$	$2/R_2$	$4/R_2$	$5/R_2$
R_2	$1/D$	0	$1/D$	$3/R_3$	$4/R_3$
R_3	-	$1/D$	0	$2/D$	$3/D$
R_4	$5/D$	$3/R_3$	$2/D$	0	$5/R_3$
R_5	-	-	$3/D$	-	0

$R_3 \rightarrow R_4.$


RIP Example (6)



	R_1	R_2	R_3	R_4	R_5
R_1	0	$1/D$	$2/R_2$	$4/R_2$	$5/R_2$
R_2	$1/D$	0	$1/D$	$3/R_3$	$4/R_3$
R_3	-	$1/D$	0	$2/D$	$3/D$
R_4	$5/D$	$3/R_3$	$2/D$	0	$5/R_3$
R_5	-	<u>$4/R_3$</u>	$3/D$	<u>$5/R_3$</u>	0

$R_3 \rightarrow R_5.$


RIP Example (7)



	R_1	R_2	R_3	R_4	R_5
R_1	0	$1/D$	$2/R_2$	$4/R_2$	$5/R_2$
R_2	$1/D$	0	$1/D$	$3/R_3$	$4/R_3$
R_3	<u>$2/R_2$</u>	$1/D$	0	$2/D$	$3/D$
R_4	$5/D$	$3/R_3$	$2/D$	0	$5/R_3$
R_5	-	$4/R_3$	$3/D$	$5/R_3$	0

$R_2 \rightarrow R_3.$

RIP Example (8)



	R_1	R_2	R_3	R_4	R_5
R_1	0	1/D	2/R ₂	4/R ₂	5/R ₂
R_2	1/D	0	1/D	3/R ₃	4/R ₃
R_3	2/R ₂	1/D	0	2/D	3/D
R_4	5/D	3/R ₃	2/D	0	5/R ₃
R_5	5/R ₃	4/R ₃	3/D	5/R ₃	0

$R_3 \rightarrow R_5.$

Outline

1 Internet Routing

2 The Bellman-Ford Algorithm

3 Arbitrage


Basic Idea

- The algorithm is named after Richard Bellman and Lester Ford, Jr., who published the algorithm in 1958 and 1956.
- It is also called Bellman-Ford-Moore because Edward F. Moore also published the algorithm in 1957.
- We initialize the graph with estimates for the shortest paths from the starting node.
- The estimate for the starting node is 0, for all other nodes it is ∞ .
- We check all the edges of the graph and we update (relax) the estimates for all nodes.
- We repeat the process $|V| - 1$ times, where $|V|$ is the number of vertices in the graph.


Bellman-Ford and RIP

- RIP is a distributed version of the Bellman-Ford algorithm.
- Each time a node receives a packet, it updates its state of knowledge regarding shortest paths, incorporating paths with ever-increasing number of edges.
- After some time, it will have update its state of knowledge for shortest paths containing up to $|V| - 1$ edges.


A Traffic Grid




Bellman-Ford Example (1)




Bellman-Ford Example (2)




Bellman-Ford Example (3)




Bellman-Ford Example (4)




Bellman-Ford Example (5)




Bellman-Ford Example (6)



Bellman-Ford Example (7)



Bellman-Ford Example (8)



The Algorithm

Algorithm: Bellman-Ford.

$\text{BellmanFord}(G, s) \rightarrow (\text{pred}, \text{dist})$

Input: $G = (V, E)$, a graph

s , the starting node

Output: pred , an array of size $|V|$ such that $\text{pred}[i]$ is the predecessor of node i in the shortest path from s

dist , an array of size $|V|$ such that $\text{dist}[i]$ is the length of the shortest path calculated from node s to i

```
1   $\text{pred} \leftarrow \text{CreateArray}(|V|)$ 
2   $\text{dist} \leftarrow \text{CreateArray}(|V|)$ 
3  foreach  $v$  in  $V$  do
4       $\text{pred}[v] \leftarrow -1$ 
5      if  $v \neq s$  then
6           $\text{dist}[v] \leftarrow \infty$ 
7      else
8           $\text{dist}[v] \leftarrow 0$ 
9  for  $i \leftarrow 0$  to  $|V|$  do
10     foreach  $(u, v)$  in  $E$  do
11         if  $\text{dist}[v] > \text{dist}[u] + \text{Weight}(G, u, v)$  then
12              $\text{dist}[v] \leftarrow \text{dist}[u] + \text{Weight}(G, u, v)$ 
13              $\text{pred}[v] \leftarrow u$ 
14 return  $(\text{pred}, \text{dist})$ 
```

Algorithm Complexity

- The initialization requires $|V|$ iterations.
- Initialization operations require constant time, i.e., $O(1)$.
- So the initialization requires time $O(|V|)$.
- The loop in lines 9–13 is repeated $|V| - 1$ times.
- Each time we check all edges, requiring time $O((|V| - 1)|E|) = O(|V||E|)$.
- The total time required by the algorithm is $O(|V| + |V||E|)$, that is, $O(|V||E|)$.

Improvements

- In its present form, the algorithm checks in each iteration all edges.
- However, in the i th iteration the algorithm finds the shortest paths consisting of up to i edges.
- So, in the $(i + 1)$ th iteration the algorithm will only need to check the edges of the nodes that were updated in the i iteration.
- Therefore, we can keep these nodes in a queue and check only them.

Queue-based Bellman-Ford

Algorithm: Queue-based Bellman-Ford.

`BellmanFordQueue(G, s) → ($pred, dist$)`


Input: $G = (V, E)$, a graph

s , the starting node

Output: $pred$, an array of size $|V|$ such that $pred[i]$ is the predecessor of node i in the shortest path from s
 $dist$, an array of size $|V|$ such that $dist[i]$ is the length of the shortest path calculated from node s to i


```
1  inqueue ← CreateArray(|V|)  
2   $Q \leftarrow$  CreateQueue()  
3  foreach  $v$  in  $V$  do  
4       $pred[v] \leftarrow -1$   
5      if  $v \neq s$  then  
6           $dist[v] \leftarrow \infty$   
7          inqueue[ $v$ ] ← FALSE  
8      else  
9           $dist[v] \leftarrow 0$   
10     Enqueue( $Q, s$ )  
11     inqueue[ $s$ ] ← TRUE  
12     while Size( $Q$ ) ≠ 0 do  
13          $u \leftarrow$  Dequeue( $Q$ )  
14         inqueue[ $u$ ] ← FALSE  
15         foreach  $v$  in AdjacencyList( $G, u$ ) do  
16             if  $dist[v] > dist[u] + \text{Weight}(G, u, v)$  then  
17                  $dist[v] \leftarrow dist[u] + \text{Weight}(G, u, v)$   
18                  $pred[v] \leftarrow u$   
19                 if not inqueue[ $v$ ] then  
20                     Enqueue( $Q, v$ )  
21                     inqueue[ $v$ ] ← TRUE  
22     return ( $pred, dist$ )
```

Bellman-Ford with Queue Example (1)




(0, 0)

Bellman-Ford with Queue Example (2)




(0, 1)	(1, 0)
--------	--------

Bellman-Ford with Queue Example (3)




(0, 2)	(1, 1)	(2, 0)
--------	--------	--------

Bellman-Ford with Queue Example (4)




Bellman-Ford with Queue Example (5)




(1, 3)	(2, 2)	(3, 1)
--------	--------	--------

Bellman-Ford with Queue Example (6)




(2, 3)	(2, 1)	(3, 2)
--------	--------	--------


Bellman-Ford with Queue Example (7)



Bellman-Ford with Queue Example (8)




Bellman-Ford with Queue Example (9)




Advantages of Bellman-Ford

- The use of the queue does not improve the worst case performance of Bellman-Ford, which remains $O(|V||E|)$.
- That is worse than the time required by Dijkstra's algorithm, which is $O(|E|\lg |V|)$.
- However, Bellman-Ford can be used with negative weights, while Dijkstra's algorithm cannot.

Graph with Negative Weights




Bellman-Ford with Negative Weights Example (1)




0

Bellman Ford with Negative Weights Example (2)




1	2	3
---	---	---

Bellman-Ford with Negative Weights Example (3)




Bellman-Ford with Negative Weights Example (4)




3


Bellman-Ford with Negative Weights Example (5)



Graphs with Negative Cycles




Bellman-Ford with Negative Cycles Example (1)




0

Bellman-Ford with Negative Cycles Example (2)




1

Bellman-Ford with Negative Cycles Example (3)




2


Bellman-Ford with Negative Cycles Example (4)



Bellman-Ford with Negative Cycles Example (5)



Bellman-Ford with Negative Cycles Example (6)



Prevention

- The first Bellman-Ford implementation will stop anyway after $|V| - 1$ repetitions of the relaxation loop.
- To make sure that Bellman-Ford stops when implemented with a queue, we must add the proper check, otherwise it will never stop.


Bellman-Ford with Negative Cycles

- We can detect a cycle if we find a path with more than $|V| - 1$ edges.
- In Bellman-Ford, we start with paths of zero edges from s .
- Then we find the paths including the neighbors of s .
- And then the paths including the neighbors of the neighbors of s , and so on.
- So we can stop after $|V| - 1$ iterations, in each of which we process one set of neighbors.


Sentinel Value

- We need to know when we finish processing one set of neighbors.
- To do that we can use a special *sentinel* value in the queue.
- In general, sentinel values are invalid values that signal some specific event.
- In our case, we will use the number $|V|$ as the sentinel value.


Bellman-Ford with Sentinel Example (1)




Bellman-Ford with Sentinel Example (2)




Bellman-Ford with Sentinel Example (3)




Bellman-Ford with Sentinel Example (4)




Bellman-Ford with Sentinel Example (5)




Bellman-Ford with Sentinel Example (6)




Bellman-Ford with Sentinel Example (7)




Bellman-Ford with Sentinel Example (8)




Bellman-Ford with Sentinel Example (9)



Bellman-Ford with Sentinel Example (10)



Bellman-Ford with Sentinel Example (11)



Queue-based Bellman-Ford with Negative Cycles Algorithm

Algorithm: Queue-based Bellman-Ford with negative cycles.

BellmanFordQueueNC($G, s \rightarrow (\text{pred}, \text{dist}, \text{ncc})$)

Input: $G = (V, E)$, a graph

s , the starting node

Output: pred , an array of size $|V|$ such that $\text{pred}[i]$ is the predecessor of node i in the shortest path from s
 dist , an array of size $|V|$ such that $\text{dist}[i]$ is the length of the shortest path calculated from node s to i
 ncc , TRUE if there is no negative cycle, FALSE otherwise.

```
1  inqueue ← CreateArray(|V|)  
2  Q ← CreateQueue()  
3  foreach  $v \in V$  do  
4       $\text{pred}[v] \leftarrow |V|$   
5      if  $v \neq s$  then  
6           $\text{dist}[v] \leftarrow \infty$   
7          inqueue[v] ← FALSE  
8      else  
9           $\text{dist}[v] \leftarrow 0$   
10 Enqueue(Q, s)  
11 inqueue[s] ← TRUE  
12 Enqueue(Q, |V|)  
13 i ← 0  
14 while Size(Q) ≠ 1 and  $i < |V|$  do  
15      $u \leftarrow \text{Dequeue}(Q)$   
16     if  $u = |V|$  then  
17          $i \leftarrow i + 1$   
18         Enqueue(Q, |V|)  
19     else  
20         inqueue[u] ← FALSE  
21         foreach  $v \in \text{AdjacencyList}(G, u)$  do  
22             if  $\text{dist}[v] > \text{dist}[u] + \text{Weight}(G, u, v)$  then  
23                  $\text{dist}[v] \leftarrow \text{dist}[u] + \text{Weight}(G, u, v)$   
24                  $\text{pred}[v] \leftarrow u$   
25                 if not inqueue[v] then  
26                     Enqueue(Q, v)  
27                     inqueue[v] ← TRUE  
28 return ( $\text{pred}, \text{dist}, i < |V|$ )
```

Outline

1 Internet Routing

2 The Bellman-Ford Algorithm

3 Arbitrage

Top Ten Most Traded Currencies by Value, April 2013

Rank	Currency	Code	% daily share
1	United States dollar	USD	87.0%
2	European Union euro	EUR	33.4%
3	Japanese yen	JPY	23.0%
4	United Kingdom pound sterling	GPB	11.8%
5	Australian dollar	AUD	8.6%
6	Swiss franc	CHF	5.2%
7	Canadian dollar	CAD	4.6%
8	Mexican peso	MXN	2.5%
9	Chinese yuan	CNY	2.2%
10	New Zealand dollar	NZD	2.0%

The total percentage (including currencies below the top ten) adds up to 200% as trades are between pairs of currencies.

Arbitrage Example

- Suppose that in London the exchange rate of euros and U.S. dollars is:

$$\text{€} 1 = \$ 1.37$$

- Suppose also that in New York City the exchange rate of U.S. dollars and euros is:

$$\$ 1 = \text{€} 0.74$$

- Then a trader can buy in London using € 1,000,000 a total of \$ 1,370,000 and send them to New York City.
- There the trader converts them back to euros:

$$\$ 1,370,000 \times 0.74 = \text{€} 1,013,800$$

- The trader made a risk-free profit of € 13,800.

How is Arbitrage Resolved

- If the exchange rate between U.S. dollars and euros in New York city is:

$$\$1 = €0.73$$

- Then with the same trades, the trader will have:

$$\$1,370,000 \times 0.73 = €1,000,100$$

- So the profit will be just €100.
- As soon as arbitrage is detected, traders will rush to take advantage of it and the exchange rates will move towards equilibrium.

General Scheme

- From U.S. dollars to Australian dollars, to Canadian dollars, and then back to U.S. Dollars:

$$1 \times (\text{USD} \rightarrow \text{AUD}) \times (\text{AUD} \rightarrow \text{CAD}) \times (\text{CAD} \rightarrow \text{USD})$$


- In general, we want to find a sequence of currency conversions:

$$a = 1 \times (c_1 \rightarrow c_2) \times (c_2 \rightarrow c_3) \times \cdots \times (c_n \rightarrow c_1)$$

where:

$$a > 1$$

Cross-currency Rates Graph



Cross-currency Graph Adjacency Matrix

	USD	EUR	JPY	GBP	AUD	CHF	CAD	MXN	CNY	NZD
USD	1	1.3744	0.009766	1.6625	0.9262	1.1275	0.9066	0.07652	0.1623	0.8676
EUR	0.7276	1	0.007106	1.2097	0.6739	0.8204	0.6596	0.05568	0.1181	0.6313
JPY	102.405	140.743	1	170.248	94.8421	115.455	92.8369	7.836	16.6178	88.8463
GBP	0.6016	0.8268	0.005875	1	0.5572	0.6782	0.5454	0.04603	0.09762	0.5219
AUD	1.0799	1.4842	0.010546	1.7953	1	1.2176	0.979	0.08263	0.1752	0.9369
CHF	0.8871	1.2192	0.008663	1.4748	0.8216	1	0.8042	0.06788	0.144	0.7696
CAD	1.1033	1.5163	0.010775	1.8342	1.0218	1.2439	1	0.08442	0.179	0.9572
MXN	13.0763	17.9724	0.1277	21.7397	12.1111	14.7435	11.8545	1	2.122	11.345
CNY	6.167	8.4761	0.06023	10.2528	5.7118	6.9533	5.5908	0.4719	1	5.3505
NZD	1.153	1.5846	0.01126	1.9168	1.0678	1.2999	1.0452	0.08822	0.1871	1

Arbitrage Detection

Problem

To spot an arbitrage opportunity we must find in the cross-currency rates graph a cycle with weights:

$$w_1, w_2, \dots, w_n$$

such that:

$$w_1 w_2 \dots w_n > 1$$

Arbitrage Detection Method

- ① If the weight between two vertices u and v is $w(u, v)$, we substitute it by the value $w'(u, v) = -\log w(u, v)$.
- ② We execute the amended Bellman-Ford to detect a negative cycle in the resulting graph.
- ③ If the algorithm reports the existence of a negative cycle, the cycle corresponds to an arbitrage opportunity.

Explanation (1)

- Suppose that the cycle consists of n edges, then the sum of the cycle weights will be:

$$w'_1 + w'_2 + \dots + w'_n < 0$$

where:

$$w'_1, w'_2, \dots, w'_n$$

are the weights.

- We have:

$$w'_1 + w'_2 + \dots + w'_n = -\log w_1 - \log w_2 - \dots - \log w_n$$

which means that in the negative cycle:

$$-\log w_1 - \log w_2 - \dots - \log w_n < 0$$

- A basic property of logarithms is that $\log(xy)$ is equal to $\log x + \log y$ and $\log(1/x)$ is equal to $\log(1/x) + \log(1/y) = -\log x - \log y$.

Explanation (2)

- So the last inequality becomes:

$$\log\left(\frac{1}{w_1} \frac{1}{w_2} \cdots \frac{1}{w_n}\right) < 0$$

- Exponentiating to remove the logarithm we get:

$$\frac{1}{w_1} \frac{1}{w_2} \cdots \frac{1}{w_n} < 10^0 = 1$$

- But this is equivalent to:

$$w_1 w_2 \cdots w_n > 1$$

That is what we wanted to find in the first place; a cyclical path with weights whose product is greater than one.