

Mathematical Concepts (2/2)

Functions, Minimization, Gradient

Fundamentals of Artificial Intelligence

Instructor: Chenhui Chu

Email: chu@i.kyoto-u.ac.jp

Teaching Assistant: Youyuan Lin

E-mail: youyuan@nlp.ist.i.kyoto-u.ac.jp

Schedule

- 1. Overview of AI and this Course (4/14)
- 2. Introduction to Python (4/21)
- 3, **4. Mathematics Concepts I, II** (4/28, **5/12**)
- 5, 6. Regression I, II (5/19, 5/26)
- 7. Classification (6/2)
- 8. Introduction to Neural Networks (6/9)
- 9. Neural Networks Architecture and Backpropagation (6/16)
- 10. Fully Connected Layers (6/23)
- 11, 12, 13. Computer Vision I, II, III (6/30, 7/7, 7/14)
- 14. Natural Language Processing (7/17)

Overview of This Course

11, 12, 13. Computer Vision
I, II, III

14. Natural language
processing

Deep Learning Applications

8. Neural network
Introduction

9. Architecture and
Backpropagation

10. Feedforward
neural networks

Deep Learning

5. Simple linear
regression

6. Multiple linear
regression

7. Classification

Basic Supervised Machine Learning

2. Python

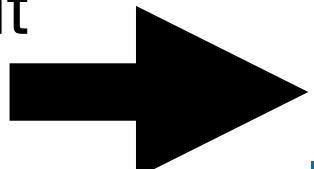
3, 4. Mathematics Concepts I, II

Fundamental of Machine Learning

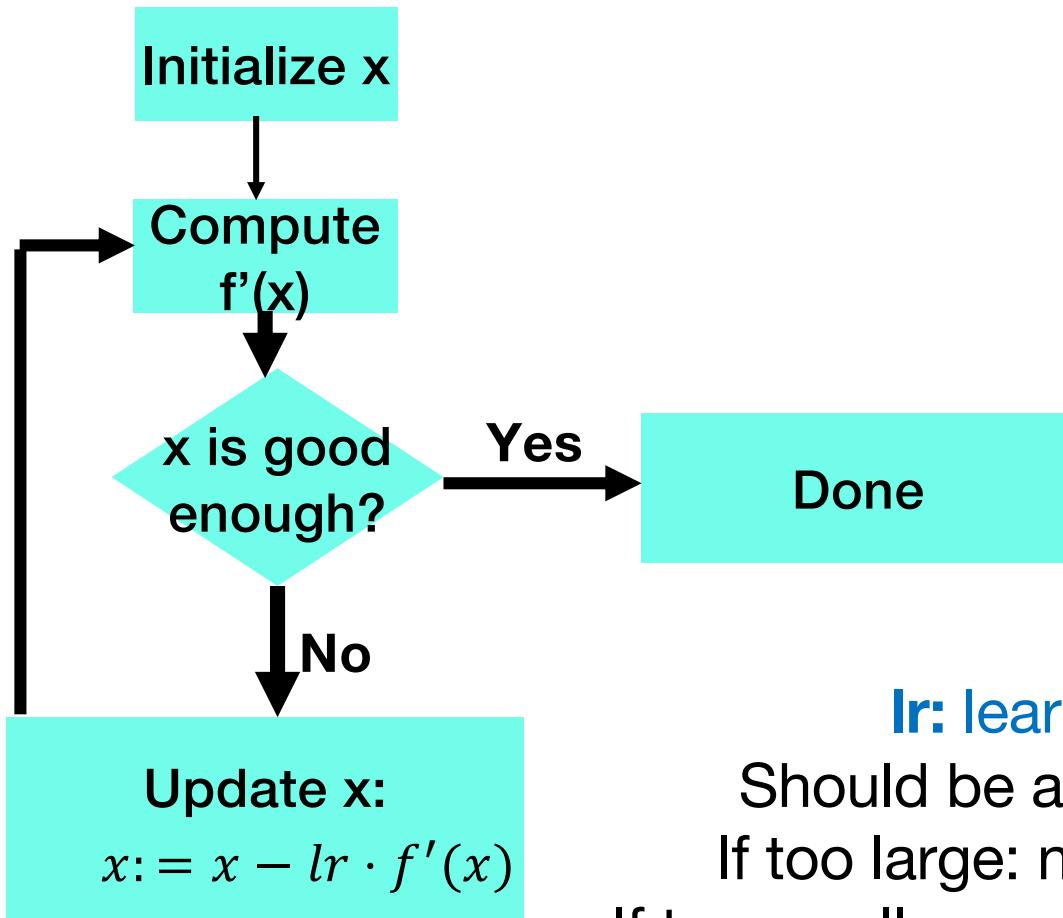
Gradient Descent Algorithm (1/4)

- This suggests some procedure for finding a minimum:
 - Start at any x (e.g., $x = 0$)
 - Compute $f'(x)$
 - If $f'(x) > 0$: Decrease x a bit
 - If $f'(x) < 0$: Increase x a bit
 - Repeat

Gradient Descent Algorithm (2/4)

- This suggests some procedure for finding a minimum:
 - Start at any x (eg. $x = 0$)
 - Compute $f'(x)$
 - If $f'(x) > 0$: Decrease x a bit
 - If $f'(x) < 0$: Increase x a bit
 - Repeat
- 
- In practice, we do this:**
- $$x := x - lr \cdot f'(x)$$
- lr:** learning rate
Should be a positive value
If too large: no convergence
If too small: very slow convergence

Gradient Descent Algorithm (3/4)



$$f(x) = x^2 - x$$

$$f'(x) = 2x - 1$$

$$\underset{x}{\operatorname{argmin}} f(x) = 0.5$$

$lr = 0.2$

$x = 0$
 $f'(x) = -1$

$x = 0.2$
 $f'(x) = -0.6$

$x = 0.32$
 $f'(x) = -0.36$

$x = 0.392$

...

...

$x = 0.493$
 $f'(x) = -0.014$

STOP?

lr: learning rate

Should be a positive value

If too large: no convergence

If too small: very slow convergence

Gradient Descent Algorithm (4/4)

- Gradient descent works well **even** when we have functions of millions of variables
 - This is why it is so useful for Machine Learning and Neural Networks
 - Other methods will not be practical in such settings
 - Convergence will depend on the choice of a good **learning rate**
 - In experiments, a good deal of time is often spent finding an optimal learning rate
 - **Too large** learning rate: **no** convergence (ie. the system learns nothing)
 - **Too small** learning rate: **slow** convergence (ie. the system takes a long time to learn)

Minimizing a Function of Several Variables

Functions of Several Variables

- A function of several variables is just that: a function which has several variables

$$f: \mathbb{R}^3 \rightarrow \mathbb{R}$$

$$f(x, y, z) = (x - y)^2 + z^2 - z$$

$$f(0,0,0) = 0$$

$$f(1,2,3) = 7$$

$$f(-1,2,2) = 11$$

$$f(0,1,1) = ?$$

$$f(2,2,0) = ?$$

- Like before, we want to find its minimum:

$$\operatorname{argmin}_{x,y,z} f(x, y, z) = (0,0,0.5)$$

Parameterized Functions (1/5)

- By fixing one of the variable, we can obtain a function with one less variable

Function of 3 variables: $f(x, y, z) = (x - y)^2 + z^2 - z$



Fixing z:

$$z = 2$$



$$f(x, y, 2) = (x - y)^2 + 4 - 2$$

Function of 2 variables: $f(x, y) = (x - y)^2 + 2$

Parameterized Functions (2/5)

- By fixing one of the variable, we can obtain a function with one less variable

Function of 3 variables: $f(x, y, z) = (x - y)^2 + z^2 - z$



Fixing y:

$$y = 2$$



$$f(x, 2, z) = (x - 2)^2 + z^2 - z$$

Function of 2 variables: $f(x, z) = (x - 2)^2 + z^2 - z$

Parameterized Functions (3/5)

- By fixing one of the variable, we can obtain a function with one less variable

Function of 3 variables: $f(x, y, z) = (x - y)^2 + z^2 - z$



Fixing y and z:

$$y = 2 \quad z = 3$$



$$f(x, 2, 3) = (x - 2)^2 + 9 - 3$$

Function of 1 variable: $f(x) = (x - 2)^2 + 6$

Parameterized Functions (4/5)

- Therefore, in this case, variables y and z can be used to describe a “family” of functions. We say they parameterize the

Function of 3 variables: $f(x, y, z) = (x - y)^2 + z^2 - z$

Fixing y and z :

$$y = 2 \qquad z = 3$$

For each value of y and z , we have one function of one variable

$$f(x, 2, 3) = (x - 2)^2 + 9 - 3$$

Function of 1 variable: $f(x) = (x - 2)^2 + 6$

Parameterized Functions (5/5)

- In such a case, we will say that f is a function *parameterized* by y and z .
- And we note the parameters separately, as subscripts

Function of 3 variables: $f_{y,z}(x) = (x - y)^2 + z^2 - z$



Fixing y and z :

$$y = 2 \qquad z = 3$$



$$f_{2,3}(x) = (x - 2)^2 + 9 - 3$$

For each value of y and z , we have one function of one variable

Function of 1 variable: $f_{2,3}(x) = (x - 2)^2 + 6$



$$f_{0,0}(x) = x^2$$

$$f_{0,2}(x) = x^2 + 2$$

Partial Derivatives (1/4)

- What is the equivalent of our “high school” derivatives when we have several variables?
- One part of the answer is ***partial derivatives***
- Partial derivatives are computed by choosing one variable and fixing the others
- In other words, we see the function of several variables as a parameterized function of one variable

Partial Derivatives (2/4)

- What is the equivalent of our “high school” derivatives when we have several variables?
- One part of the answer is ***partial derivatives***
- Partial derivatives are computed by choosing one variable and fixing the others
- In other words, we see the function of several variables as a parameterized function of one variable $f(x, y, z) = (x - y)^2 + z^2 - z$
- Indeed, if we choose y , and fix x and z , we can see $f(x, y, z)$ as a function of one variable and compute its derivative

$$\frac{\partial f}{\partial x} = 2(x - y)$$

$$\frac{\partial f}{\partial y} = 2(y - x)$$

$$\frac{\partial f}{\partial z} = 2z - 1$$

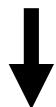
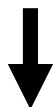
Partial Derivatives (3/4)

$$f(x, y, z) = (x - y)^2 + z^2 - z$$

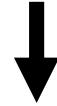
↓

Fix other variables

$$f_{y,z}(x) = (x - y)^2 + z^2 - z \quad | \quad f_{x,z}(y) = (x - y)^2 + z^2 - z \quad | \quad f_{x,y}(z) = (x - y)^2 + z^2 - z$$



Compute derivative



$$f'_{y,z}(x) = 2(x - y)$$

$$f'_{x,z}(y) = 2(y - x)$$

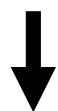
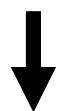
$$f'_{x,y}(z) = 2z - 1$$

Partial Derivatives (4/4)

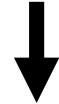
$$f(x, y, z) = (x - y)^2 + z^2 - z$$

Fix other variables

$$f_{y,z}(x) = (x - y)^2 + z^2 - z \quad | \quad f_{x,z}(y) = (x - y)^2 + z^2 - z \quad | \quad f_{x,y}(z) = (x - y)^2 + z^2 - z$$



Compute derivative



$$f'_{y,z}(x) = 2(x - y)$$

$$f'_{x,z}(y) = 2(y - x)$$

$$f'_{x,y}(z) = 2z - 1$$

In practice, we use this notation for partial derivatives:

$$\frac{\partial f}{\partial x} = 2(x - y)$$

$$\frac{\partial f}{\partial y} = 2(y - x)$$

$$\frac{\partial f}{\partial z} = 2z - 1$$

Compute the Partial Derivatives

$$f(x, y, z) = (x - y)^2 + z^2 - z$$

$$\frac{\partial f}{\partial x} =$$

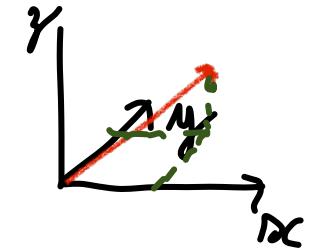
$$\frac{\partial f}{\partial y} =$$

$$\frac{\partial f}{\partial z} =$$

Vectors (1/3)

- What are vectors?
- You probably have used vectors in physics classes to represent force and speed
 - 3-dimensional vectors: [2.3, 4.5, -1]
- In machine learning, we also use them a lot
- Except that they can have more than 3 dimensions
 - 5-dimensional vector: [-1, 3, 4.1, 5.2, 4]
 - We often note the set of all n-dimensional vectors \mathbb{R}^n

$$[1.2, 1.4, 1, -1, -1] \in \mathbb{R}^5$$



Vectors (2/3)

- For now, we only need to know the following about vectors:

- A n-dimensional vector is a list of n numbers

- We can add 2 vectors (if they have the same dimension)

$$[2.1, 3.4, 1.1, 3.2] + [-1, 2.1, 3.1, -2] = [1.1, 5.5, 4.2, 1.2]$$

$$[2.1, 3.4] + [-1, 2.1, 3.1, -2] = \text{X}$$

- We can multiply a vector by a number

$$0.5 \times [2, 3, -1, -2] = [1, 1.5, -0.5, -1]$$

Vectors (3/3)

- We will usually denote a vector by a letter with an arrow on it: \vec{x}
- We denote the i^{th} component of \vec{x} by x_i
- If $\vec{x} = [1, 2.2, -1, 4]$
 - Then we have $x_0=1, x_1=2.2, x_2=-1, x_3=4$

Vectors and Numpy

- In Python, Numpy arrays are a convenient way to represent vectors

$$\vec{x} = [1, 5, -2, 0.5]$$

- `x = np.array([1, 5, -2, 0.5])`
- `x[0] == x0 == 1`
- `x[1] == x1 == 5`
- Vector operations: `x+0.5*y`

```
In [981]: x = np.array([1, 5, -2, 0.5])
y = np.array([2,2,10,10])
z = np.array([3, -3, 0])
print(x)
print(y)
print(z)

[ 1.   5.  -2.   0.5]
[ 2.   2.  10.  10.]
[ 3.  -3.   0.]
```

```
In [982]: print(len(x), len(y), len(z))

4 4 3
```

```
In [983]: print(x[1], y[2], z[0], y[0])

5.0 10 3 2
```

```
In [984]: print(x+y)

[ 3.   7.   8.  10.5]
```

```
In [985]: print(x+0.5*y)

[2.   6.   3.   5.5]
```

```
In [986]: print(y+z)
```

```
-----  
ValueError  
    /usr/local/lib/python3.6/dist-packages/numpy/core/_asarray.py:143: UserWarning:
```

Vectors and Multivariate Functions

- For now, we have represented the variables of a multivariate function with the letters x, y, z as in: $f(x, y, z) = (x - y)^2 + z^2 - z$
- In practice, we can have any number of variables. So it is more convenient to use: x_0 (instead of x) , x_1 (instead of y), x_2 (instead of z), $x_3 \dots x_n$ (if we need more than 3 variables) $f(x_0, x_1, x_2) = (x_0 - x_1)^2 + x_2^2 - x_2$
- We can also use a vectorial notation to represent all of the variables as one vector variable: $\vec{x} = [x_0, x_1, x_2]$ $f(\vec{x}) = (x_0 - x_1)^2 + x_2^2 - x_2$
- So, keep in mind that the 3 following expressions actually refer to the same function: $f(x, y, z) = (x - y)^2 + z^2 - z$
 $f(x_0, x_1, x_2) = (x_0 - x_1)^2 + x_2^2 - x_2$
 $f(\vec{x}) = (x_0 - x_1)^2 + x_2^2 - x_2$

Gradient (1/2)

- The partial derivatives become the component of a vector we call the ***gradient***

$$\text{grad} \cdot f(x, y, z) = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right]$$

- For example:
- $$f(x, y, z) = (x - y)^2 + z^2 - z$$
- $$\text{grad} \cdot f(x, y, z) = [2(x - y), 2(y - x), 2z - 1]$$

Gradient (2/2)

$$\text{grad} \cdot f(x, y, z) = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right]$$

- In this case, the function has 3 variables. Therefore the gradient is a vector of size 3
- If the gradient has n variables, it is a vector of size n
- More precisely, the gradient of f is itself a function that return a vector

$$f: \mathbb{R}^n \rightarrow \mathbb{R}$$

$$f(x_1, x_2, \dots, x_n)$$

$$\text{grad} \cdot f: \mathbb{R}^n \rightarrow \mathbb{R}^n$$

$$\text{grad} \cdot f(x_1, x_2, \dots, x_n) = [g_1, \dots, g_n]$$

Interpreting the Gradient

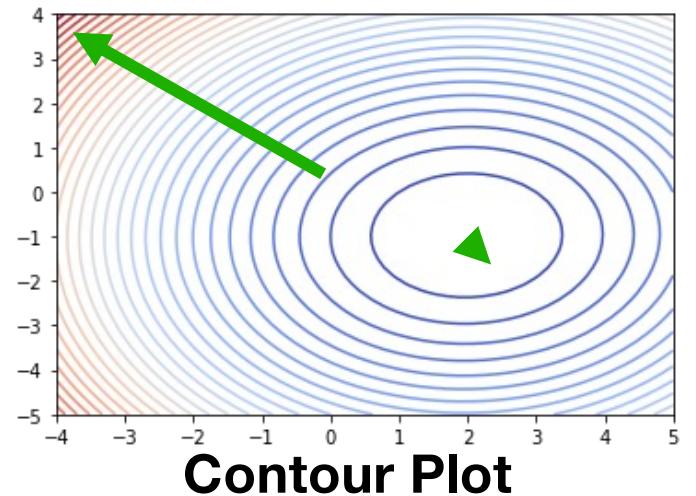
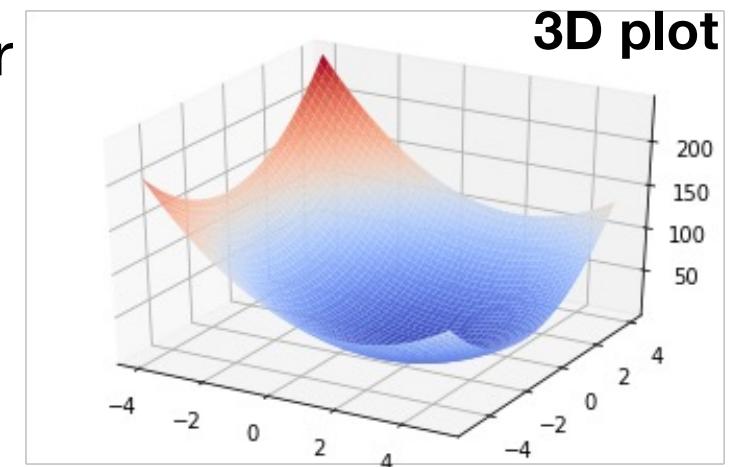
- At a given point, the gradient is the direction for which the value of the function increase fastest
- Therefore, in general, it points in the direction **opposite to the minimum**

$$f(x, y) = 4(x - 2)^2 + 4(y + 1)^2 - 0.1xy$$

$$\text{grad} \cdot f(x, y) = [8(x - 2) - 0.1y, 8(y + 1) - 0.1x]$$

$$\text{grad} \cdot f(0, 0) = [-16, 16]$$

$$\text{grad} \cdot f(2, -1) = [0.1, -0.2]$$



Gradient Descent

- Because we know that the gradient point in a direction opposite to the minimum, we can use the same idea as in the case of one variable

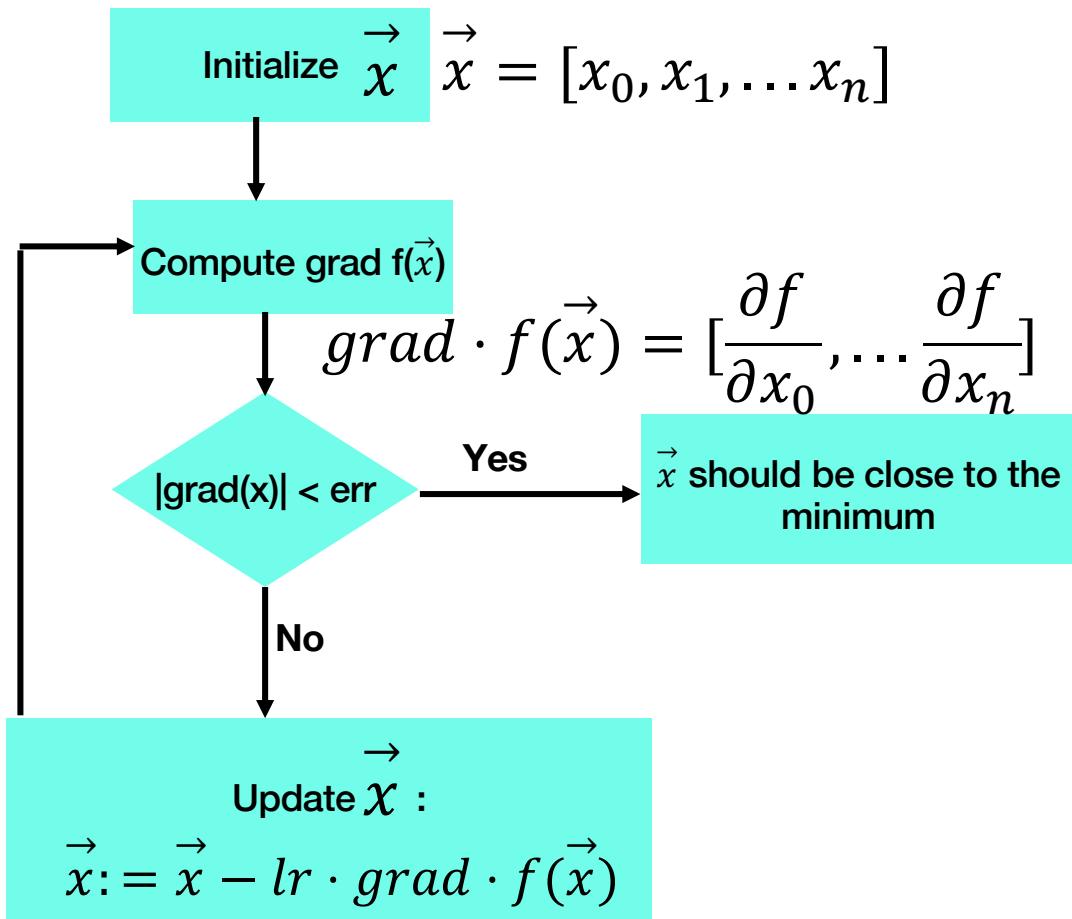
One variable:

$$x := x - lr \cdot f'(x)$$

Multiple variables:

$$\vec{x} := \vec{x} - lr \cdot \text{grad} \cdot f(\vec{x})$$

Gradient Descent Algorithm



$$f(\vec{x}) = (x_0 - x_1)^2 + x_2^2 - x_2$$

$$\text{grad} \cdot f(\vec{x}) = [2(x_0 - x_1), 2(x_1 - x_0), 2x_2 - 1]$$

$$lr = 0.2$$

$$\vec{x} = [0, 1, 0]$$

$$\text{grad} \cdot f(\vec{x}) = [-2, 2, -1]$$

$$\vec{x} = [0.4, 0.6, 0.2]$$

$$\text{grad} \cdot f(\vec{x}) = [-0.4, 0.4, -0.6]$$

$$\vec{x} = [0.41, 0.43, 0.51]$$

$$\text{grad} \cdot f(\vec{x}) = [-0.04, 0.04, 0.01]$$

What is the Equivalent of Second Derivative for Multivariate Functions?

$$\begin{vmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial x \partial z} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} & \frac{\partial^2 f}{\partial y \partial z} \\ \frac{\partial^2 f}{\partial x \partial z} & \frac{\partial^2 f}{\partial y \partial z} & \frac{\partial^2 f}{\partial z^2} \end{vmatrix}$$

- It is the Hessian Matrix:
- But thankfully, we will not need to use it
- But for your information, this would be the equivalent of the “High School” minimization when we have several variables:

To minimize $f(x, y, z)$:

1. Compute gradient of $f(x, y, z)$
2. Compute hessian of $f(x)$
3. Find x, y, z such that $\text{grad } f(x, y, z) = 0$
4. If hessian of $f(x, y, z)$ is definite positive then (x, y, z) is a local minimum of f

Gradient Descent Algorithm

- You can see that, in the case of the gradient descent, the algorithm is the same for univariate functions and multivariate functions
- It is a simple algorithm, but it scales very well
- There exists many variations of it:
 - Gradient Descent with Momentum
 - Stochastic Gradient Descent
 - Adagrad, Adadelta, Adam, ...

Gradient Descent with Momentum

- Compute a “gradient with momentum” at each iteration:

$$gm_t = 0.6grad \cdot f(\vec{x}) + 0.4gm_{t-1}$$

Update \vec{x} :
 $\vec{x} := \vec{x} - lr \cdot gm_t$

Stochastic Gradient Descent (1/2)

- What happens if the gradient is noisy?
- That is, we can only compute a value that is equal to the true gradient “on average”?
 - A bit like if you are drunk and trying to get home

Stochastic Gradient Descent (2/2)

- What happens if the gradient is noisy?
- That is, we can only compute a value that is equal to the true gradient “on average”?
 - A bit like if you are drunk and trying to get home
- It turns out it works.
 - But you have to decrease your learning rate over time to stabilize
 - Convergence will be slower
- Very interesting because a noisy gradient can be million times faster to compute than a “true” gradient

$$lr = \frac{lr_0}{\sqrt{(t + 1)}}$$

Optimization Libraries

- You can also minimize a function by using a specialized library
- It gives you access to more sophisticated minimization algorithms
- However, these more sophisticated algorithms do not scale as well as Gradient Descent
 - Which is one reason for Gradient Descent and its variants are still the main tool for large scale Machine Learning (In particular, Deep Learning)

```
In [103]: scipy.optimize.minimize(f_numpy, np.array([0,0]), jac=grad_f_numpy, method="L-BFGS-B")
```

```
Out[103]:      fun: 0.1969057665260197
      hess_inv: <2x2 LbfgsInvHessProduct with dtype=float64>
                  jac: array([-3.88578059e-16,  2.77555756e-17])
      message: b'CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL'
        nfev: 5
         nit: 4
      status: 0
     success: True
           x: array([ 1.9878106 , -0.97515237])
```

Google Colab Notebook

- Let us check gradient descent in practice with Google Colab notebook

<https://shorturl.at/dTCs9>

Report

- Submit Exercise 1 and 2 **in pdf** via PandA
 - Submission due: **next lecture**
 - Name the pdf file as **student id_name**.

Exercise 1

- Compute the partial derivatives of:

$$f(x, y, z) = xyz - z^2 - y^2$$

$$f(x, y, z) = e^{x+y} - \log(z)$$

Exercise 2

$$\vec{x} = [1.5, -2.0, 5]$$

$$\vec{y} = [2, 2, 10, 10]$$

$$\vec{z} = [3, -3, 0]$$

- Dimensions of $\vec{x}, \vec{y}, \vec{z}$?
- Values of x_1, y_2, z_0, y_0 ?
- Compute: $\vec{x} + \vec{y}$ $\vec{x} + 0.5 \times \vec{y}$ $\vec{y} + \vec{z}$