

# Regression II

Fundamentals of Artificial Intelligence

Instructor: Chenhui Chu

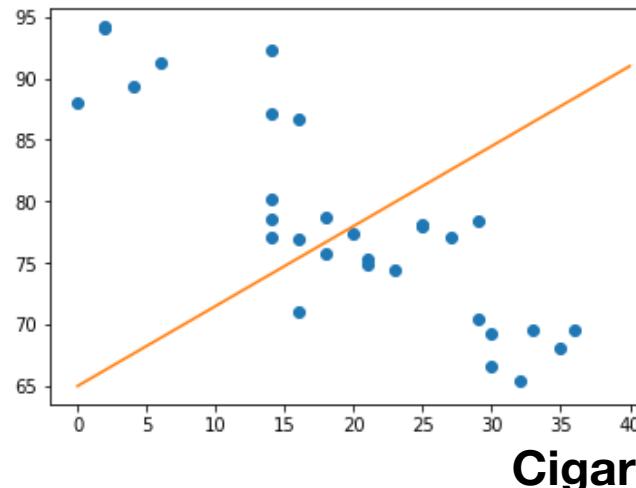
Email: chu@i.kyoto-u.ac.jp

Teaching Assistant: Youyuan Lin

E-mail: youyuan@nlp.ist.i.kyoto-u.ac.jp

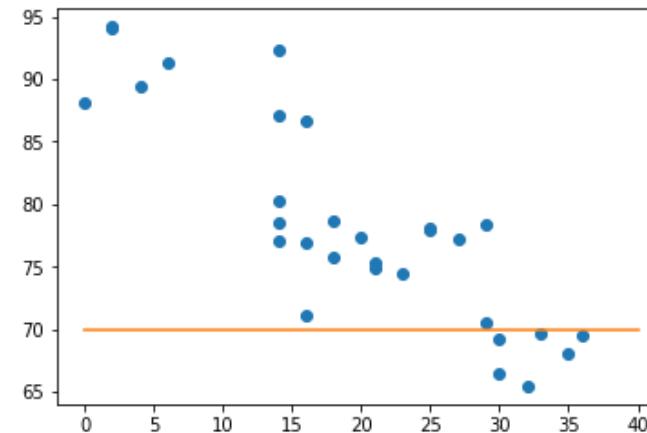
# Linear Regression

Age

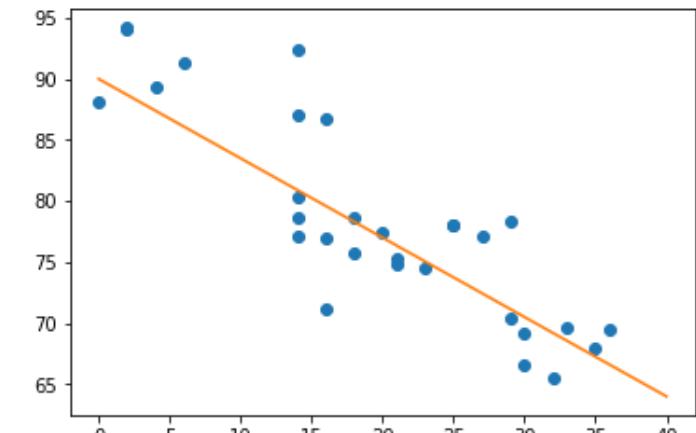


$$\theta_0 \approx 65 \quad \theta_1 \approx 0.7$$

Cigarettes



$$\theta_0 \approx 70 \quad \theta_1 \approx 0$$



$$\theta_0 \approx 90 \quad \theta_1 \approx -0.7$$

So, which one  
is best?

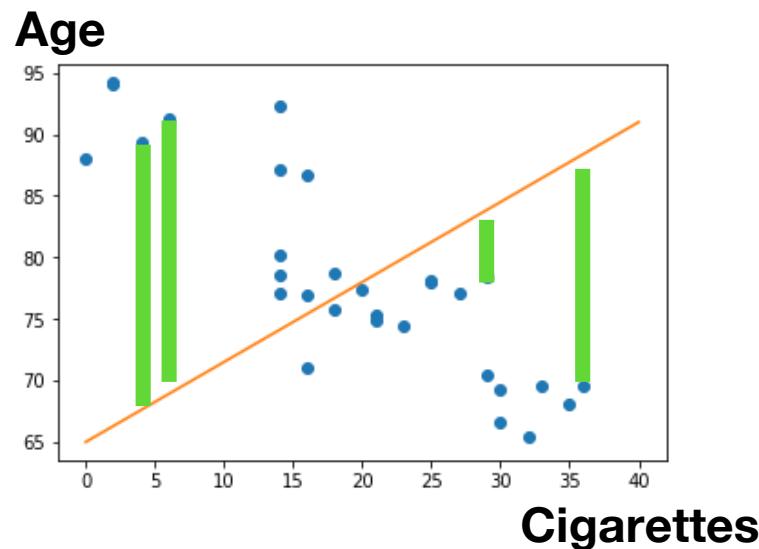


$$age = \theta_0 + \theta_1 \times cig$$

**How to know which values of the parameters  $\theta_0, \theta_1$  we should use?**

# Mean Squared Error

- Finally, we take average of the squared error for all examples



$$\theta_0 \approx 65 \quad \theta_1 \approx 0.7$$

$$\text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (\text{error}_i)^2$$

$$\text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (f(x_i) - y_i)^2$$

N: total number of examples

$x_i$ : number of cigarettes smoked by person i

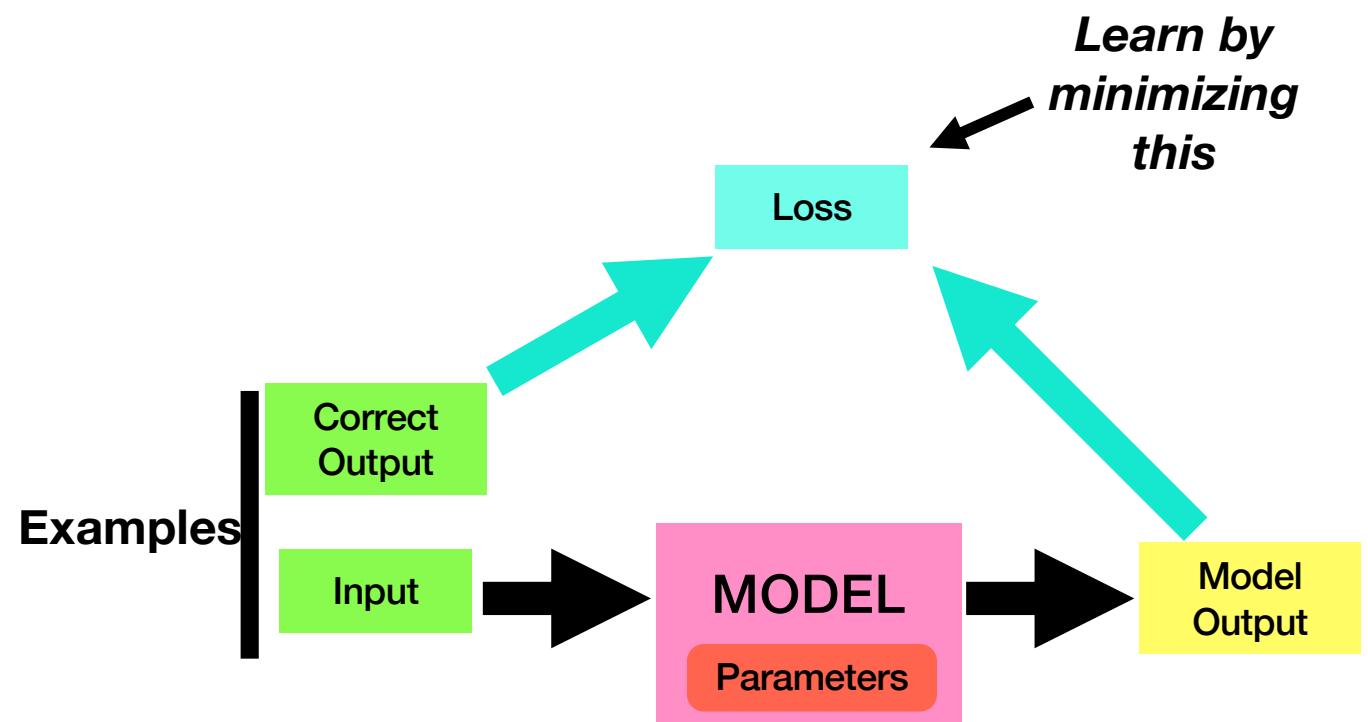
$y_i$ : age person i died

$f(x_i)$ : prediction of our model

In our case:  $f(x_i) \sim age_{\theta_0, \theta_1}(cig)$

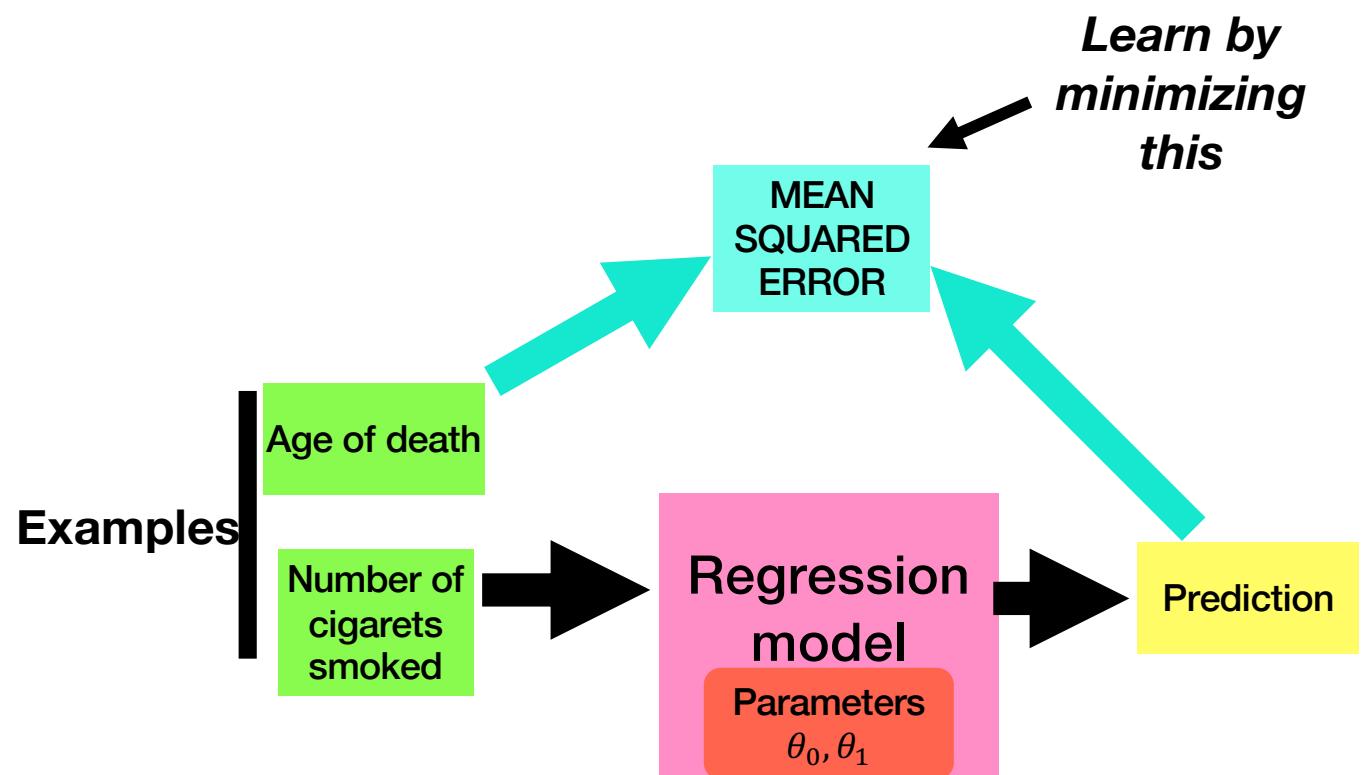
# Supervised Learning (1/2)

- In supervised learning, we usually have:
  - A **MODEL**: a “parameterized” function that takes input and produce output
  - A **Loss**: A function that computes how different the model output is from the correct output
  - **Examples** of input and correct output



# Supervised Learning (2/2)

- In supervised learning, we usually have:
  - A **MODEL**: a “parameterized” function that takes input and produce output
  - A **Loss**: A function that computes how different the model output is from the correct output
  - **Examples** of input and correct output



$$age_{\theta_0, \theta_1}(cig) = \theta_0 + \theta_1 \times cig$$

# Stochastic Gradient Descent

- Stochastic gradient descent says that we can replace the average of the gradient over all examples by the gradient given by a randomly chosen example
- Slower Convergence
- But if we have one million examples: one million times faster to compute!

$$\theta_0 := \theta_0 - lr \times \frac{2}{N} \cdot \sum_i (\theta_0 + \theta_1 \times x_i - y_i)$$



**Choose example  $i$  randomly**

$$\theta_0 := \theta_0 - lr \times 2 \cdot (\theta_0 + \theta_1 \times x_i - y_i)$$

In practice, we often average over a few examples (instead of just one). This is called mini-batch gradient descent

# Schedule

- 1. Overview of AI and this Course (4/14)
- 2. Introduction to Python (4/21)
- 3, 4. Mathematics Concepts I, II (4/28, 5/12)
- 5, **6. Regression I, II (5/19, 5/26)**
- 7. Classification (6/2)
- 8. Introduction to Neural Networks (6/9)
- 9. Neural Networks Architecture and Backpropagation (6/16)
- 10. Fully Connected Layers (6/23)
- 11, 12, 13. Computer Vision I, II, III (6/30, 7/7, 7/14)
- 14. Natural Language Processing (7/17)

# Overview of This Course

11, 12, 13. Computer Vision  
I, II, III

14. Natural language  
processing

## Deep Learning Applications

8. Neural network  
Introduction

9. Architecture and  
Backpropagation

10. Feedforward  
neural networks

## Deep Learning

5. Regression I

6. Regression II

7. Classification

## Basic Supervised Machine Learning

2. Python

3, 4. Mathematics Concepts I, II

## Fundamental of Machine Learning

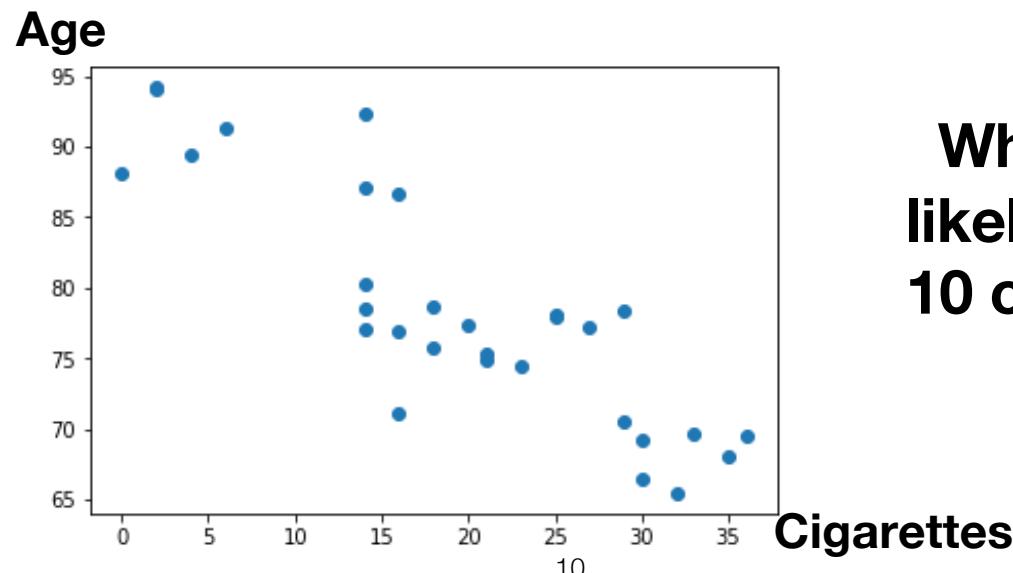
# Today

- We expand on this by considering:
  - More than one input feature
  - More complex functions
- We will have a look at the important concept of overfitting

# Other Factors?

- The number of cigarettes smoked is **not the only** important **factor** for **predicting** the age of death
  - Physical fitness, biological Sex, Wealth....

daily cigarettes	age of death
32.0	73
7.0	88
30.0	82
17.0	85
27.0	76
15.0	84
20.0	72
28.0	77
....	.....



**What age am I most likely to die if I smoke 10 cigarettes per day?**

# Adding More Information

- Let us suppose now that, for the same 30 persons, we **also** have their BMI index and know their sex
- Vocabulary Note: In Machine Learning, we often call a “feature” or “feature function” each of this piece of information about an example

	daily cigarettes	bmi	is male	age of death
0	5.0	18.5	1.0	79.8
1	9.0	45.1	0.0	56.8
2	38.0	14.2	0.0	61.4
3	12.0	48.5	1.0	37.5
4	34.0	19.2	0.0	68.4
5	5.0	38.6	0.0	69.3
6	31.0	33.8	1.0	54.8
7	25.0	33.6	1.0	63.0
8	24.0	45.2	1.0	39.3
9 ...	....	....	....	....

$$bmi = \frac{weight}{height^2}$$

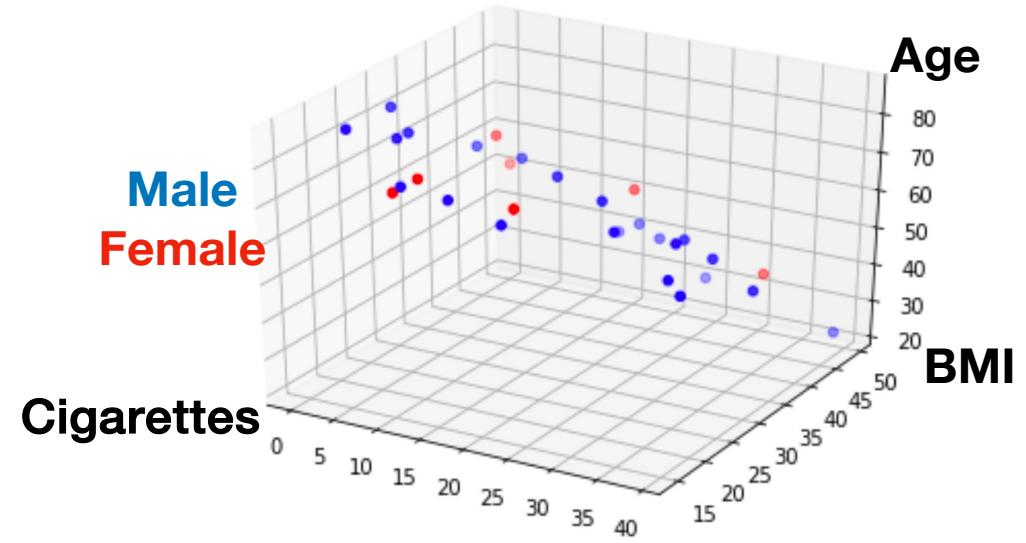
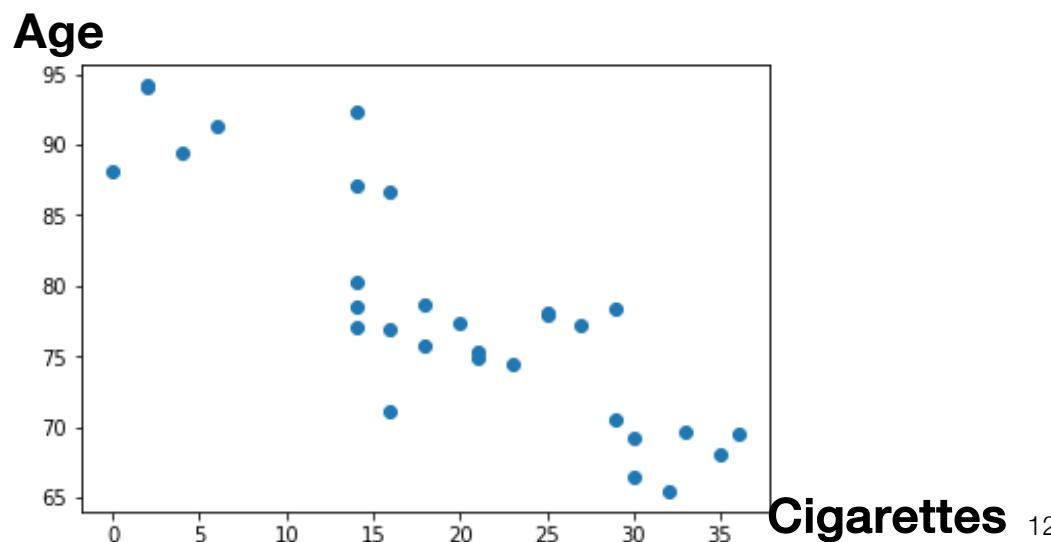
Is male: 1 if the person is a male, 0 if female

**Note: most of the time, this is how we represent “categorical data” in Machine Learning: a feature equal to 1 if the example belongs to the category in question, and equal to zero otherwise**

# Visual Representation of the Examples

	daily cigarettes	age of death
0	5.0	79.8
1	9.0	56.8
2	38.0	61.4
3	12.0	37.5
7	25.0	63.0
8	24.0	39.3
9	...	....

	daily cigarettes	bmi	is male	age of death
0	5.0	18.5	1.0	79.8
1	9.0	45.1	0.0	56.8
2	38.0	14.2	0.0	61.4
3	12.0	48.5	1.0	37.5
7	25.0	33.6	1.0	63.0
8	24.0	45.2	1.0	39.3
9	...	....	....	....



# Linear Regression with More Than One Feature

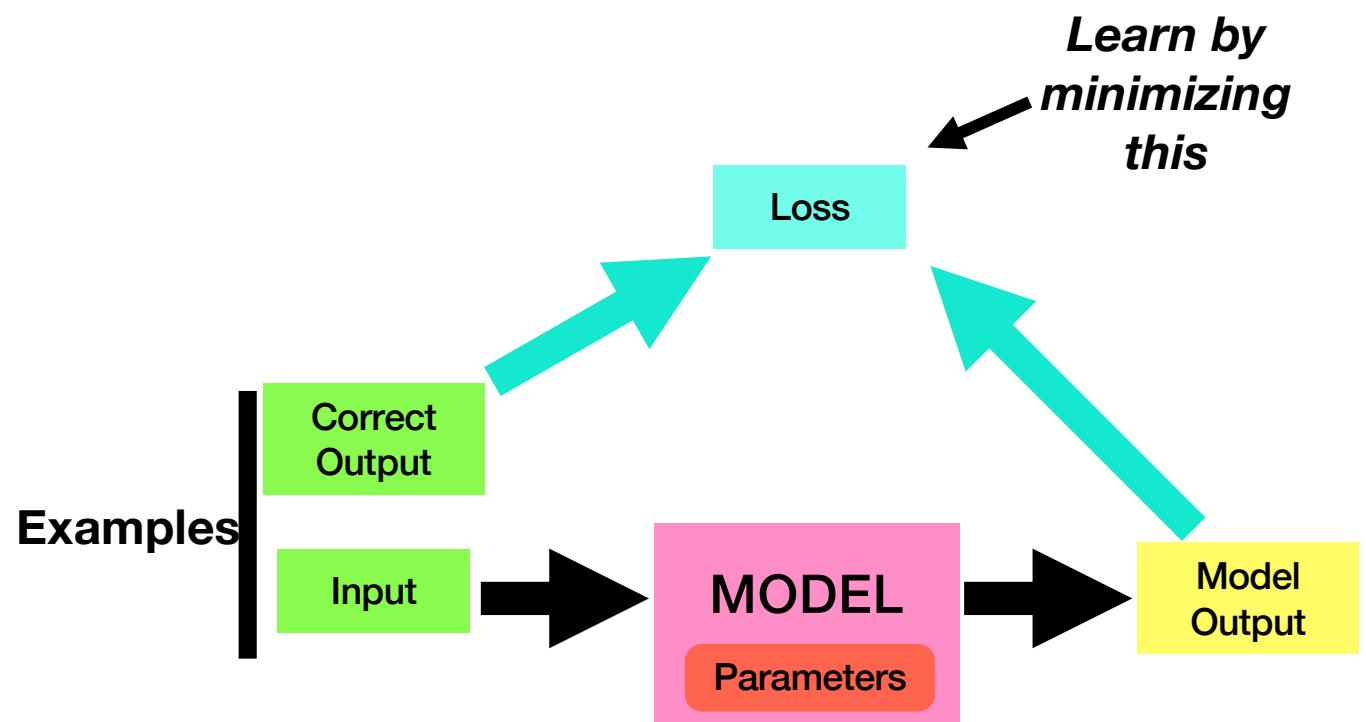
- We can still suppose that the *age of death* is a **linear function** of the *features*

$$age = \theta_0 + \theta_1 \times cig \rightarrow age = \theta_0 + \theta_1 \times cig + \theta_2 \times bmi + \theta_3 \times ismale$$

- We now have 2 more parameters (because we have 2 more features)
- For a total of 4 parameters
- But finding the parameters will be done in exactly the same way as in the case with one feature

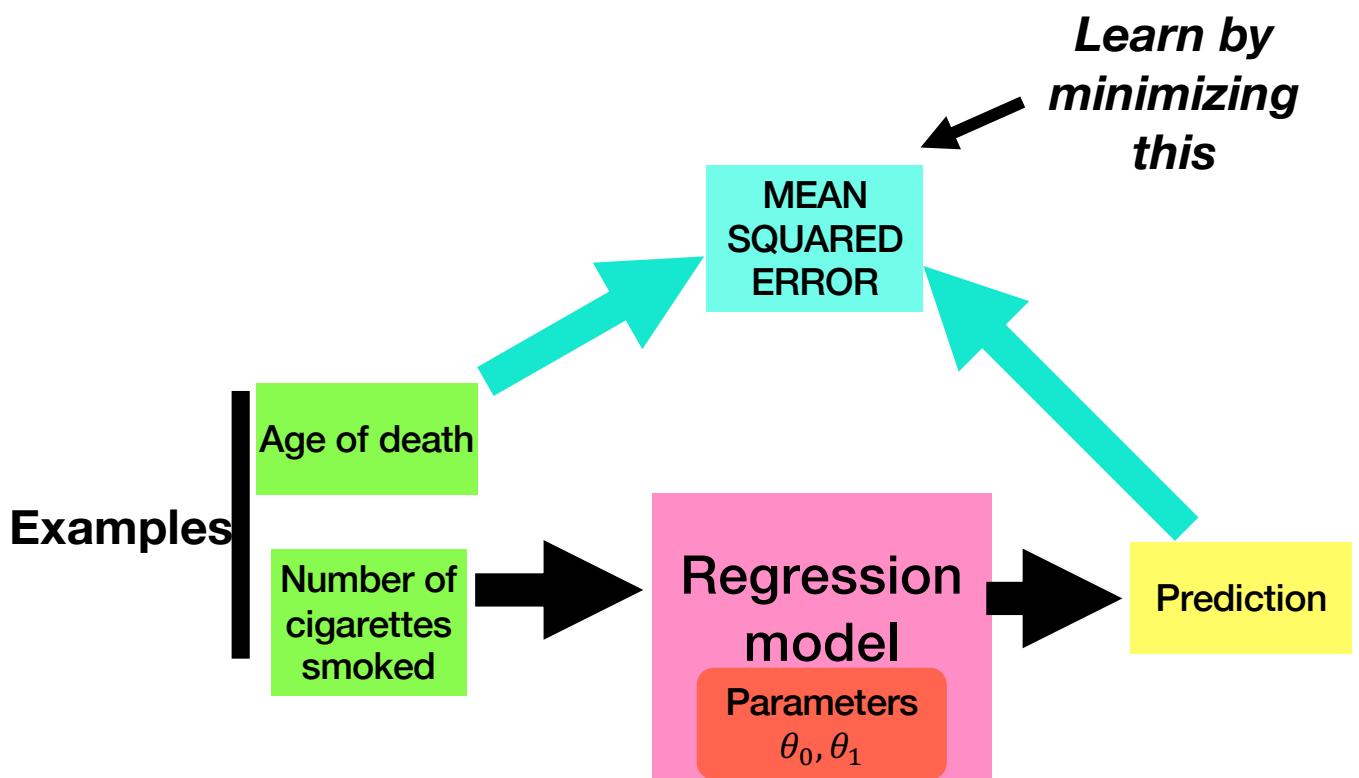
# Supervised Learning (1/3)

- In supervised learning, we usually have:
  - A **MODEL**: a “parameterized” function that takes input and produce output
  - A **Loss**: A function that computes how different the model output is from the correct output
  - **Examples** of input and correct output



# Supervised Learning (2/3)

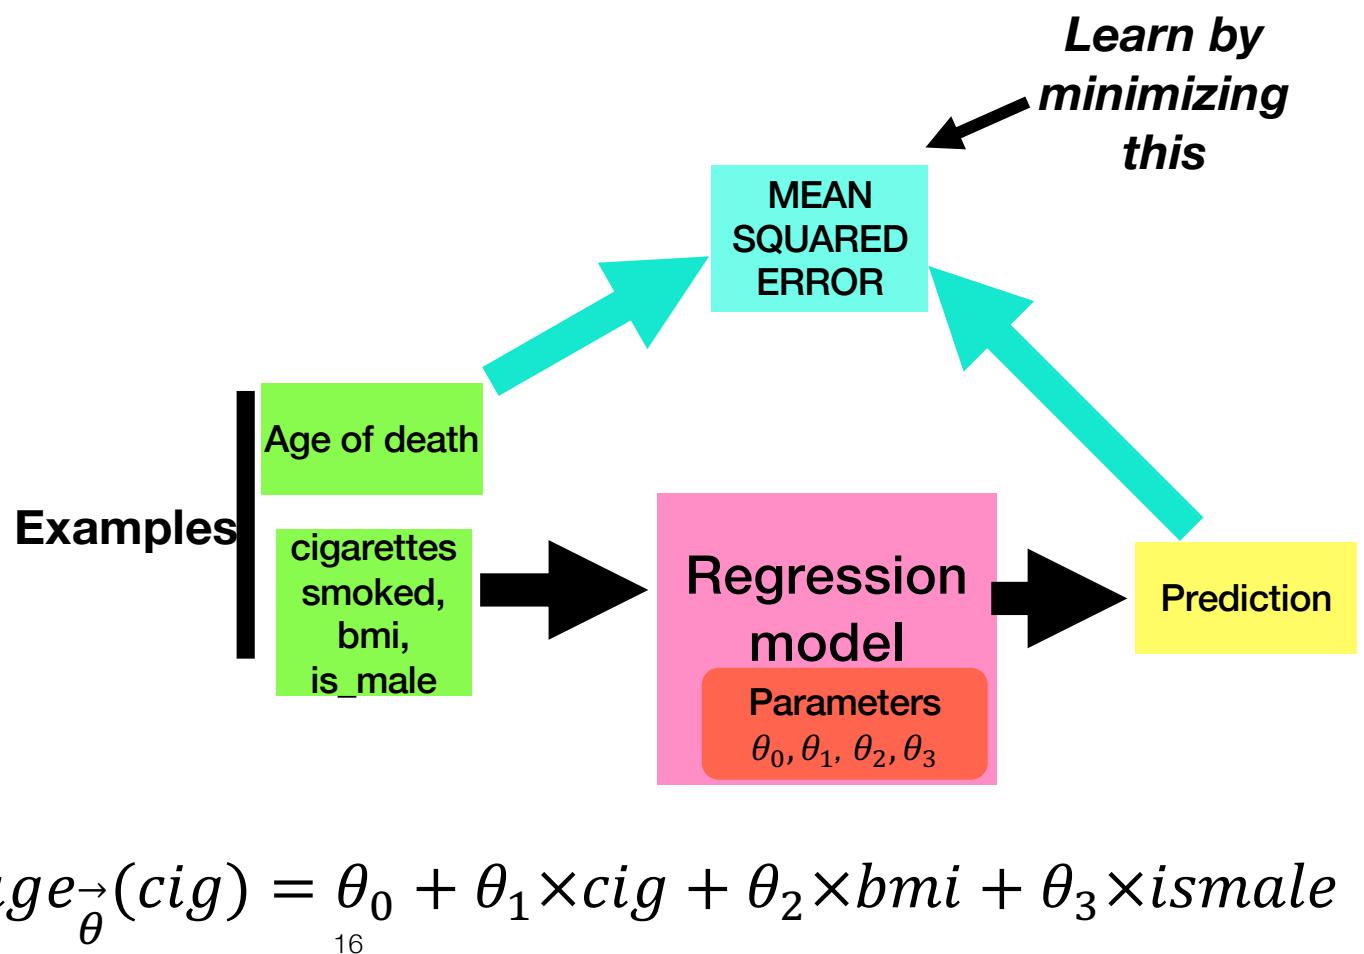
- In supervised learning, we usually have:
  - A **MODEL**: a “parameterized” function that takes input and produce output
  - A **Loss**: A function that computes how different the model output is from the correct output
  - **Examples** of input and correct output (cigarettes smoked, age of death)



$$age_{\theta_0, \theta_1}(cig) = \theta_0 + \theta_1 \times cig$$

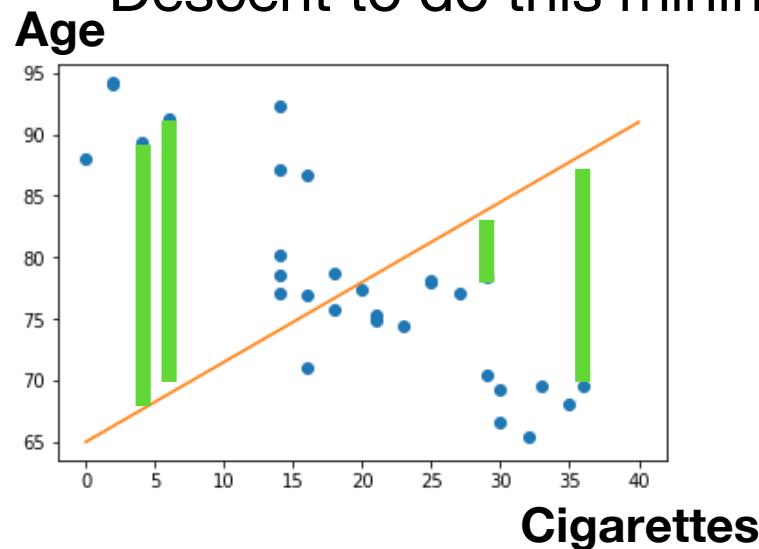
# Supervised Learning (3/3)

- In supervised learning, we usually have:
- A **MODEL**: a “parameterized” function that takes input and produce output
- A **Loss**: A function that computes how different the model output is from the correct output
- **Examples** of input and correct output (cigarettes smoked, bmi, is\_male, age of death)



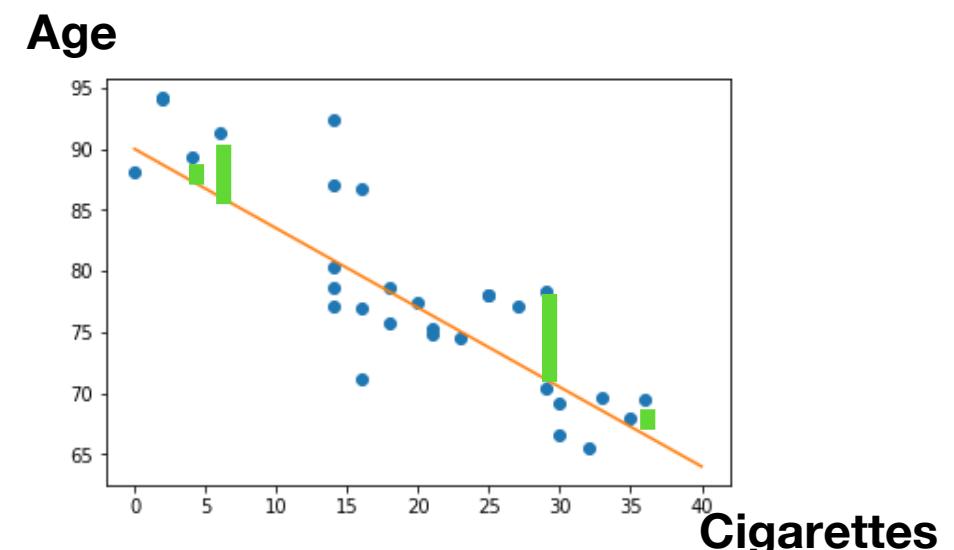
# Reminder: Loss with One Feature

- We saw that we could find the parameters of this linear relation by minimizing the Mean Squared Distance between the prediction of the model and the actual value; We saw we could use Gradient Descent to do this minimization



$$\theta_0 \approx 65 \quad \theta_1 \approx 0.7$$

*MeanSquaredError = 294.7*



$$\theta_0 \approx 90 \quad \theta_1 \approx -0.7$$

*MeanSquaredError = 18.3*

# Linear Regression with More Than One Feature

- The **loss** is defined as before: the average of the squared distance between model prediction and real example values

$$\text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (\underline{\text{model}(cig_i, bmi_i, ismale_i)} - \underline{\text{age}_i})^2$$

**Model prediction for example i**   **Real value for example i**

	daily cigarettes	bmi	is male	age of death
0	5.0	18.5	1.0	79.8
1	9.0	45.1	0.0	56.8
2	38.0	14.2	0.0	61.4
3	12.0	48.5	1.0	37.5
7	25.0	33.6	1.0	63.0
8	24.0	45.2	1.0	39.3
9	...	....	....	....

$$\text{age}_{\text{model}} = \theta_0 + \theta_1 \times \text{cig} + \theta_2 \times \text{bmi} + \theta_3 \times \text{ismale}$$

# Small Exercise: Compute the Loss

- The loss is defined as before: the average of the squared distance between model prediction and real example values

$$\text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (\underline{\text{model}(cig_i, bmi_i, ismale_i)} - \underline{\text{age}_i})^2$$

**Model prediction for example i**   **Real value for example i**

	daily cigarettes	bmi	is male	age of death
0	5.0	18	0.0	80
1	9.0	45	1.0	57
2	38.0	16	0.0	61

$$\text{age}_{\text{model}} = \theta_0 + \theta_1 \times \text{cig} + \theta_2 \times \text{bmi} + \theta_3 \times \text{ismale}$$
$$\theta_0 = 90 \quad \theta_1 = -1 \quad \theta_2 = -0.1 \quad \theta_3 = -6$$

Compute the model prediction for each of the examples.  
Then, compute the loss.

$$MeanSquaredError = \frac{1}{N} \cdot \sum_i (\underline{\text{model}(cig_i, bmi_i, ismale_i)} - \underline{\text{age}_i})^2$$

**Model prediction for example i Real value for example i**

	daily cigarettes	bmi	is male	age of death
0	5.0	18	0.0	80
1	9.0	45	1.0	57
2	38.0	16	0.0	61

$$\text{age}_{\text{model}} = \theta_0 + \theta_1 \times \text{cig} + \theta_2 \times \text{bmi} + \theta_3 \times \text{ismale}$$

$$\theta_0 = 90 \quad \theta_1 = -1 \quad \theta_2 = -0.1 \quad \theta_3 = -6$$

# What is the Gradient (1/2)?

- To do it a bit differently than last time: let us note the error on example  $i$  as:

$$error_i = model(cig_i, bmi_i, ismale_i) - age_i$$

Then our loss is equal to:

$$MeanSquaredError = \frac{1}{N} \cdot \sum_i (error_i)^2$$

(The mean squared distance is actually also often called the mean squared error)

Then our gradient is equal to:

$$\frac{\partial}{\partial \theta_k} MeanSquaredDistance = \frac{1}{N} \cdot \sum_i 2 \times error_i \times \frac{\partial}{\partial \theta_k} error_i$$

Using linearity and the fact that

$$\frac{d}{dx} [f(x)]^2 = 2 \times f(x) \times \frac{d}{dx} f(x)$$

# What is the Gradient (2/2)?

$$\frac{\partial}{\partial \theta_k} MeanSquaredError = \frac{1}{N} \cdot \sum_i 2 \times error_i \times \frac{\partial}{\partial \theta_k} error_i$$

$$error_i = model(cig_i, bmi_i, ismale_i) - age_i$$

$$age_{model} = \theta_0 + \theta_1 \times cig + \theta_2 \times bmi + \theta_3 \times ismale$$

$$\frac{\partial}{\partial \theta_0} error_i = 1$$

$$\frac{\partial}{\partial \theta_1} error_i = cig_i$$

$$\frac{\partial}{\partial \theta_2} error_i = bmi_i$$

$$\frac{\partial}{\partial \theta_3} error_i = ismale_i$$



(we know have 4 parameters, so gradient is a 4-dimensional vector)

$$\begin{aligned} & \frac{1}{N} \cdot \sum_i 2 \times error_i \times 1 \\ & \frac{1}{N} \cdot \sum_i 2 \times error_i \times cig_i \\ & \frac{1}{N} \cdot \sum_i 2 \times error_i \times bmi_i \\ & \frac{1}{N} \cdot \sum_i 2 \times error_i \times ismale_i \end{aligned}$$

# Small Exercise:

## Compute the Gradient

- Let us compute the gradient for this examples and  $\theta_k$ :

$$\text{error}_i = \text{model}(\text{cig}_i, \text{bmi}_i, \text{ismale}_i) - \text{age}_i$$

$$\begin{aligned} \text{age}_{\text{model}} &= \theta_0 + \theta_1 \times \text{cig} + \theta_2 \times \text{bmi} + \theta_3 \times \text{ismale} \\ \theta_0 &= 90 \quad \theta_1 = -1 \quad \theta_2 = -0.1 \quad \theta_3 = -6 \end{aligned}$$

	daily cigarettes	bmi	is male	age of death
0	5.0	18	0.0	80
1	9.0	45	1.0	57
2	38.0	16	0.0	61

Note you already computed the errors when you computed the loss in previous exercise

$$\text{gradient} = \left[ \begin{array}{l} \frac{1}{N} \cdot \sum_i 2 \times \text{error}_i \times 1 \\ \frac{1}{N} \cdot \sum_i 2 \times \text{error}_i \times \text{cig}_i \\ \frac{1}{N} \cdot \sum_i 2 \times \text{error}_i \times \text{bmi}_i \\ \frac{1}{N} \cdot \sum_i 2 \times \text{error}_i \times \text{ismale}_i \end{array} \right]$$

- Let us compute the gradient for this examples and  $\theta_k$ :

$$\theta_0 = 90 \quad \theta_1 = -1 \quad \theta_2 = -0.1 \quad \theta_3 = -6$$

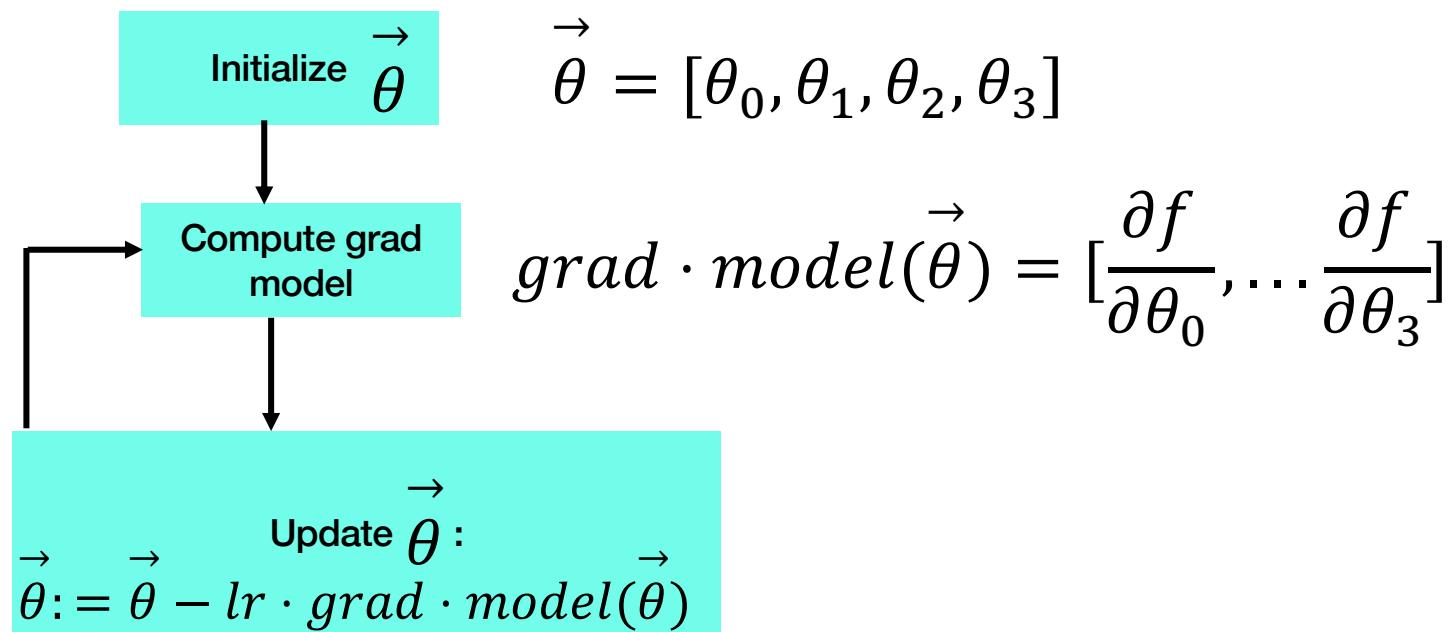
	daily cigarettes	bmi	is male	age of death
0	5.0	18	0.0	80
1	9.0	45	1.0	57
2	38.0	16	0.0	61

Note you already computed the errors when you computed the loss in previous exercise

$$gradient = \begin{array}{l} \frac{1}{N} \cdot \sum_i 2 \times error_i \times 1 \\ \frac{1}{N} \cdot \sum_i 2 \times error_i \times cig_i \\ \frac{1}{N} \cdot \sum_i 2 \times error_i \times bmi_i \\ \frac{1}{N} \cdot \sum_i 2 \times error_i \times ismale_i \end{array}$$

# Gradient Descent

- Now that we know how to compute the gradient, we can find the optimal parameters with gradient descent:



# Feature Scaling (1/2)

- We can see that our features have very different **ranges and means**: “daily cigarettes” is from 0 to 50, “bmi” is from 15 to 50, “is\_male” is from 0 to 1
- Usually, features having different ranges make learning/gradient descent more difficult

	daily cigarettes	bmi	is male	age of death
0	5.0	18.5	1.0	79.8
1	9.0	45.1	0.0	56.8
2	38.0	14.2	0.0	61.4
3	12.0	48.5	1.0	37.5
7	25.0	33.6	1.0	63.0
8	24.0	45.2	1.0	39.3
9	...	....	....	....

**The solution is to scale all the features to the same range**

$$new = \frac{old - mean(old)}{max(old) - min(old)}$$

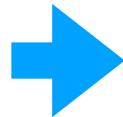
**This way, all features have ranges between -1 and 1, and mean equal to zero**

# Feature Scaling (2/2)

- After the features have been scaled, we apply Gradient Descent as usual:

$$new = \frac{old - mean(old)}{max(old) - min(old)}$$

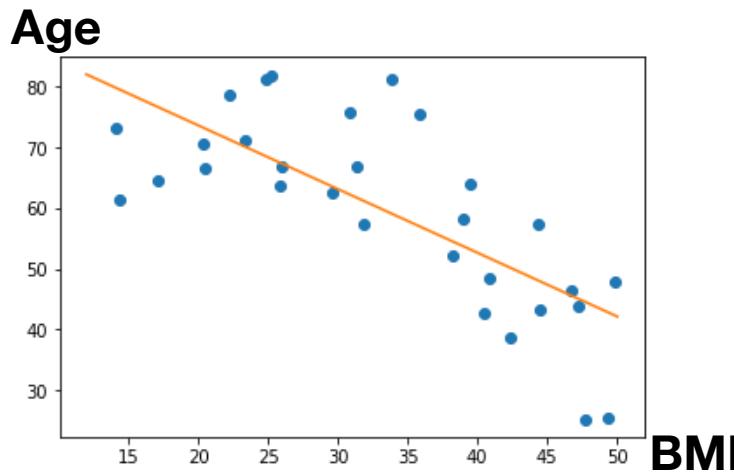
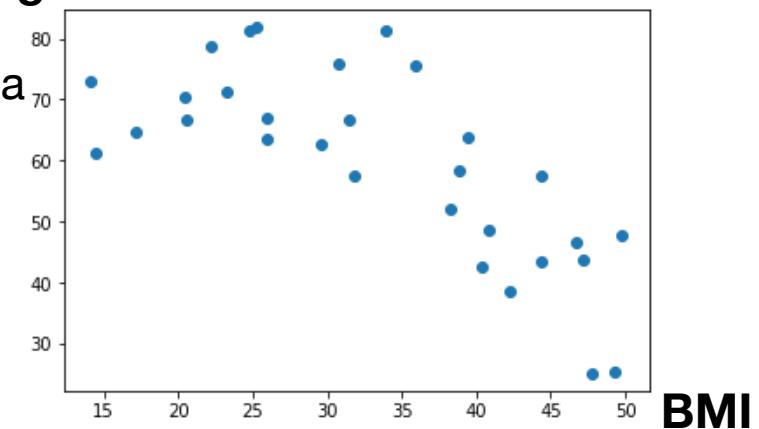
	daily cigarettes	bmi	is male	age of death
0	24.0	44.4	1.0	43.4
1	37.0	47.7	1.0	25.2
2	11.0	14.1	0.0	73.1
3	33.0	49.3	1.0	25.5
4	27.0	20.4	0.0	70.5
5	27.0	38.2	1.0	52.1
6	6.0	22.2	1.0	78.6
7	31.0	17.1	0.0	64.6
8	28.0	23.3	0.0	71.3
9	15.0	31.4	1.0	66.8



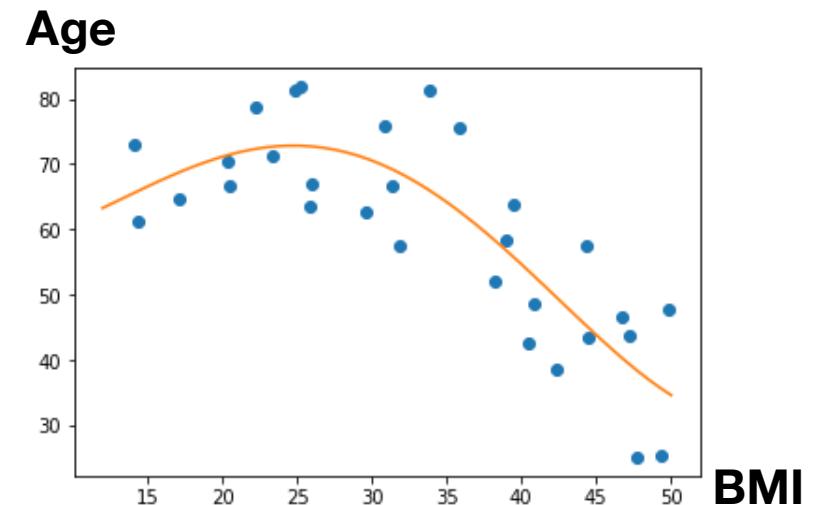
	daily cigarettes	bmi	is male	age of death
0	0.064957	0.313072	0.433333	43.4
1	0.398291	0.405509	0.433333	25.2
2	-0.268376	-0.535668	-0.566667	73.1
3	0.295726	0.450327	0.433333	25.5
4	0.141880	-0.359197	-0.566667	70.5
5	0.141880	0.139402	0.433333	52.1
6	-0.396581	-0.308777	0.433333	78.6
7	0.244444	-0.451634	-0.566667	64.6
8	0.167521	-0.277965	-0.566667	71.3
9	-0.165812	-0.051074	0.433333	66.8

# Learning More Complicated Functions

- So far, we have only tried to learn linear functions of the data
- We might want to learn more complicated functions
- For example, having a high BMI reduce life expectancy
- But having a very low BMI also reduces life expectancy



We need a more complex model!



# Learning More Complex Functions: “Expanding the Feature Space”

- One neat trick to learn more complex functions: Create additional features from existing features; Then apply linear regression
- Example: From the feature bmi, we add  $bmi^2$ ,  $bmi^3$  and  $bmi^4$ :

	bmi	age of death
0	44.4	43.4
1	47.7	25.2
2	14.1	73.1
3	49.3	25.5
4	20.4	70.5
5	38.2	52.1
6	...	...



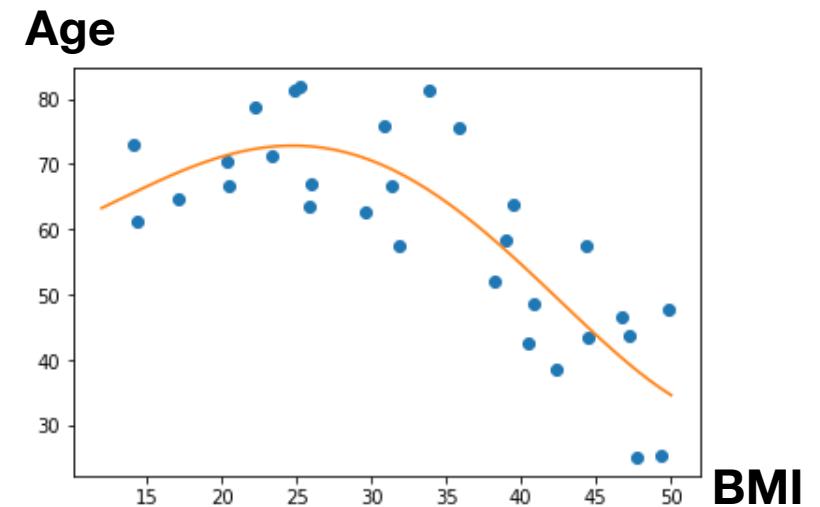
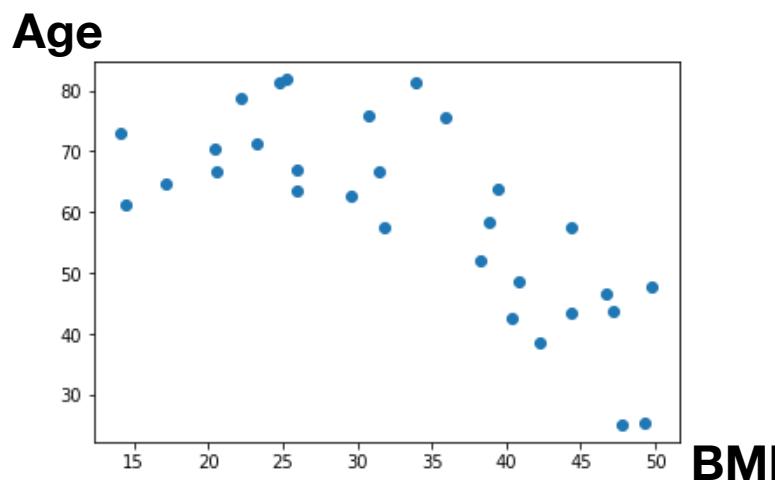
$$age_{model} = \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 + \theta_3 \times bmi^3 + \theta_4 \times bmi^4$$

	bmi	bmi <sup>2</sup>	bmi <sup>3</sup>	bmi <sup>4</sup>	age of death
0	44.4	1971.36	87528.384	3.88E+06	43.4
1	47.7	2275.29	108531.333	5.17E+06	25.2
2	14.1	198.81	2803.221	3.95E+04	73.1
3	49.3	2430.49	119823.157	5.90E+06	25.5
4	20.4	416.16	8489.664	1.73E+05	70.5
5	38.2	1459.24	55742.968	2.12E+06	52.1
6	...	...	...	...	...

# Learning More Complicated Functions

- This way, we can learn much more complex functions
- After finding the optimal parameters by gradient descent on the Mean Squared Error:

$$age_{model} = \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 + \theta_3 \times bmi^3 + \theta_4 \times bmi^4$$



# Learning Functions That are Too Complicated (1/3)

- Now, we might think:

Cool. I am going to add as many variations of the features as I can. Then my loss (Mean Squared Error) will drop to zero!



	bmi	bmi^2	bmi^3	bmi^4	bmi^5	sin(bmi)	log(bmi)	...	age of death
0	44.4	1971.36	87528.384	3.886260E+06	1.725500E+08	0.405662	3.793239...		43.4
1	47.7	2275.29	108531.333	5.176945E+06	2.469403E+08	-0.544766	3.864931...		25.2
2	14.1	198.81	2803.221	3.952542E+04	5.573084E+05	0.999309	2.646175...		73.1
3	49.3	2430.49	119823.157	5.907282E+06	2.912290E+08	-0.822324	3.897924...		25.5
4	20.4	416.16	8489.664	1.731891E+05	3.533059E+06	0.999793	3.015535...		70.5
5	38.2	1459.24	55742.968	2.129381E+06	8.134237E+07	0.480205	3.642836...		52.1
6	22.2	492.84	10941.048	2.428913E+05	5.392186E+06	-0.207336	3.100092...		78.6
7	17.1	292.41	5000.211	8.550361E+04	1.462112E+06	-0.984065	2.839078...		64.6
8	...	...	...	...	...	...	...	...	...

$$\begin{aligned}age_{model} = & \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 + \theta_3 \times bmi^3 + \theta_4 \times bmi^4 \\& + \theta_5 \times bmi^5 + \theta_6 \times sin(bmi) + \theta_7 \times log(bmi) + \dots\end{aligned}$$

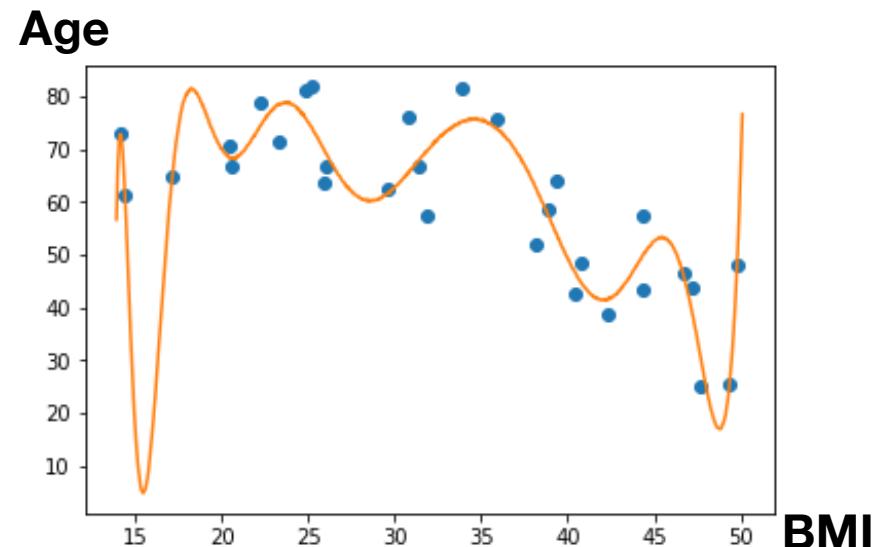
# Learning Functions That are Too Complicated (2/3)

- The problem: this is what you are going to get:



Cool. My loss (Mean Squared Error) is lower than before. My model is better now!

Or is it really better???



$MeanSquaredError = 1.3$

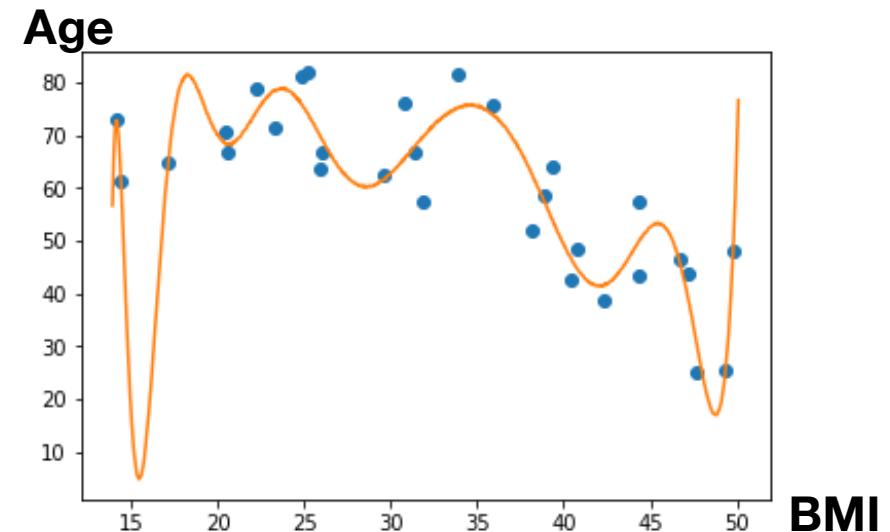
# Learning Functions That are Too Complicated (3/3)

- The problem: this is what you are going to get:



Cool. My loss (Mean Squared Error) is lower than before. My model is better now!

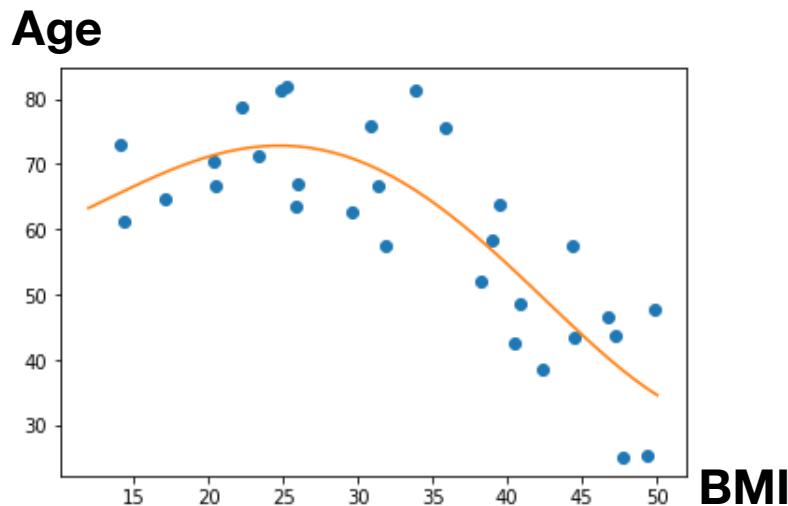
Or is it really better???



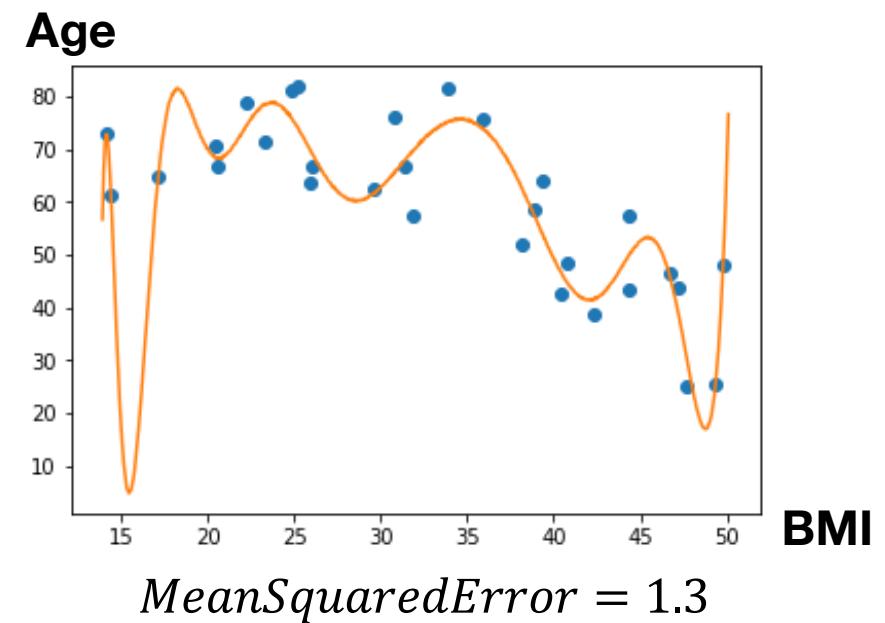
The model predict that somebody with BMI 48 will die at 18  $MeanSquaredError = 1.3$   
But somebody with BMI 51 is predicted to die at 80

# Overfitting (1/2)

- Minimizing the loss does not always give the best model....
- This phenomenon is called overfitting



*MeanSquaredError = 11.2*



*MeanSquaredError = 1.3*

# Overfitting (2/2)

- **Overfitting** is to Machine Learning what **Rote Learning** is to Human Learning
- Instead of understanding the data, the model just memorized all of the examples
- If we ask it to predict the age of death of an example it has seen, it will give very good prediction
- But it will give very bad prediction as soon as we ask him to make a prediction for someone that was not in the example data.
- Very similar to a student that memorize without understanding the answer to a set of exercises in a class. He will do very well in the exam if the exam contains the exercises he studied, but very bad if the exercises are a bit different.

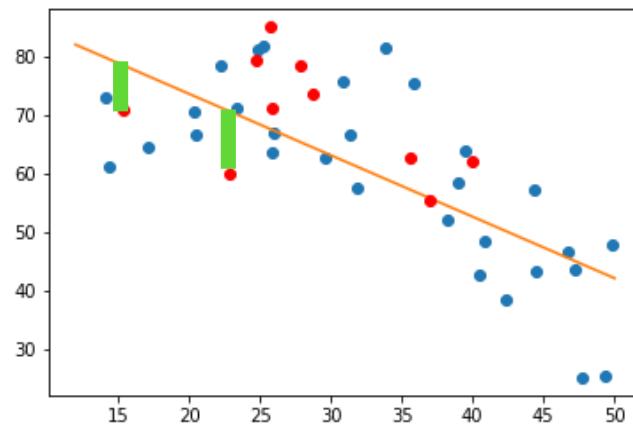
# How to Detect Overfitting (1/3)?

- Pretty much like we do with humans: We evaluate them with different exercises than the ones they were trained for.
- In practice, it means we separate our example data in 2:
  - Training Data
  - Test Data
- We use the training data to do the Learning (we train the model)
- We use the test data to evaluate the quality of the learning (we test the model)

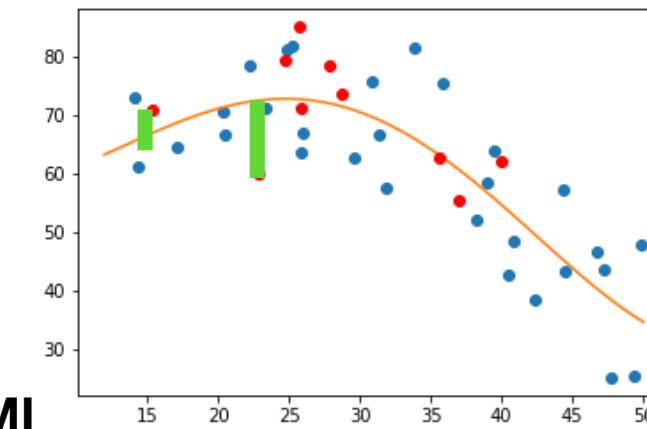
# How to Detect Overfitting (2/3)?

- In blue: training examples
- In red: test examples

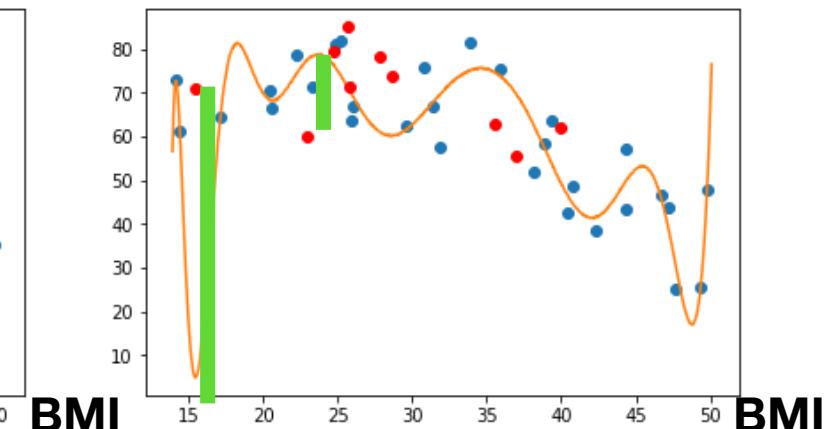
Age



Age



Age



MeanSquaredError = 18.6

MeanSquaredError = 11.2

MeanSquaredError = 1.3

TestMeanSquaredError = 17.3

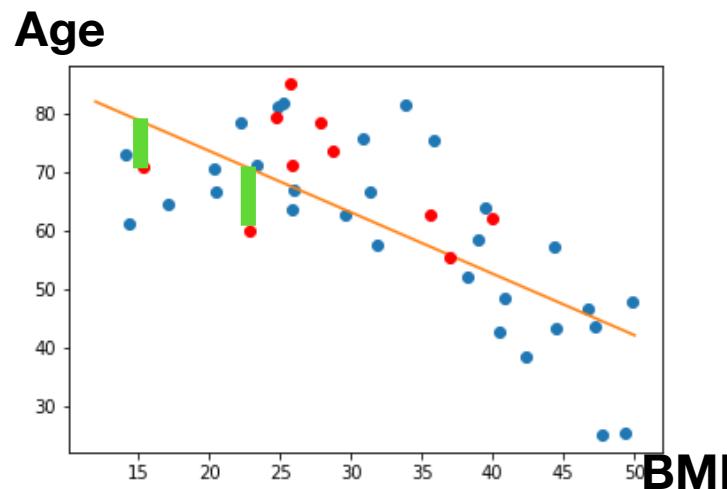
TestMeanSquaredError = 11.8

TestMeanSquaredError = 24.5

# How to Detect Overfitting (3/3)?

- If we see a model with very low training loss and high test loss: it is overfitting!

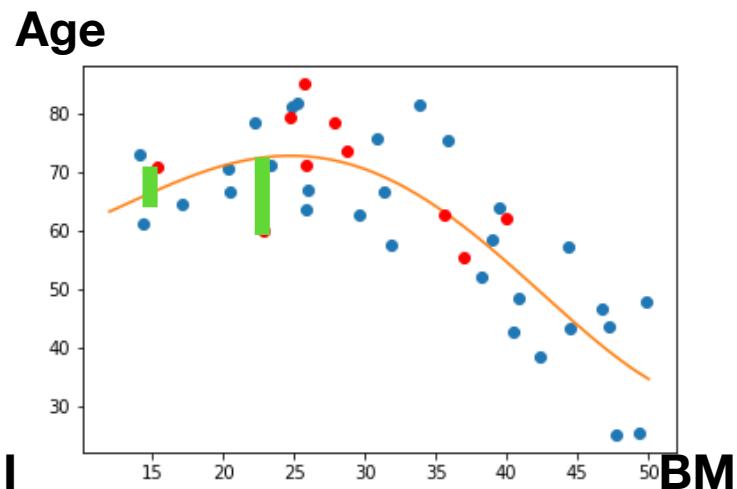
**Linear Model**



MeanSquaredError = 18.6

TestMeanSquaredError = 17.3

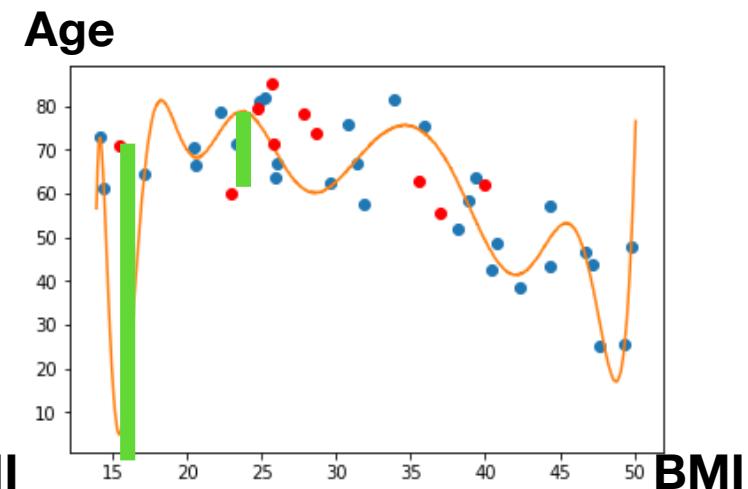
**More Complex Model**



MeanSquaredError = 11.2

TestMeanSquaredError = 11.8

**“Too Complex” Model**



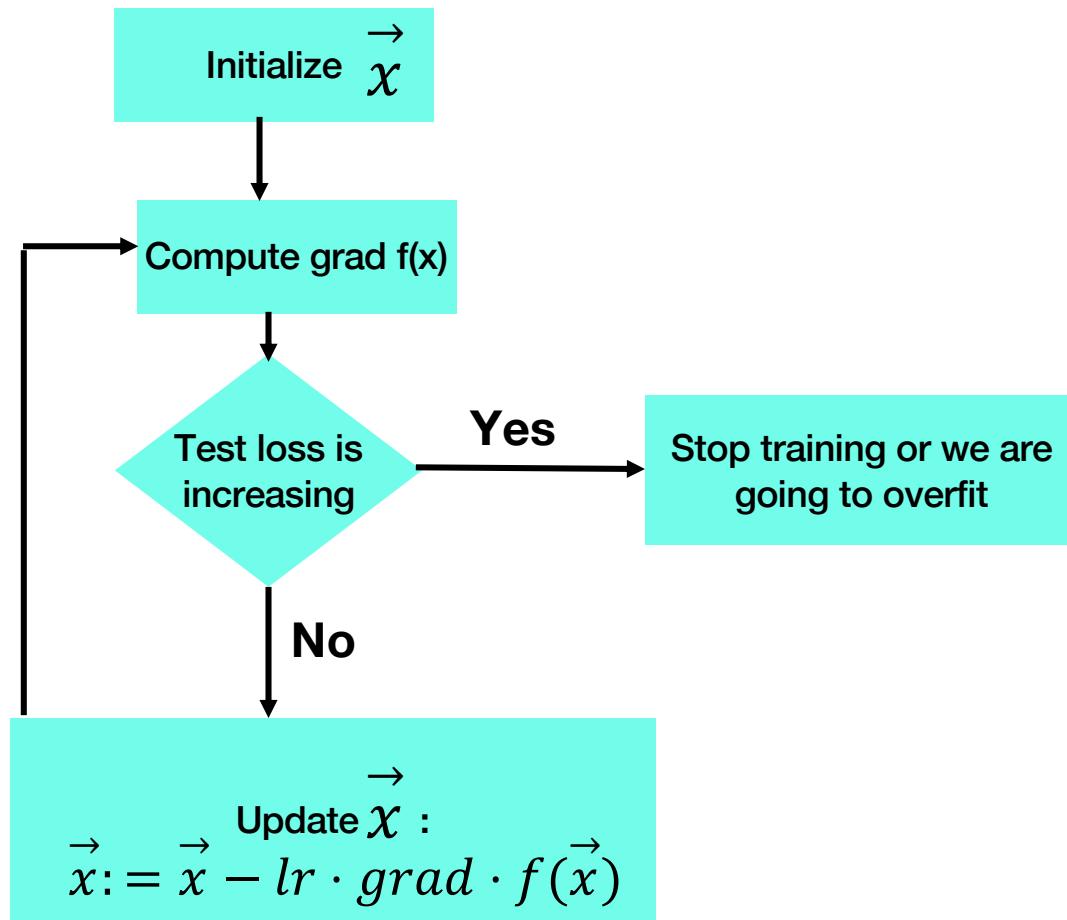
MeanSquaredError = 1.3

TestMeanSquaredError = 24.5

# How to Prevent Overfitting: Early Stopping (1/2)

- A very **simple** and very **efficient** method for preventing overfitting is called “***early stopping***”
- During the gradient descent, we check the test loss at each iteration. When the test loss starts to increase (or stay unchanged) for a few iteration, we stop the training

# How to Prevent Overfitting: Early Stopping (2/2)



# How to Prevent Overfitting: Capacity (1/5)

- Another method is to reduce the **capacity** of the model
- Roughly speaking, the capacity of a model is its ability to adapt to a large number of examples
- The capacity of a model will increase with the number of parameters

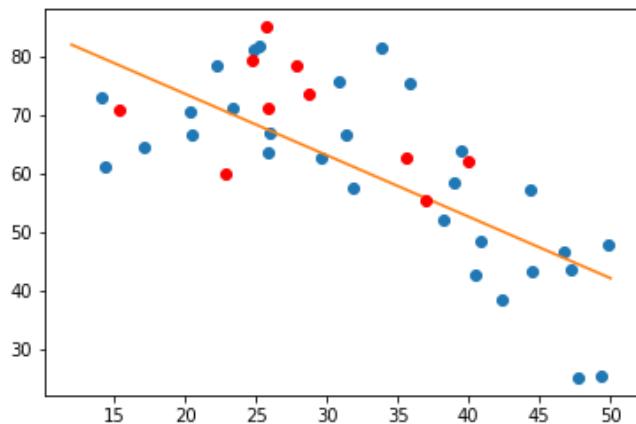
# How to Prevent Overfitting: Capacity (2/5)

- Models with high capacity can learn more complicated relations
- But they overfit more easily
- A model with very high capacity has the ability to memorize all the training examples (the rote learning problem)
- If we reduce the capacity of a model, we can make it less prone to overfitting

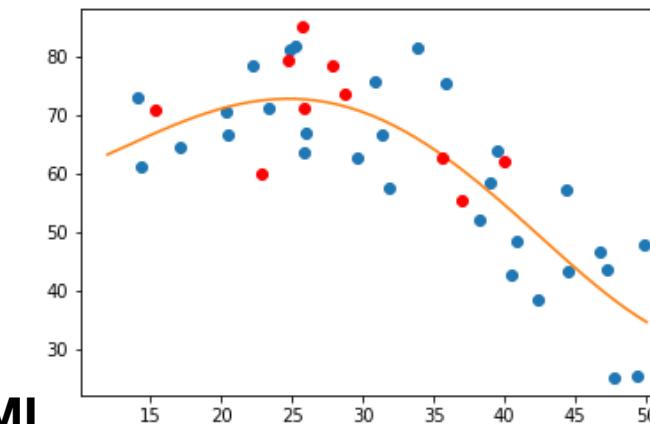
# How to Prevent Overfitting: Capacity (3/5)

Increasing Capacity

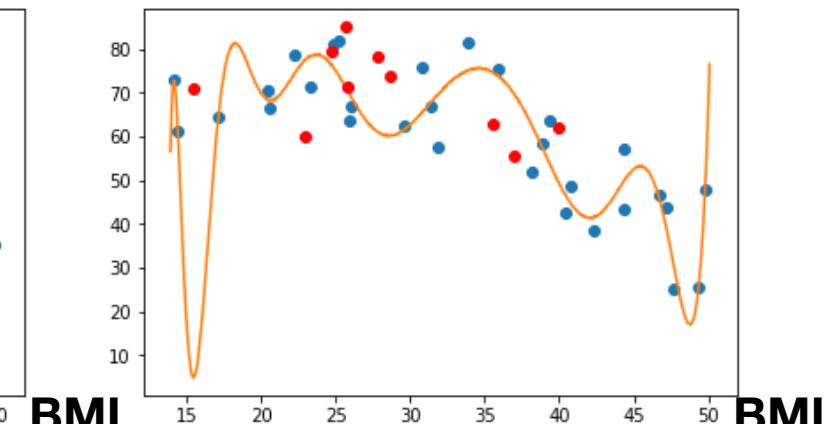
Age



Age



Age



$$age = \theta_0 + \theta_1 \times bmi$$

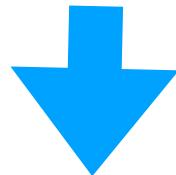
$$age_{model} = \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 \\ + \theta_3 \times bmi^3 + \theta_4 \times bmi^4$$

$$age_{model} = \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 \\ + \theta_3 \times bmi^3 + \theta_4 \times bmi^4 + \theta_5 \times bmi^5 \\ + \theta_6 \times \sin(bmi) + \theta_7 \times \log(bmi) + \dots$$

# How to Prevent Overfitting: Capacity (4/5)

- The simplest way to reduce capacity is to remove some parameters in the model:

$$\begin{aligned}age_{model} = & \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 + \theta_3 \times bmi^3 + \theta_4 \times bmi^4 \\& + \theta_5 \times bmi^5 + \theta_6 \times \sin(bmi) + \theta_7 \times \log(bmi) + \dots\end{aligned}$$



$$age_{model} = \theta_0 + \theta_1 \times bmi + \theta_2 \times bmi^2 + \theta_3 \times bmi^4 + \theta_4 \times bmi^5$$

# How to Prevent Overfitting: Capacity (5/5)

- Another way to reduce the capacity is “force” the model to use small values for the parameters
  - By default, the parameters  $\theta_k$  can take any value: -100 000, 0.1, 1 000 000
  - If we forbid the model to use very high values for  $\theta_k$ , we reduce its capacity: it cannot adapt to data as well as before
- Most practical way to forbid high values: “L2 Regularization”

# L2 Regularization

- We note  $\vec{|\theta|}^2$  the sum of the square of all parameters  $\theta_k$  (this is called the “L2 Norm”)
- Then we add this quantity to the loss we want to minimize:

$$Loss = MeanSquaredError = \frac{1}{N} \cdot \sum_i (model(cig_i, bmi_i, ismale_i) - age_i)^2$$



$$Loss = \frac{1}{N} \cdot \sum_i (model(cig_i, bmi_i, ismale_i) - age_i)^2 + \lambda \vec{|\theta|}^2$$

Then we apply Gradient Descent to this new loss

# Next

- Next time, we will consider Classification Problems:
  - Predicting if some symptoms are the sign of a disease or not...
  - Predicting if an image represents a cat or a dog
  - Predicting if a text is in French, German or Japanese

# Report

- Submit a report “summarizing what overfitting is and ways to detect and prevent overfitting” **in pdf** via PandA
  - Submission due: **next lecture**
  - Name the pdf file as **student id\_name**.

# Google Colab Notebook

<https://shorturl.at/Iy52R>