

Neural Networks with Fully Connected Layers in Practice

Fundamentals of Artificial Intelligence

Instructor: Chenhui Chu

Email: chu@i.kyoto-u.ac.jp

Teaching Assistant: Youyuan Lin

E-mail: youyuan@nlp.ist.i.kyoto-u.ac.jp

Schedule

- 1. Overview of AI and this Course (4/14)
- 2. Introduction to Python (4/21)
- 3, 4. Mathematics Concepts I, II (4/28, 5/12)
- 5, 6. Regression I, II (5/19, 5/26)
- 7. Classification (6/2)
- 8. Introduction to Neural Networks (6/9)
- 9. Neural Networks Architecture and Backpropagation (6/16)
- **10. Fully Connected Layers (6/23)**
- 11, 12, 13. Computer Vision I, II, III (6/30, 7/7, 7/14)
- 14. Natural Language Processing (7/17)

Overview of This Course

11, 12, 13. Computer Vision
I, II, III

14. Natural Language
Processing

Deep Learning Applications

8. Neural network
Introduction

9. Architecture and
Backpropagation

10. Feedforward
Neural Networks

Deep Learning

5. Regression I

6. Regression II

7. Classification

Basic Supervised Machine Learning

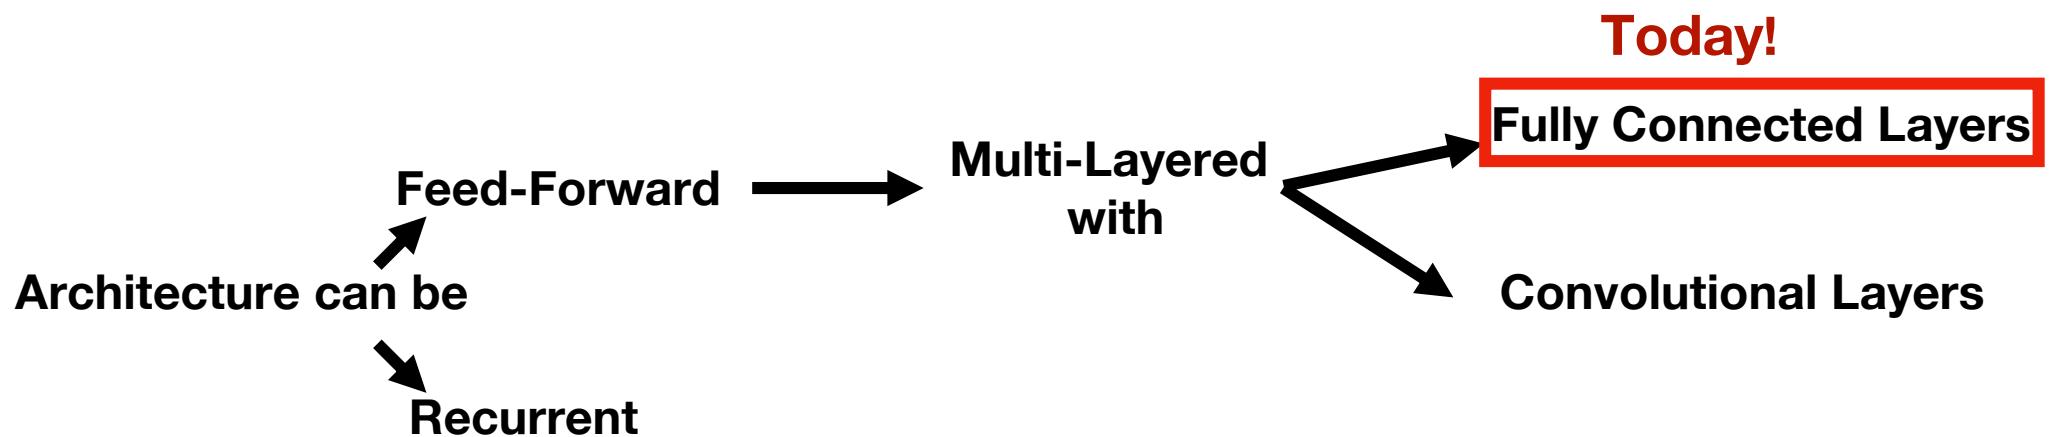
2. Python

3, 4. Mathematics Concepts I, II

Fundamental of Machine Learning

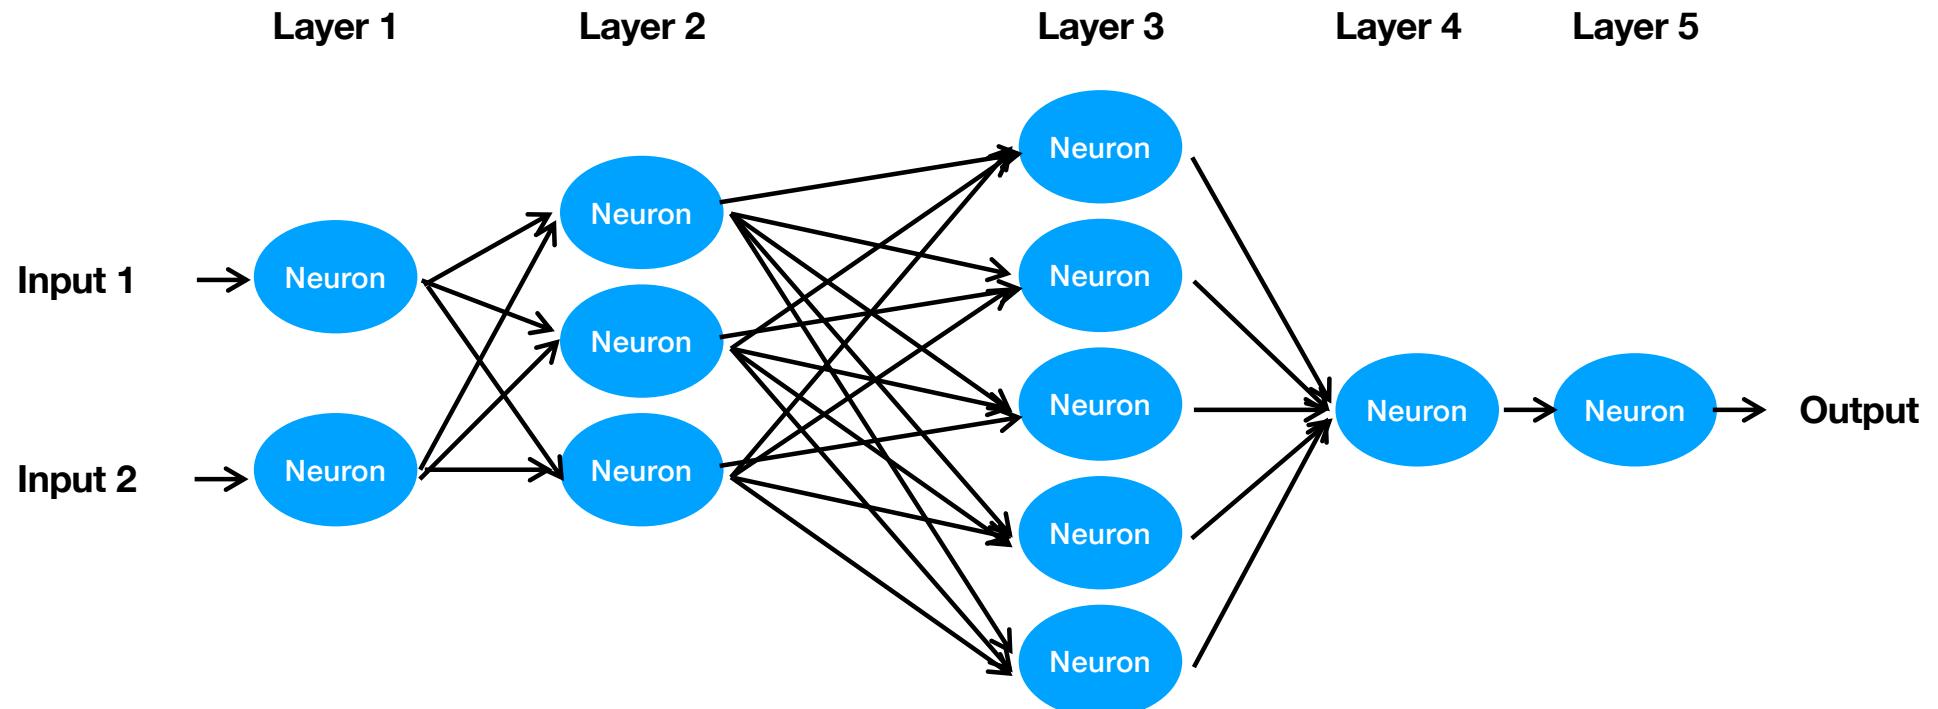
Neural Network Architectures

- We are still here:



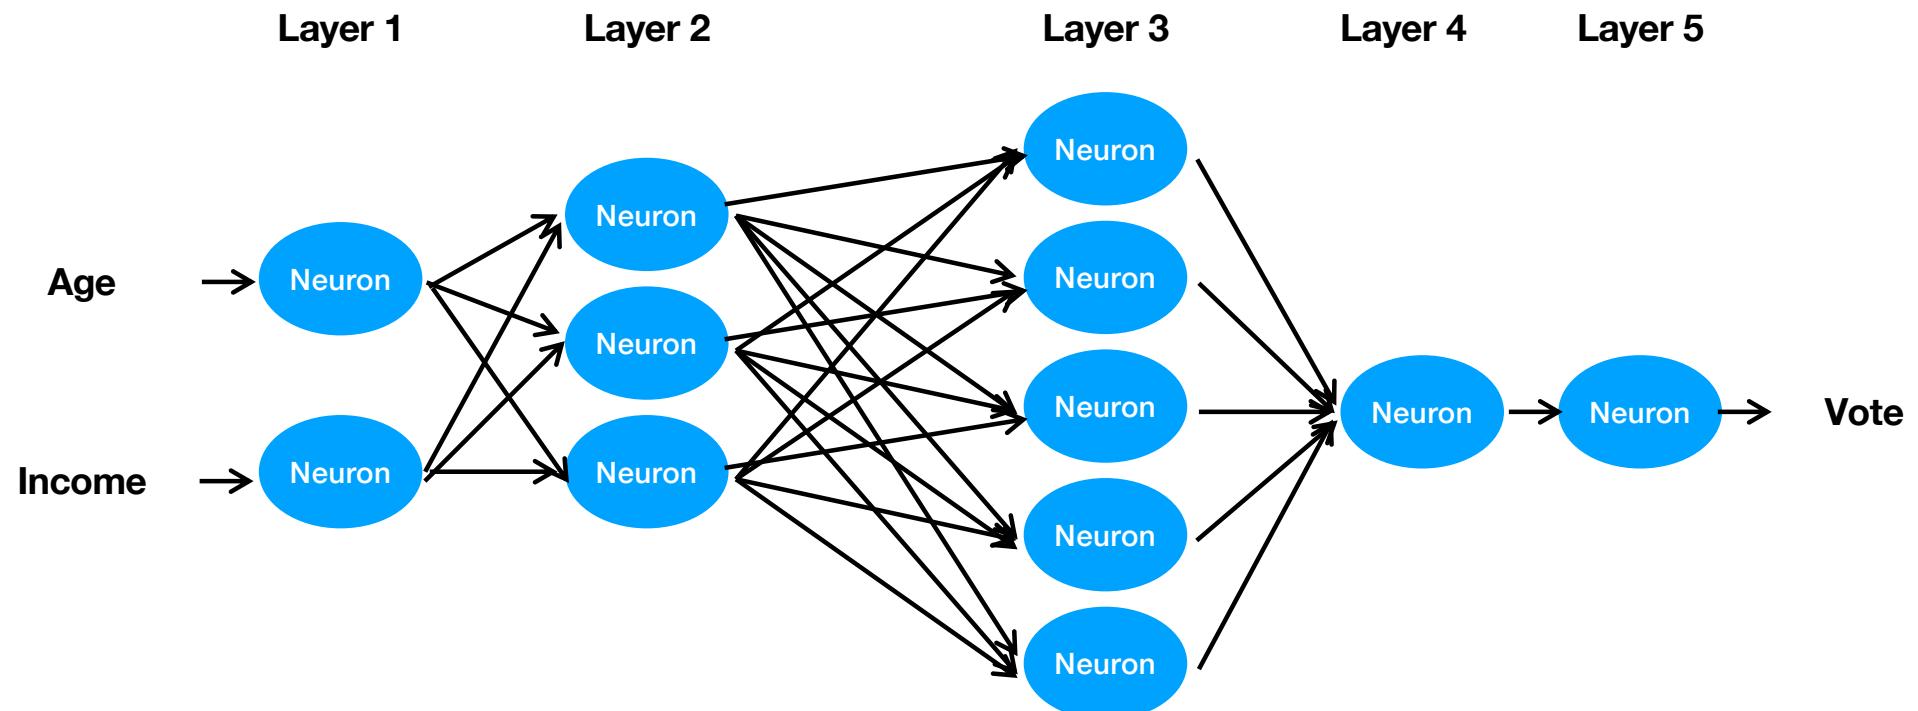
Feed-Forward Networks with Fully Connected Layers (1/2)

- Therefore, we are going to consider this type of Neural Network:



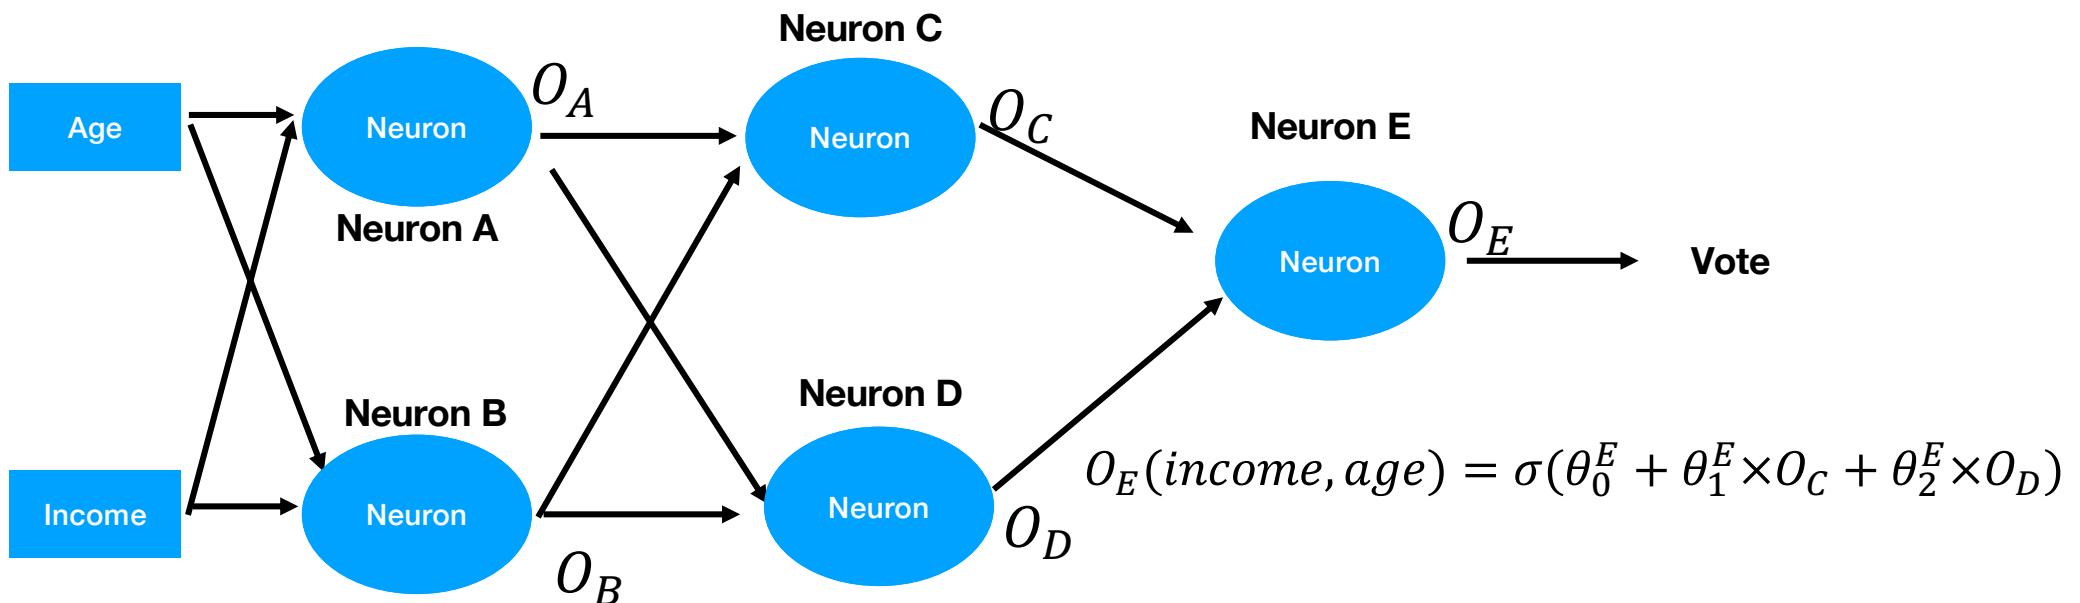
Feed-Forward Networks with Fully Connected Layers (2/2)

- Therefore, we are going to consider this type of Neural Network:



Keeping in Mind What This Type of Graph Mean (1/3)

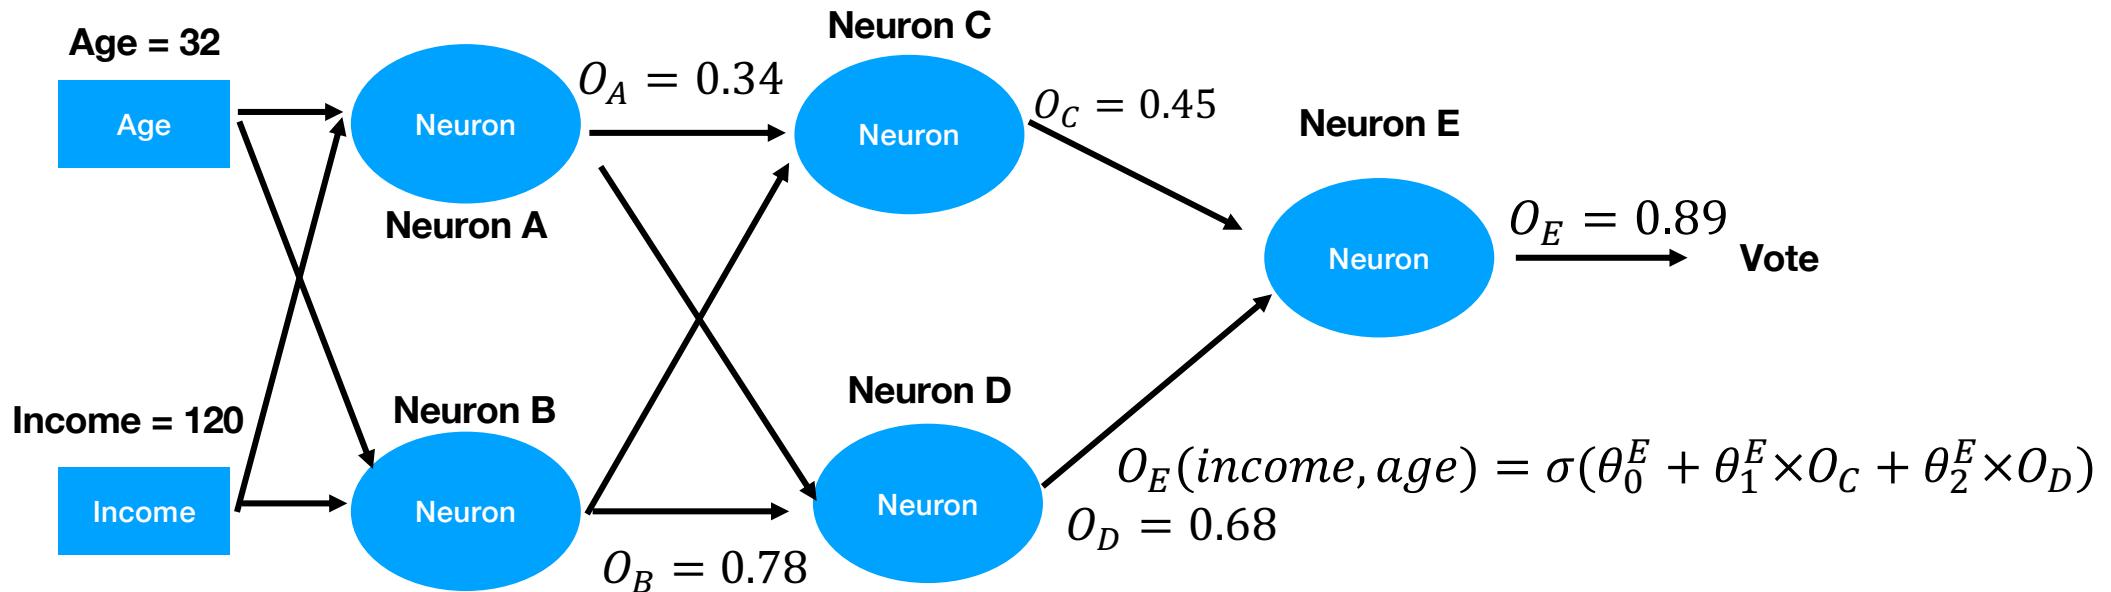
$$O_A(\text{income}, \text{age}) = \sigma(\theta_0^A + \theta_1^A \times \text{income} + \theta_2^A \times \text{age}) \quad O_C(\text{income}, \text{age}) = \sigma(\theta_0^C + \theta_1^C \times O_A + \theta_2^C \times O_B)$$



$$O_B(\text{income}, \text{age}) = \sigma(\theta_0^B + \theta_1^B \times \text{income} + \theta_2^B \times \text{age}) \quad O_D(\text{income}, \text{age}) = \sigma(\theta_0^D + \theta_1^D \times O_A + \theta_2^D \times O_B)$$

Keeping in Mind What This Type of Graph Mean (2/3)

$$O_A(\text{income}, \text{age}) = \sigma(\theta_0^A + \theta_1^A \times \text{income} + \theta_2^A \times \text{age}) \quad O_C(\text{income}, \text{age}) = \sigma(\theta_0^C + \theta_1^C \times O_A + \theta_2^C \times O_B)$$

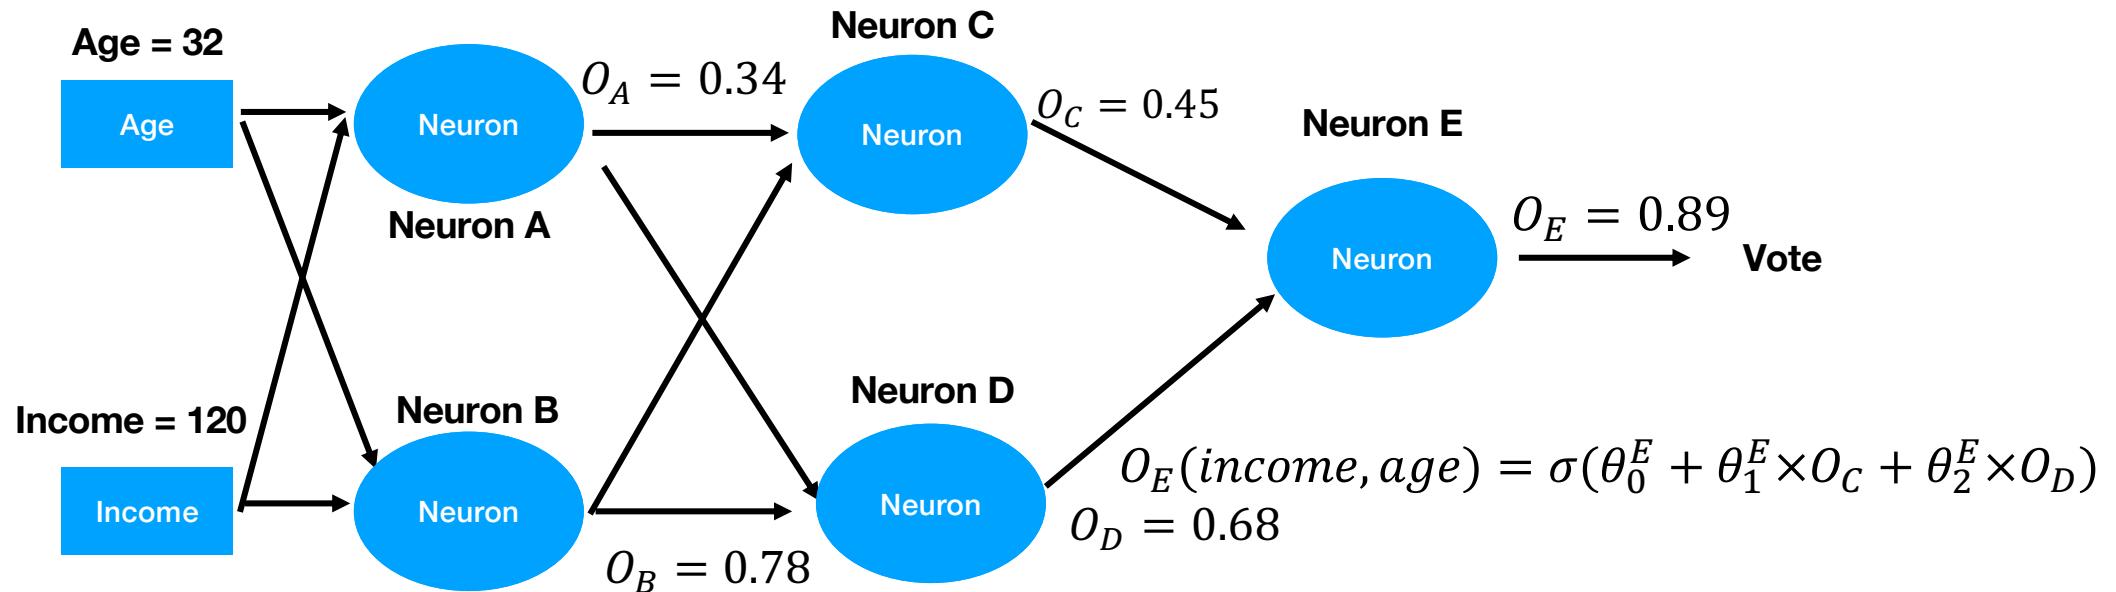


$$O_B(\text{income}, \text{age}) = \sigma(\theta_0^B + \theta_1^B \times \text{income} + \theta_2^B \times \text{age}) \quad O_D(\text{income}, \text{age}) = \sigma(\theta_0^D + \theta_1^D \times O_A + \theta_2^D \times O_B)$$

Keeping in Mind What This Type of Graph Mean (3/3)

→ Each Neural Network architecture defines a function of the input with parameters θ

$$O_A(\text{income}, \text{age}) = \sigma(\theta_0^A + \theta_1^A \times \text{income} + \theta_2^A \times \text{age}) \quad O_C(\text{income}, \text{age}) = \sigma(\theta_0^C + \theta_1^C \times O_A + \theta_2^C \times O_B)$$

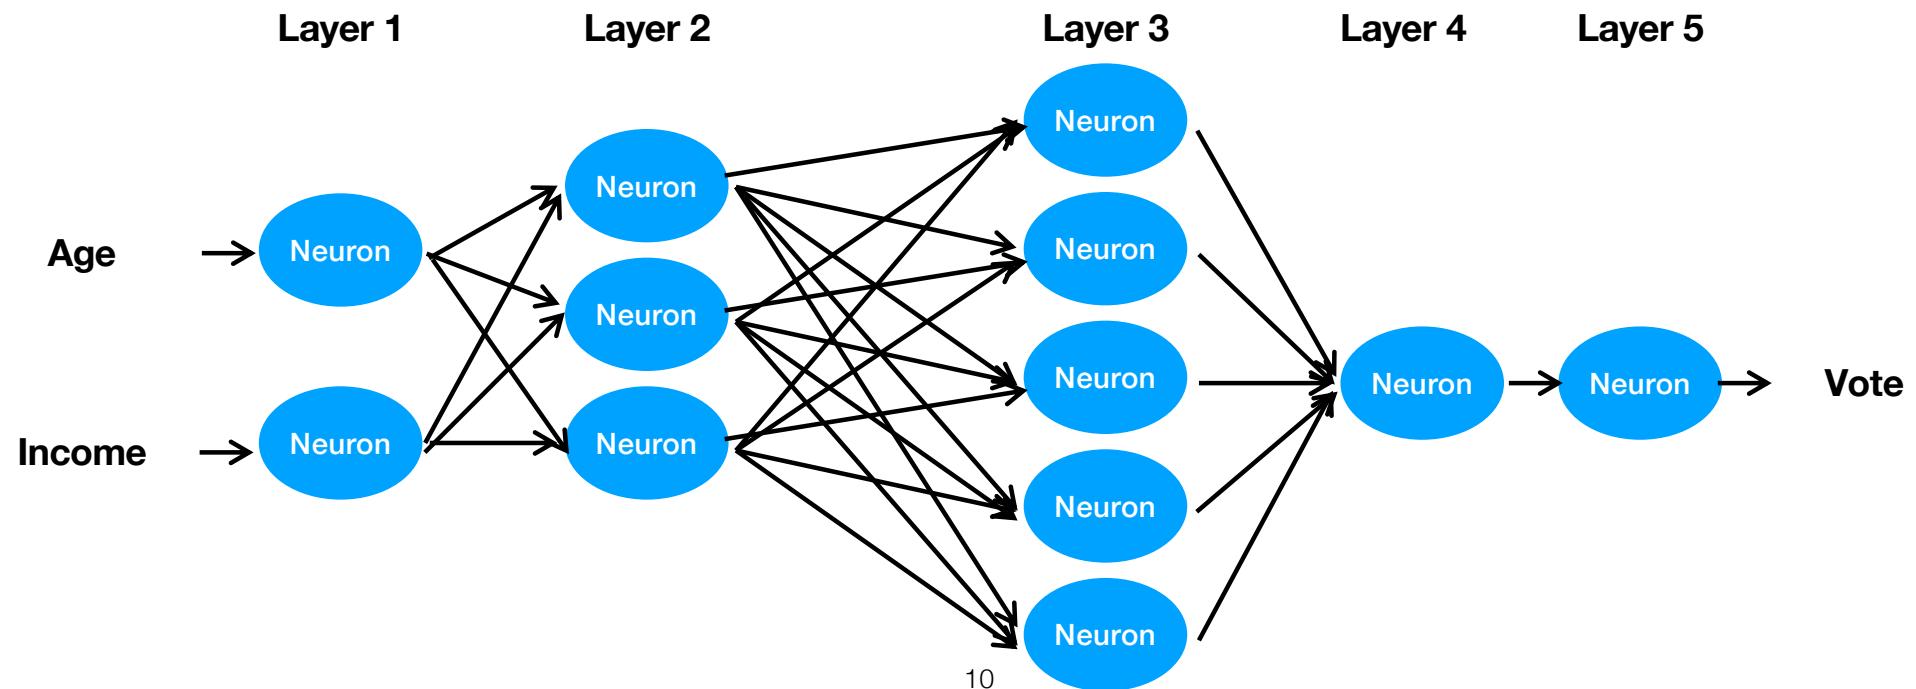


$$O_B(\text{income}, \text{age}) = \sigma(\theta_0^B + \theta_1^B \times \text{income} + \theta_2^B \times \text{age}) \quad O_D(\text{income}, \text{age}) = \sigma(\theta_0^D + \theta_1^D \times O_A + \theta_2^D \times O_B)$$

Feed-Forward Networks with Fully Connected Layers

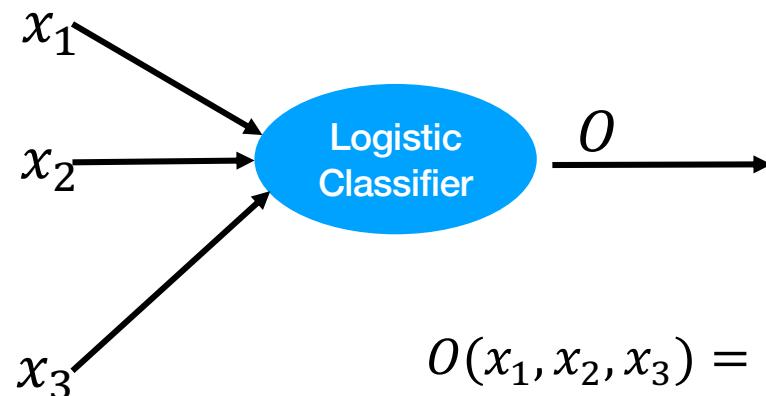
→ Each Neural Network architecture defines a function of the input with parameters θ

- Therefore, this is just a visual way of defining a complicated parameterized function of **Vote** given **Age** and **Income**:



Parameters (1/2)

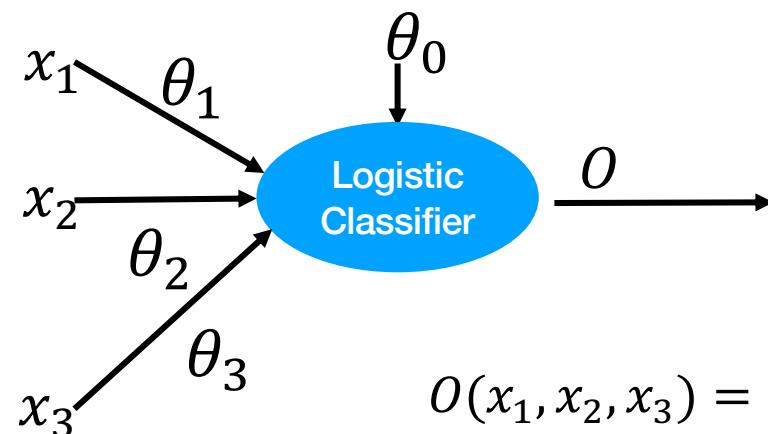
- If a neuron has N inputs, it has N+1 parameters



$$O(x_1, x_2, x_3) = \sigma(\theta_0 + \theta_1 \times x_1 + \theta_2 \times x_2 + \theta_3 \times x_3)$$

Parameters (2/2)

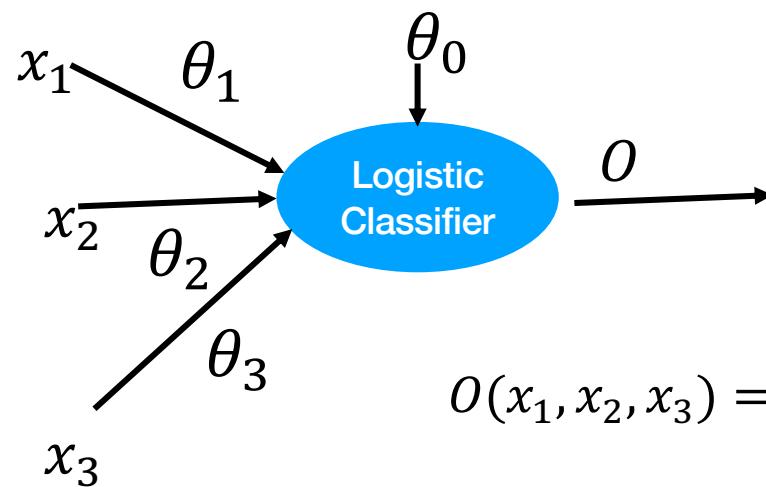
- If a neuron has N inputs, it has $N+1$ parameters
 - Visually, we can associate a parameter to each input, and show θ_0 separately



$$\theta(x_1, x_2, x_3) = \sigma(\theta_0 + \theta_1 \times x_1 + \theta_2 \times x_2 + \theta_3 \times x_3)$$

Parameters: Terminology (1/2)

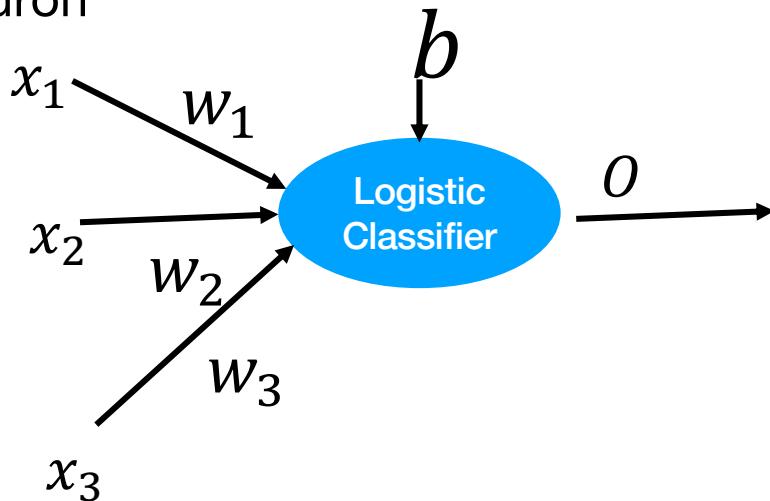
- $\theta_1, \theta_2, \theta_3$ are often called the **weights** of the neuron
- θ_0 is often called the **bias** of the neuron



$$O(x_1, x_2, x_3) = \sigma(\theta_0 + \theta_1 \times x_1 + \theta_2 \times x_2 + \theta_3 \times x_3)$$

Parameters: Terminology (2/2)

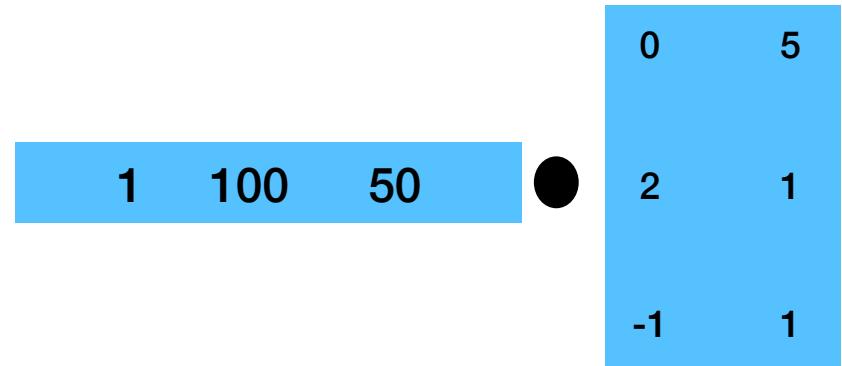
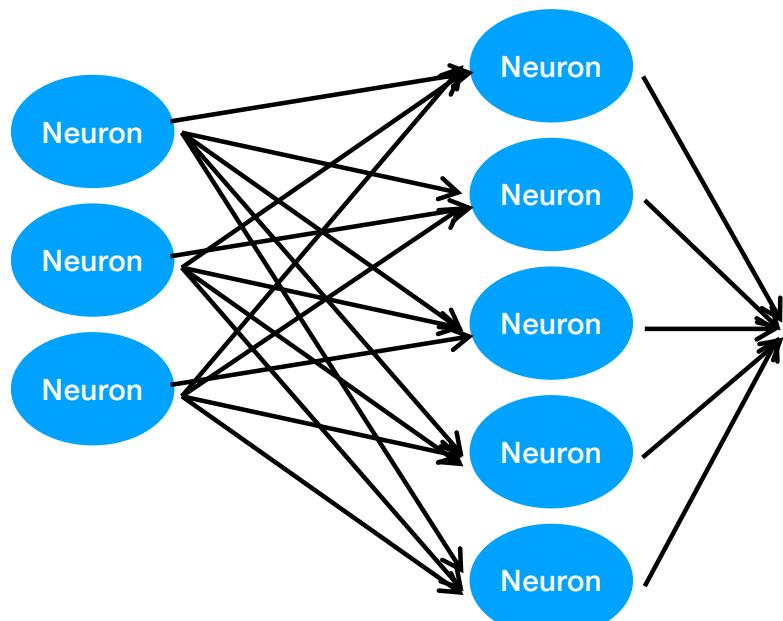
- $\theta_1, \theta_2, \theta_3$ are often called the **weights** of the neuron
 - They are therefore often also noted w_1, w_2, w_3
- θ_0 is often called the **bias** of the neuron
 - It is often noted **b**



$$O(x_1, x_2, x_3) = \sigma(b + w_1 \times x_1 + w_2 \times x_2 + w_3 \times x_3)$$

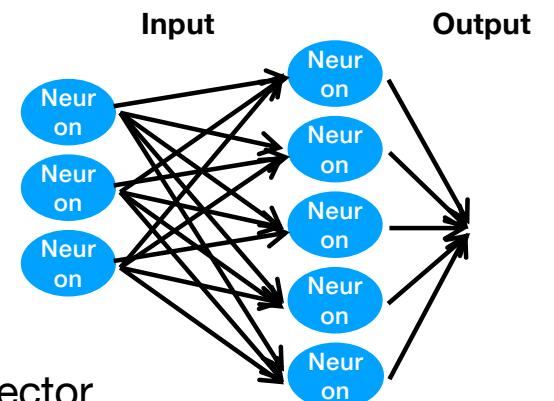
“Connectionist” View vs. “Mathematical” View (1/2)

- “Connections” viewpoint
 - *Fully Connected Layers*
- Math operations viewpoint
 - *Matrix multiplications*



“Connectionist” View vs. “Mathematical” View (2/2)

- In practice, computing the output of a fully connected layer is equivalent to computing:
 - a Matrix multiplication
 - followed by an activation
- Equivalences:
 - Input of the layer (= output of the previous layer) \leftrightarrow row vector
 - Output of the layer (= set of output of each neuron in a layer) \leftrightarrow row vector
 - Weights of a neuron \leftrightarrow column vector
 - Set weights of all neurons in a fully connected layer \leftrightarrow matrix
 - Batched input of a layer (= several inputs at once) \leftrightarrow matrix



Linear Algebra

- In order to discuss the way Neural Networks are actually implemented, we need to discuss some mathematical concepts
- We discuss:
 - Vector scalar products (a.k.a inner product)
 - Matrices
 - Matrix-vector multiplication
 - Matrix-Matrix multiplication
- Warning: it is a bit ambitious to explain that in 30 minutes. Hopefully you understand most of it.

Row Vectors and Column Vectors

- We have seen that **vectors** are just a “list” of numbers
- We are actually going to distinguish 2 “types” of vectors: **row vectors** and **column vectors**

**(5-dimensional)
Column vector:**

0.2
3.6
2.1
5.3
-2.2

**(5-dimensional)
Row vector:**

0.2 3.6 2.1 5.3 -2.2

Column Vectors: Add

- Column vectors of the same dimension can be added

(5-dimensional) Column vector (5-dimensional) Column vector (5-dimensional) Column vector

0.2		1.1		1.3
3.6		-2.0		1.6
2.1	+	1.1	=	3.2
5.3		-5.3		0.0
-2.2		-1.0		-3.2

Column Vectors: Multiply

- Column vectors can be multiplied by a number

(5-dimensional) (5-dimensional)
Column vector **Column vector**

$$2.0 \times \begin{bmatrix} 1.1 \\ -2.0 \\ 1.1 \\ -5.3 \\ -1.0 \end{bmatrix} = \begin{bmatrix} 2.2 \\ -4.0 \\ 2.2 \\ -10.6 \\ -2.0 \end{bmatrix}$$

Row Vectors

- Row vectors have the same operations as Column vectors
- They can be added:

$$\begin{array}{c} \text{(5-dimensional)} \\ \text{Row vector} \\ \hline 0.2 \ 3.6 \ 2.1 \ 5.3 \ -2.2 \end{array} + \begin{array}{c} \text{(5-dimensional)} \\ \text{Row vector} \\ \hline 1.1 \ -2.0 \ 1.1 \ -5.3 \ -1.0 \end{array} = \begin{array}{c} \text{(5-dimensional)} \\ \text{Row vector} \\ \hline 1.3 \ 1.6 \ 3.2 \ 0.0 \ -3.2 \end{array}$$

- They can be multiplied by a number:

$$2.0 \times \begin{array}{c} \hline 0.2 \ 3.6 \ 2.1 \ 5.3 \ -2.2 \end{array} = \begin{array}{c} \hline 0.4 \ 7.2 \ 4.2 \ 10.6 \ -4.4 \end{array}$$

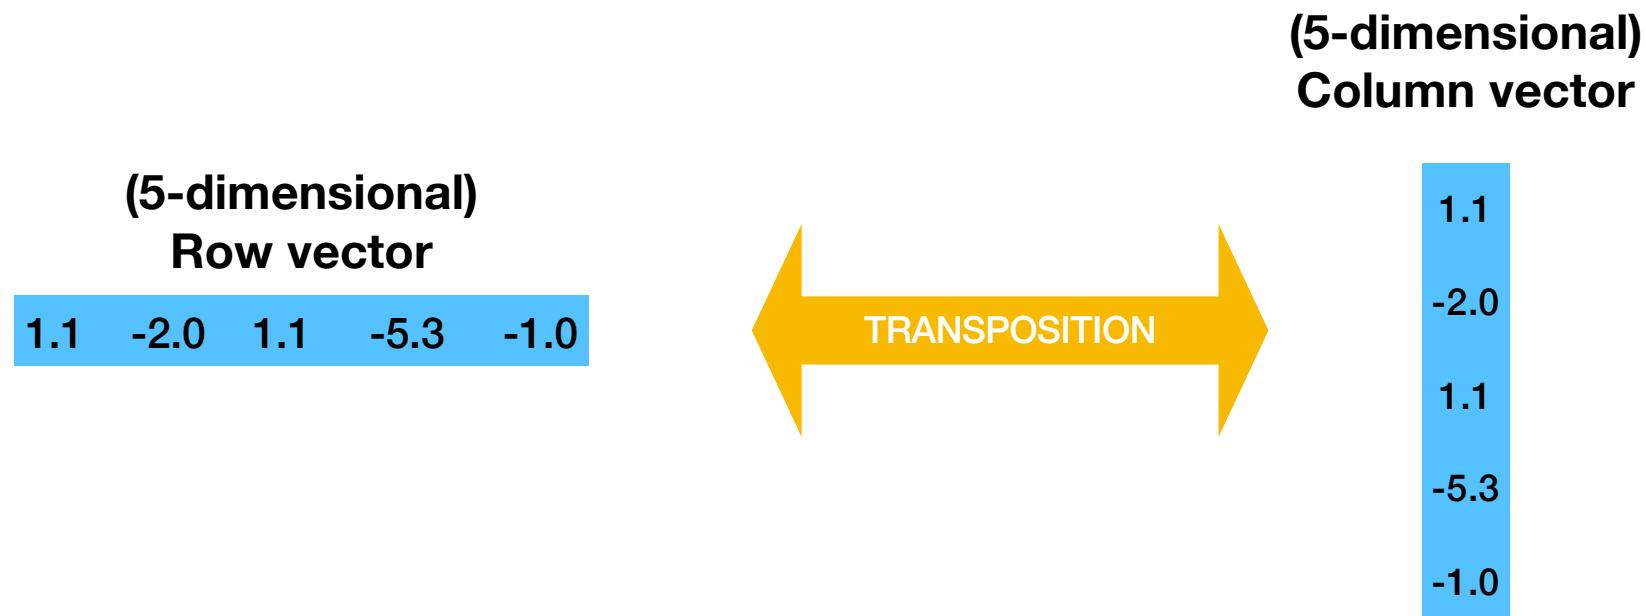
Row Vectors and Column Vectors

- Row vectors and column vectors cannot be added together

$$\begin{array}{c} \text{(5-dimensional)} \\ \text{Column vector} \\ \\ \text{(5-dimensional)} \\ \text{Row vector} \\ \begin{array}{ccccc} 0.2 & 3.6 & 2.1 & 5.3 & -2.2 \end{array} \end{array} + \begin{array}{c} 2.2 \\ -4.0 \\ 2.2 \\ -10.6 \\ -2.0 \end{array} = \text{X}$$

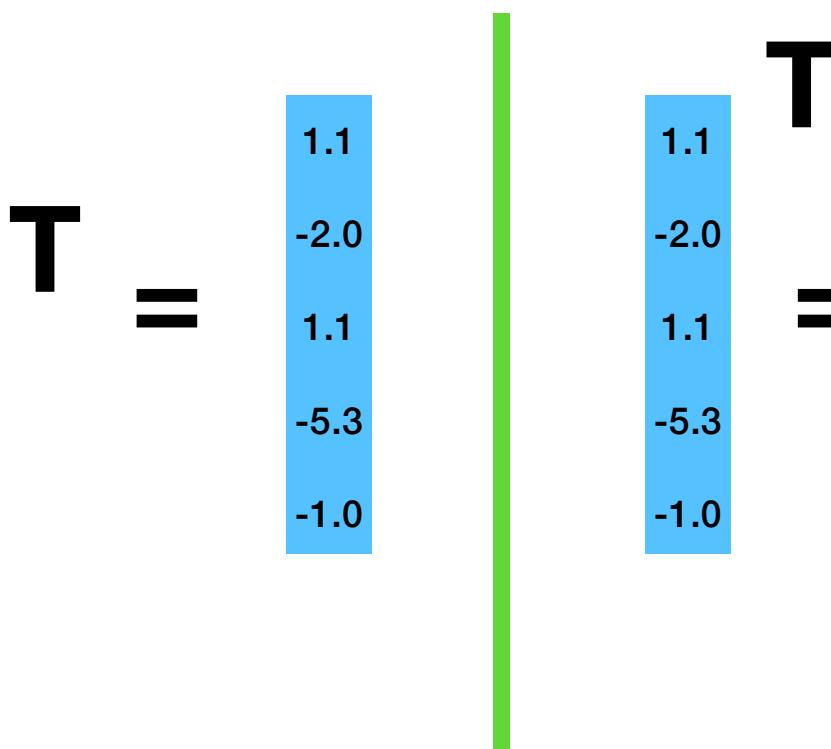
Transposition (1/2)

- However, *row vectors* can be transformed in *column vectors* by an operation called **transposition** (and vice-versa)



Transposition (2/2)

- Transposition is usually noted with a T in exponent:

$$\begin{matrix} 1.1 & -2.0 & 1.1 & -5.3 & -1.0 \end{matrix} \quad \mathbf{T} = \quad \begin{matrix} 1.1 \\ -2.0 \\ 1.1 \\ -5.3 \\ -1.0 \end{matrix} \quad \mathbf{T} = \quad \begin{matrix} 1.1 & -2.0 & 1.1 & -5.3 & -1.0 \end{matrix}$$


Inner product (1/4)

- We can compute the *inner product* of a *row vector* and a *column vector* to obtain a single number
- It consists in taking the product of the number dimension by dimension and then taking the sum

(3-dimensional)
Column vector

(3-dimensional)
Row vector

$$\begin{matrix} 2 & 0.0 & -1.0 \end{matrix} \bullet \begin{matrix} 1.5 \\ 1.0 \\ 2.0 \end{matrix} = 2 \times 1.5 + 0 \times 1 + -1 \times 2 = 1.0$$

Inner product (2/4)

- The vectors should have the **same dimension**
- You should have the row vector on the left, and the column vector on the right
 - (if it is the opposite, the operation is called *outer product* and gives a different result)

**(3-dimensional)
Column vector**

**(3-dimensional)
Row vector**

2	0.0	-1.0
---	-----	------

●

1.5
1.0
2.0

= **1.0**

Inner Product (3/4)

$$\begin{bmatrix} 2 & 0.0 & -1.0 \end{bmatrix} \bullet \begin{bmatrix} 1.5 \\ 1.0 \\ 2.0 \end{bmatrix} = 1.0$$

You can try to visualize this as the column vector lying down on the row vector to produce the number

Inner Product (4/4)

$$\begin{array}{ccc} \downarrow & \downarrow & \downarrow \\ \begin{array}{ccc} 1.5 & 1.0 & 2.0 \\ \times & \times & \times \\ 2 & 0.0 & -1.0 \end{array} & \bullet & = 1.0 \\ 3.0 + 0.0 + -2.0 & \rightarrow & 1.0 \end{array}$$

You can try to visualize this as the column vector lying down on the row vector to produce the number

Why Inner Product is Important? (1/2)

- It allows us to **express linear functions efficiently**
- And remember that linear functions are one of the **fundamental components** of Machine Learning

Linear regression

$$f(x, y) = \theta_0 + \theta_1 \times x + \theta_2 \times y$$

Logistic Classifier

$$\begin{aligned} score(x, y) &= \theta_0 + \theta_1 \times x + \theta_2 \times y \\ prediction &= \sigma(score) \end{aligned}$$

Neuron (same formula as Logistic Classifier)

$$output(x, y) = \sigma(\theta_0 + \theta_1 \times x + \theta_2 \times y)$$

Why Inner Product is Important? (2/2)

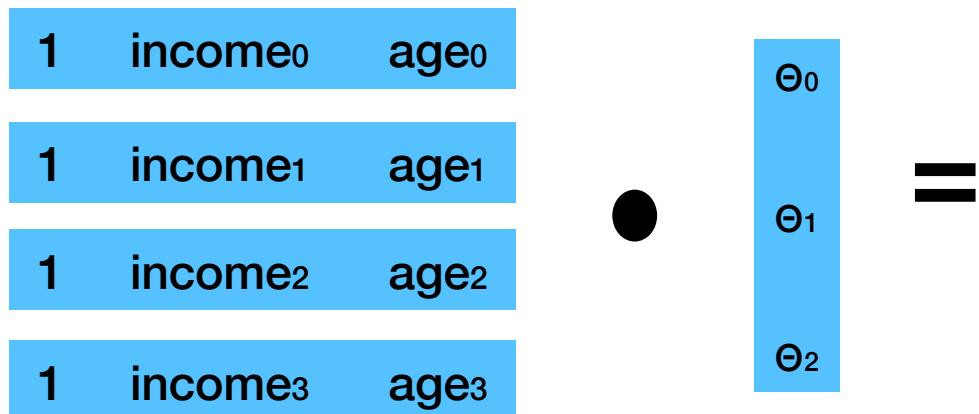
$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$

$$\begin{matrix} 1 & \text{income} & \text{age} \end{matrix} \bullet \begin{matrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{matrix} = \theta_0 + \theta_1 \times income + \theta_2 \times age$$

$$vote(age, income) = \sigma \left(\begin{matrix} 1 & \text{income} & \text{age} \end{matrix} \bullet \begin{matrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{matrix} \right)$$

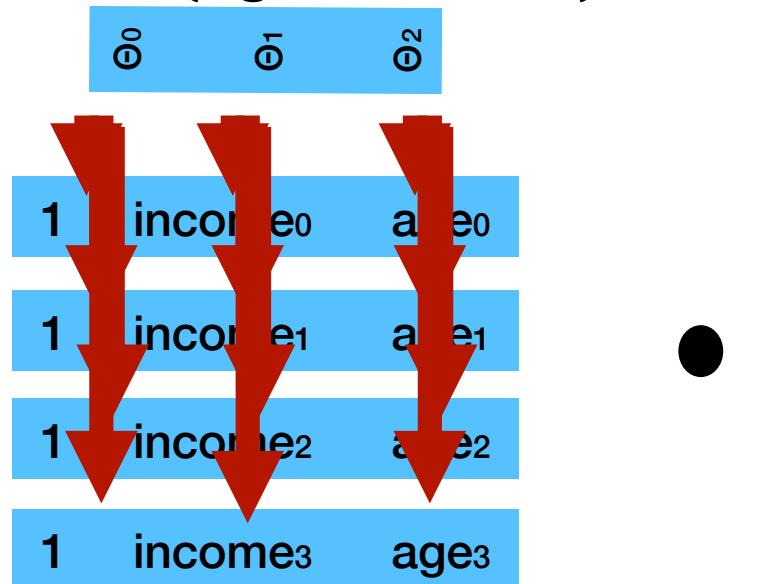
Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$



Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$



• =

$$\theta_0 + \theta_1 \times income_0 + \theta_2 \times age_0$$

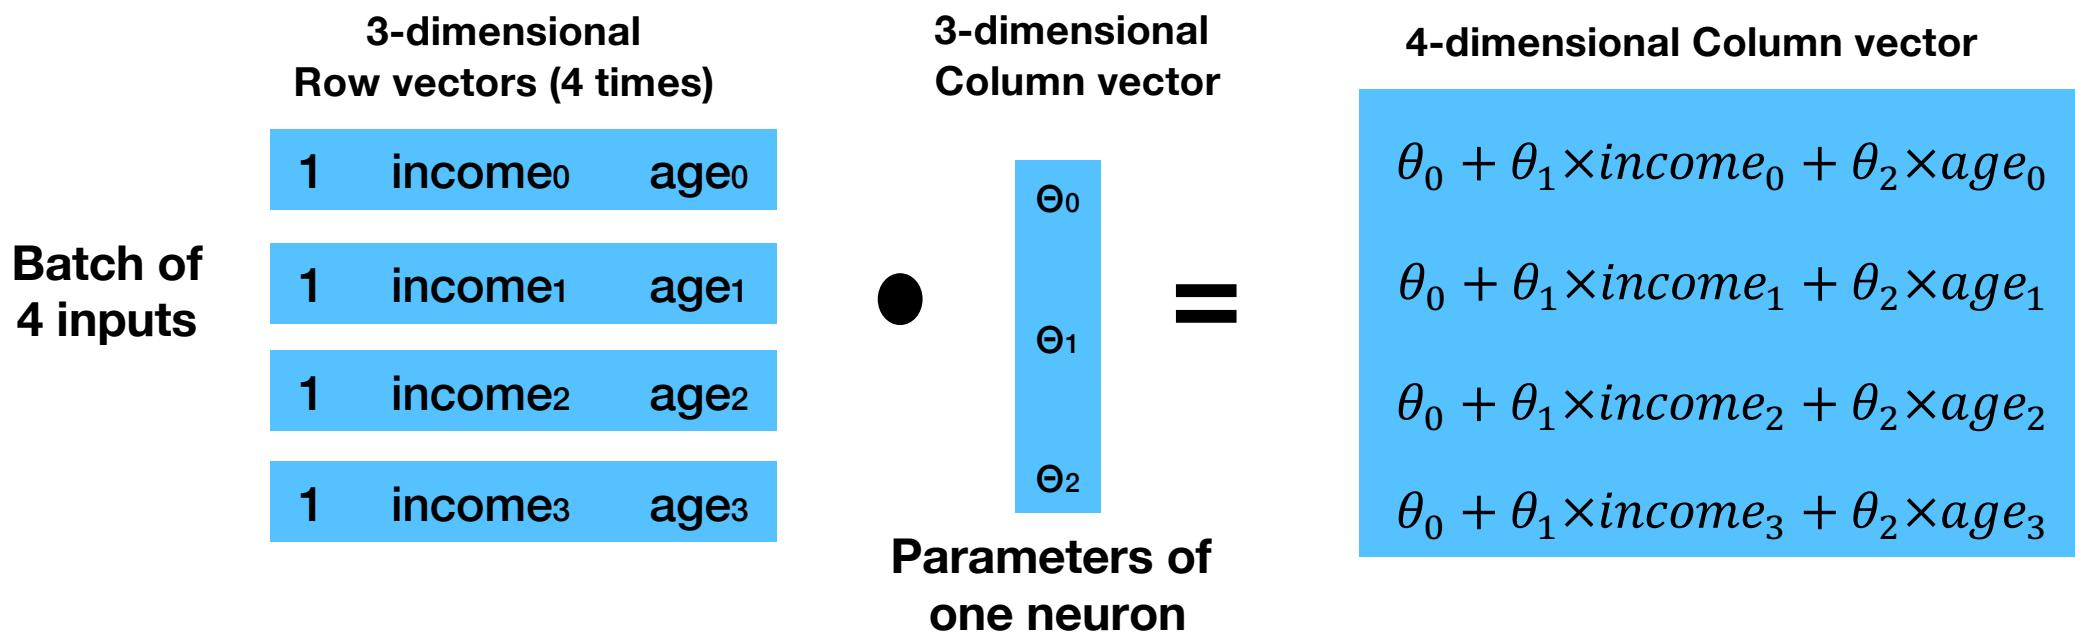
$$\theta_0 + \theta_1 \times income_1 + \theta_2 \times age_1$$

$$\theta_0 + \theta_1 \times income_2 + \theta_2 \times age_2$$

$$\theta_0 + \theta_1 \times income_3 + \theta_2 \times age_3$$

Representing Many Inner Product at Once

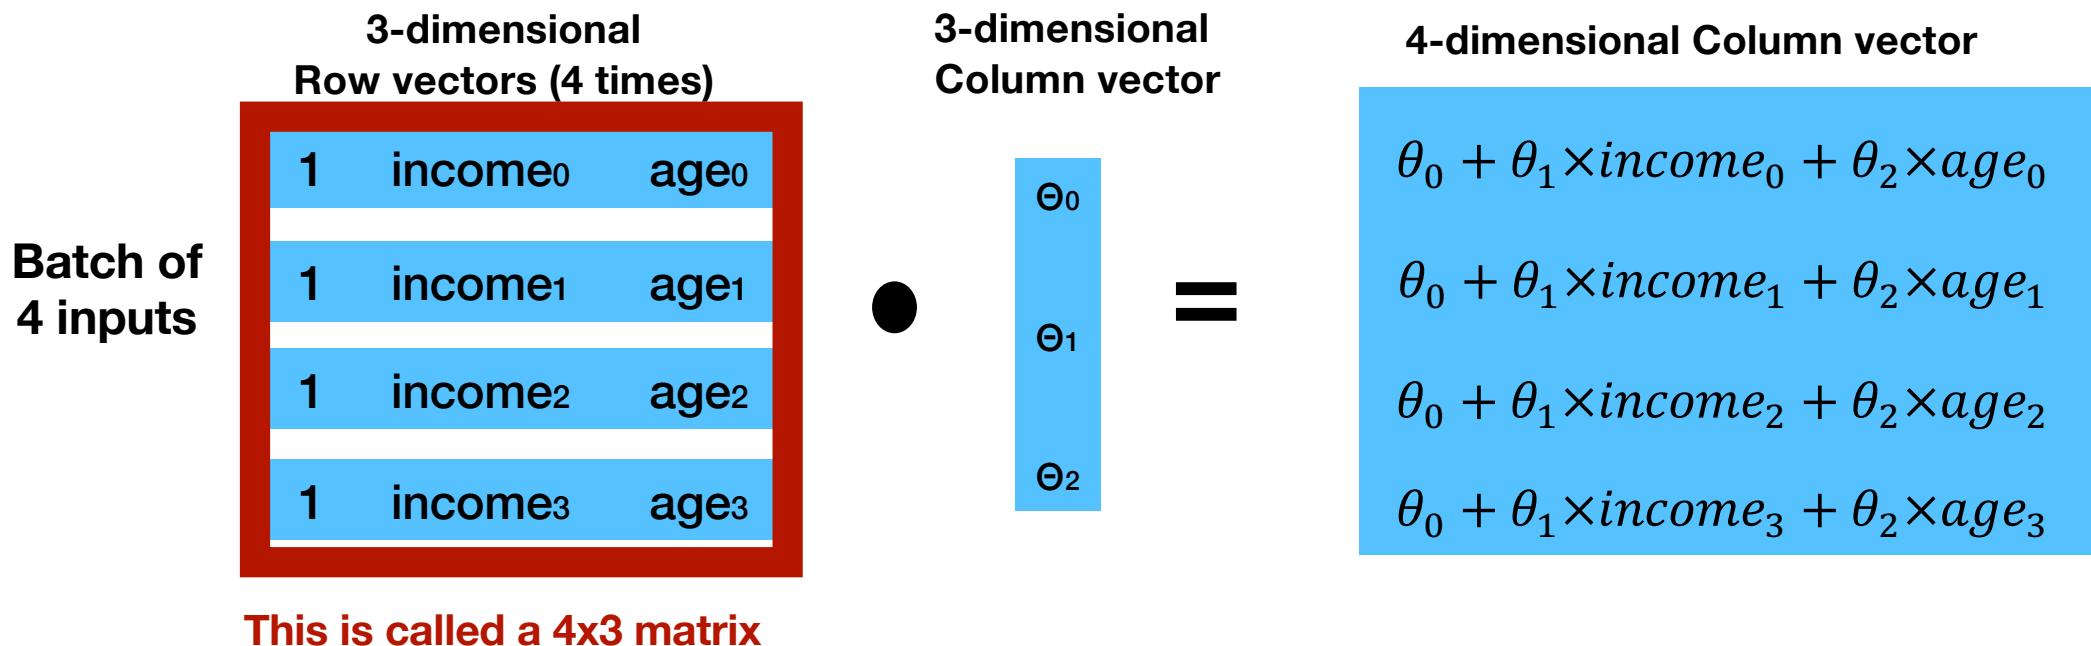
$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$



This represents the computation of the score output of a single neuron for 4 inputs at the same time!

Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$



Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$

3-dimensional
Row vectors (4 times)

1	income ₀	age ₀
1	income ₁	age ₁
1	income ₂	age ₂
1	income ₃	age ₃

3-dimensional
Column vector

θ_0
θ_1
θ_2



4-dimensional Column vector

$$\begin{aligned} & \theta_0 + \theta_1 \times income_0 + \theta_2 \times age_0 \\ & \theta_0 + \theta_1 \times income_1 + \theta_2 \times age_1 \\ & \theta_0 + \theta_1 \times income_2 + \theta_2 \times age_2 \\ & \theta_0 + \theta_1 \times income_3 + \theta_2 \times age_3 \end{aligned}$$

This is called a 4x3 matrix

Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$

Batch of
2 inputs

3-dimensional
Row vectors
(2 times)

1	10	20
1	30	40

3-dimensional
Column vector

1
2
-1

2-dimensional Column vector

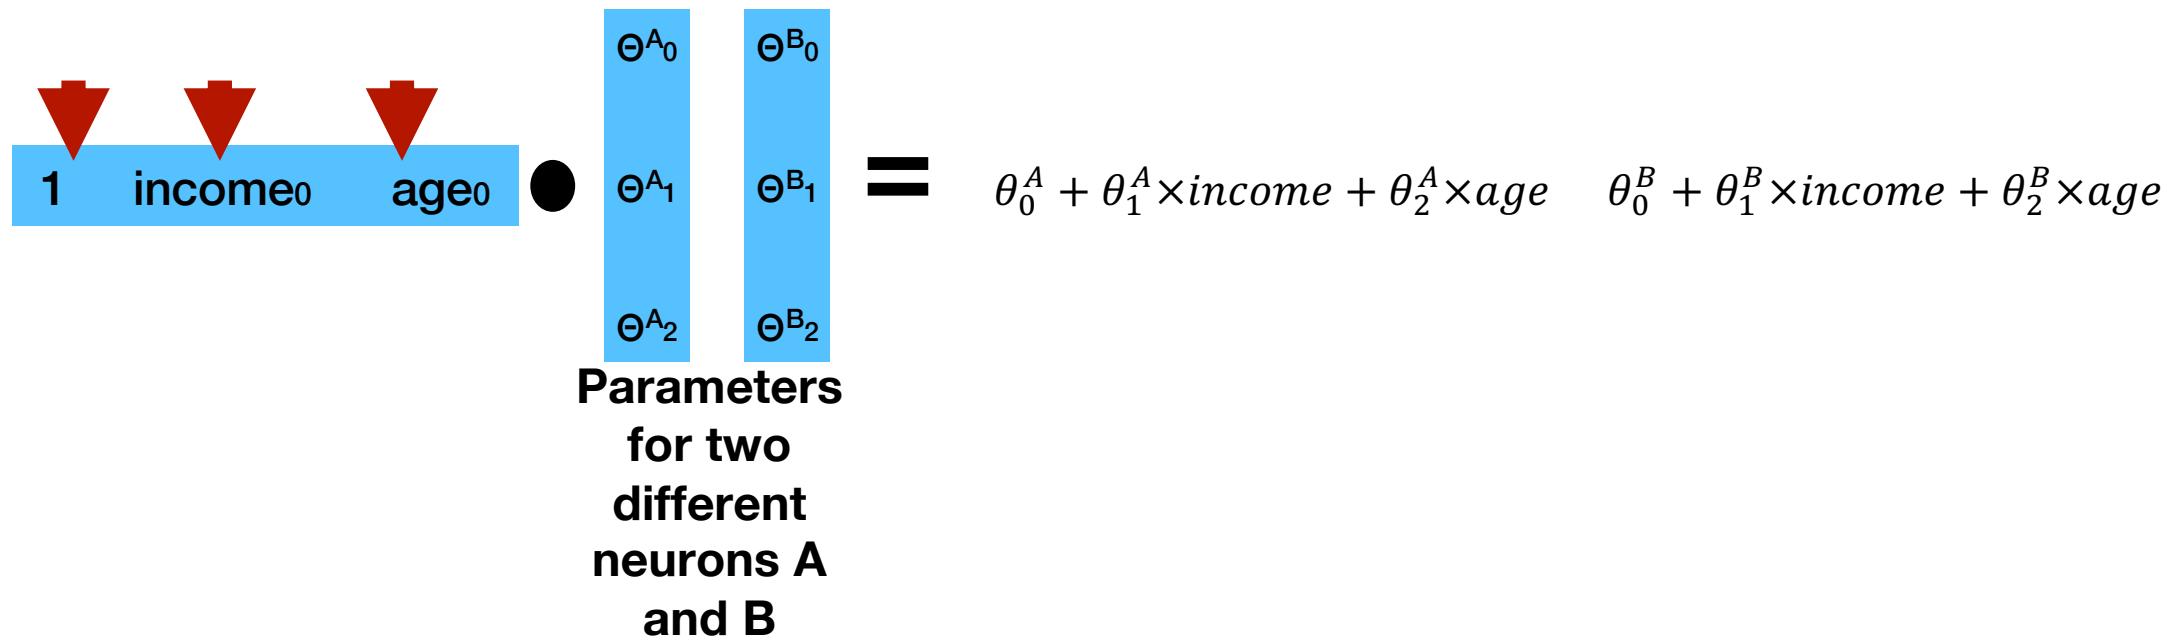
=

?

Parameters of one
Neuron

Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times income + \theta_2 \times age)$$



Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$

3-dimensional
Row vectors

1 income₀ age₀

3-dimensional
Column vector
(2 times)

Θ^A_0	Θ^B_0
Θ^A_1	Θ^B_1
Θ^A_2	Θ^B_2

2-dimensional Row vector

$\theta_0^A + \theta_1^A \times income + \theta_2^A \times age$ $\theta_0^B + \theta_1^B \times income + \theta_2^B \times age$

This is called a
3x2 matrix

Representing Many Inner Product at Once

$$vote(age, income) = \sigma(\theta_0 + \theta_1 \times age + \theta_2 \times income)$$

3-dimensional
Column vector
(2 times)

3-dimensional
Row vectors

1 100 50

0	5
2	1
-1	1

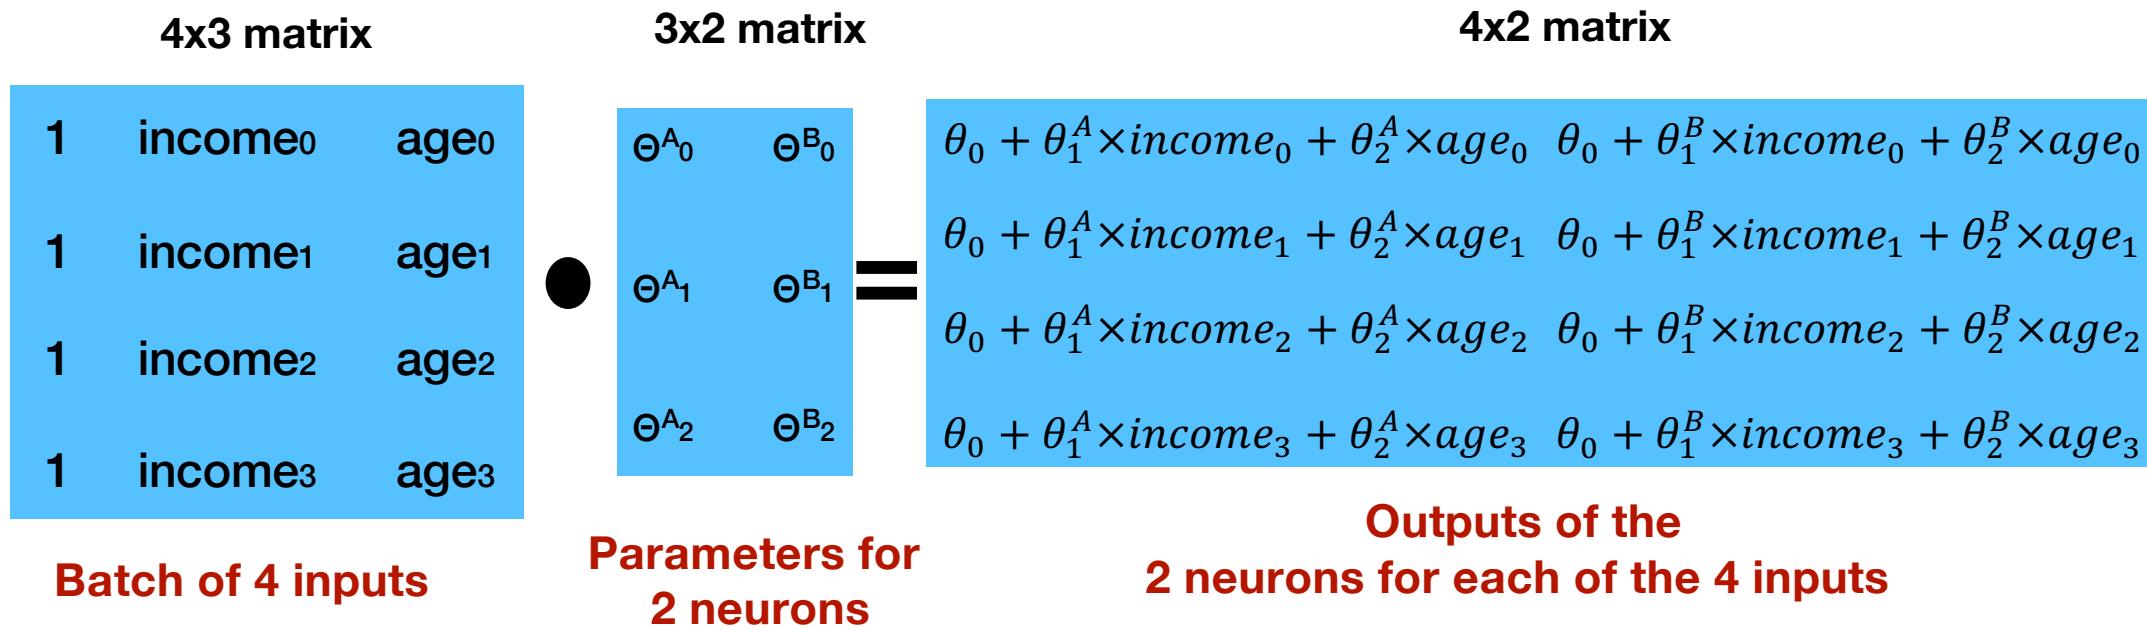
2-dimensional Row vector

Batch of
1 input

Parameters of two
Neurons

Representing Many Inner Product at Once

- If we combine the 2 aspects of having several inputs and several neurons, we have what is called a *matrix multiplication*



Matrix Multiplication (1/4)

4x3 matrix		
1	income ₀	age ₀
1	income ₁	age ₁
1	income ₂	age ₂
1	income ₃	age ₃

•

3x2 matrix	
Θ^{A_0}	Θ^{B_0}
Θ^{A_1}	Θ^{B_1}
Θ^{A_2}	Θ^{B_2}

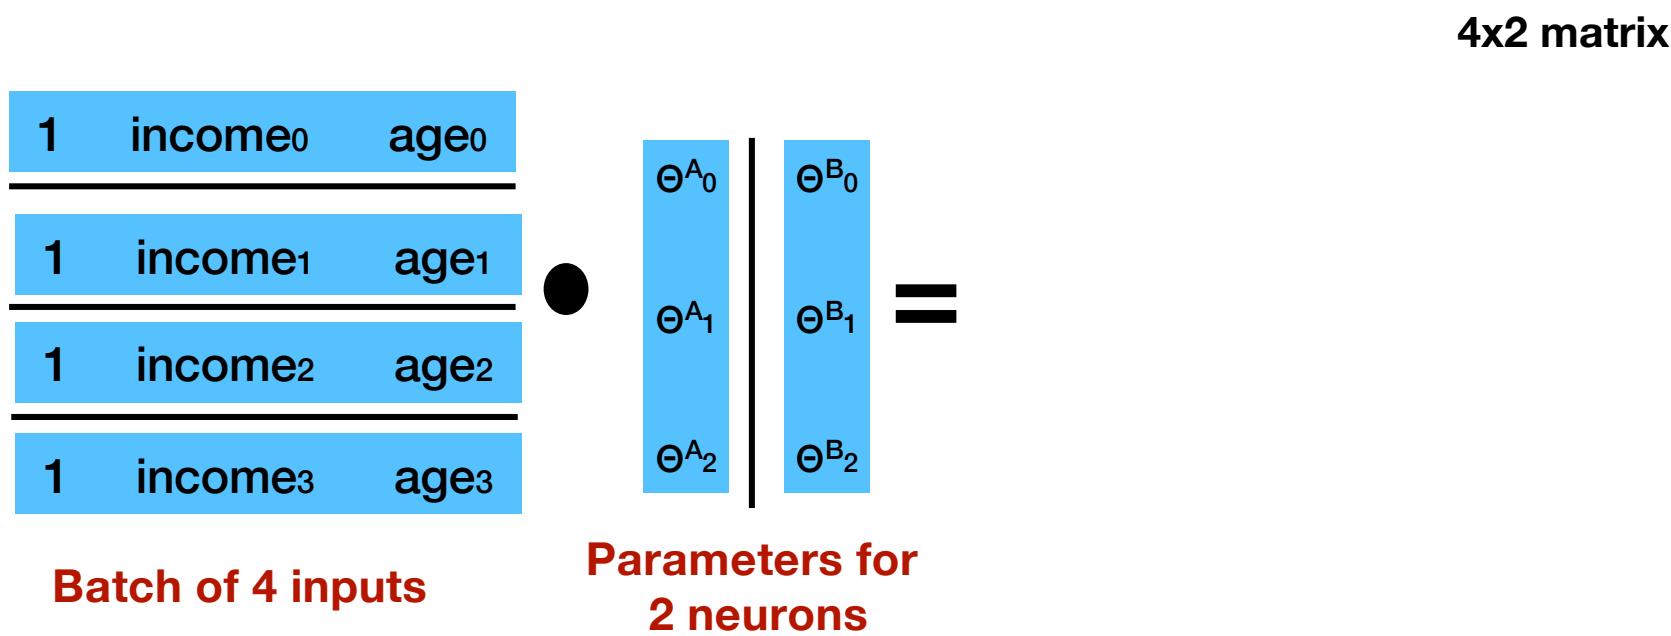
=

4x2 matrix	
------------	--

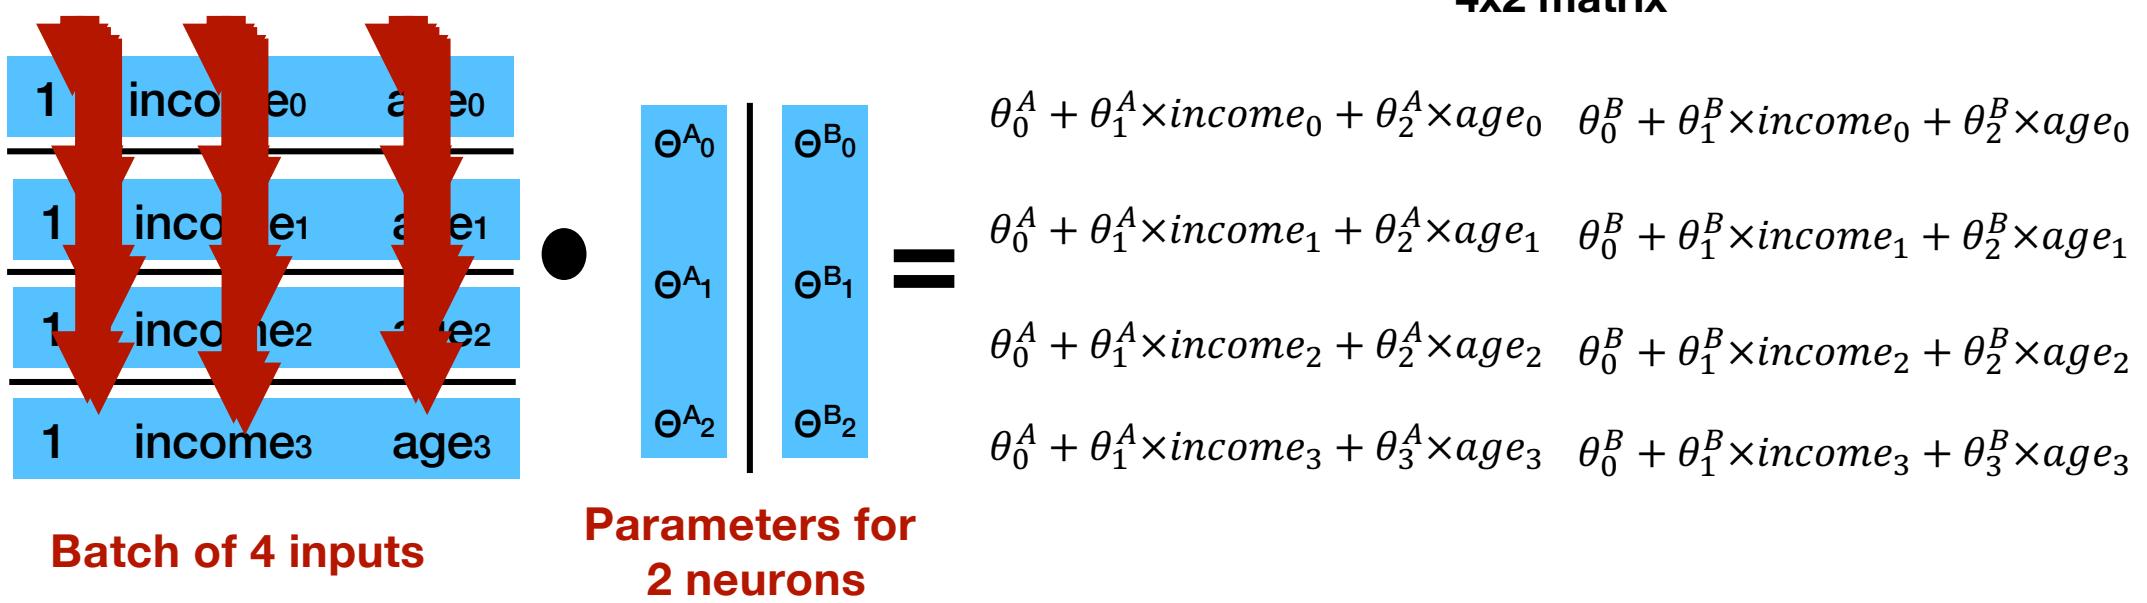
Batch of 4 inputs

Parameters for 2 neurons

Matrix Multiplication (2/4)

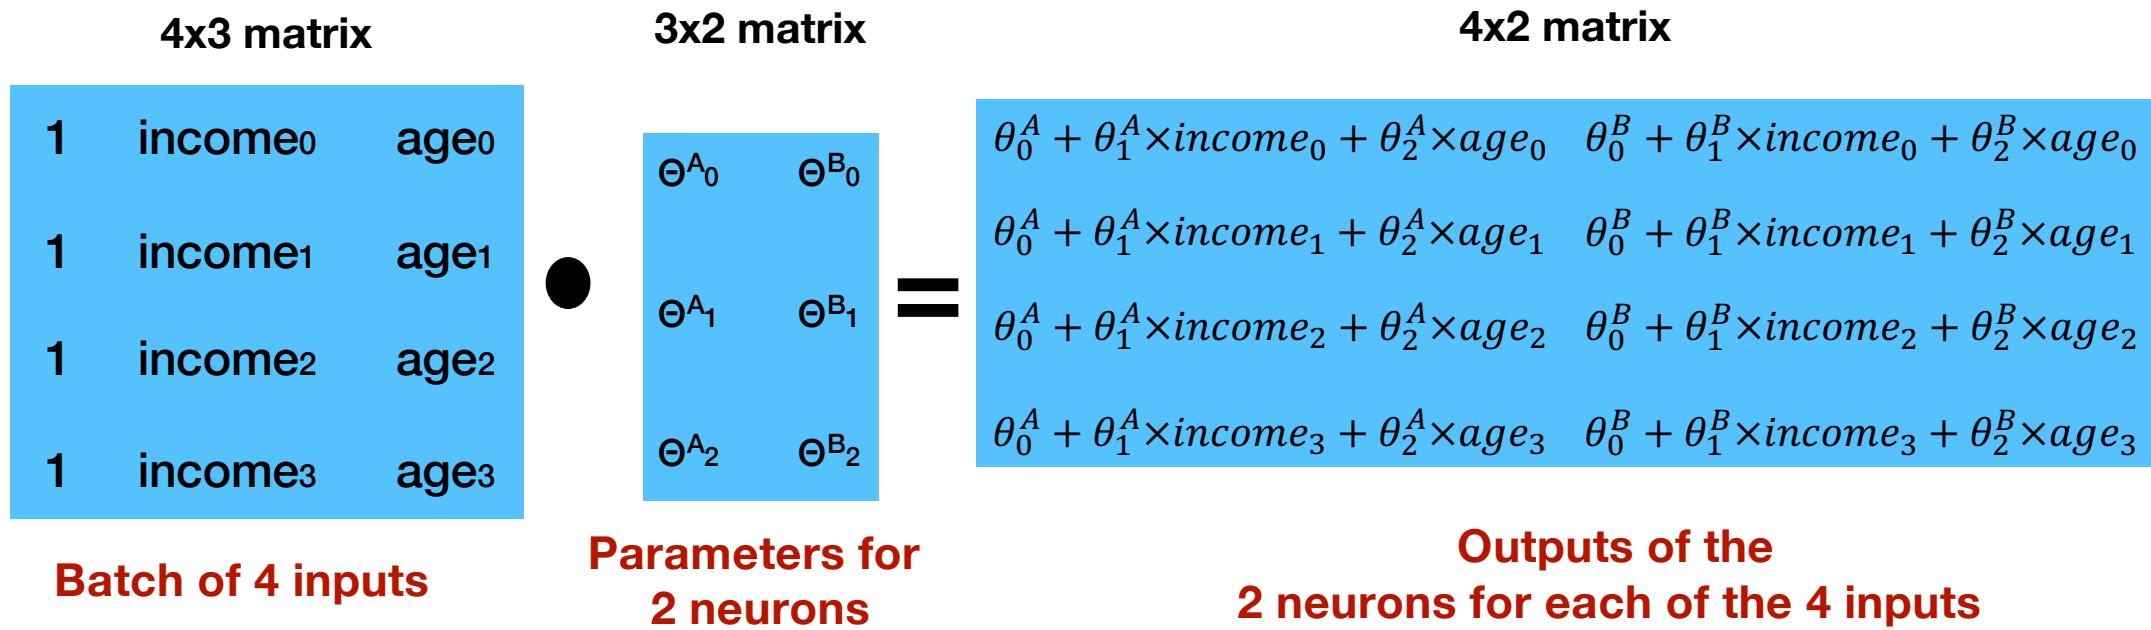


Matrix Multiplication (3/4)



Matrix Multiplication (4/4)

- We represent many inputs as a matrix
- We represent many neurons as a matrix
- Matrix multiplication compute the output score of all the neurons for all the inputs



Exercise 1

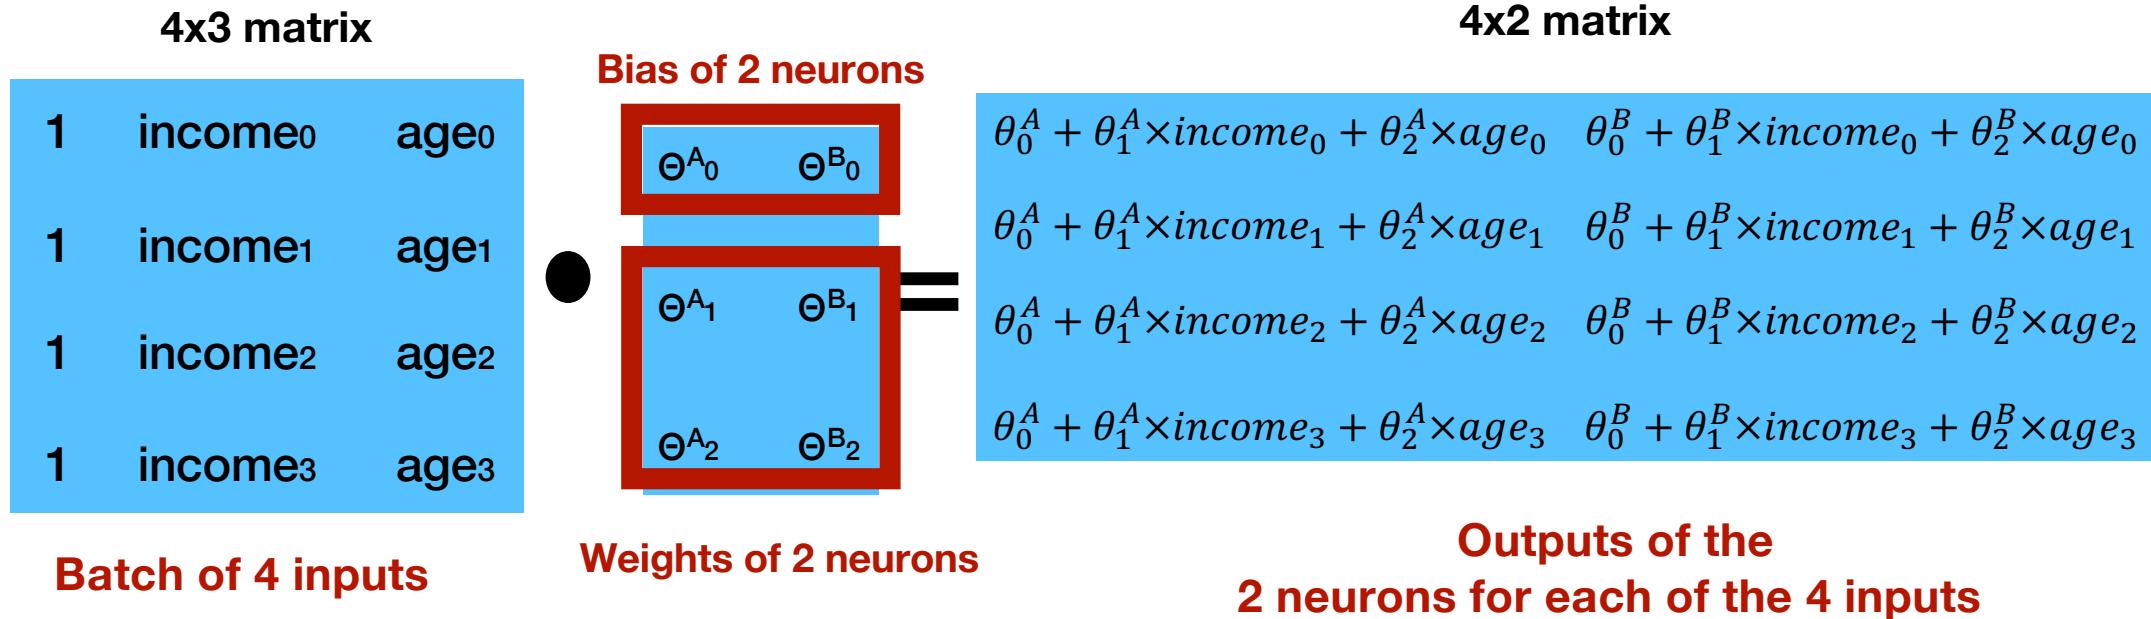
$$\begin{array}{c} \text{2x3 matrix} \\ \\ \begin{bmatrix} 1 & 100 & 50 \\ 1 & 10 & 20 \end{bmatrix} \end{array} \bullet \begin{array}{c} \text{3x2 matrix} \\ \\ \begin{bmatrix} 1 & 0 \\ 2 & 1 \\ 1 & -1 \end{bmatrix} \end{array} = \begin{array}{c} \text{2x2 matrix} \\ \\ \text{Outputs of the} \\ \text{2 neurons for each of the 2 inputs} \end{array}$$

Batch of 2 inputs

Parameters for 2 neurons

Computing the Output of Many Neurons for Many Inputs (1/2)

- In practice, we separate the **weights** from the **bias**



Computing the Output of Many Neurons for Many Inputs (2/2)

4x3 matrix

income ₀	age ₀
income ₁	age ₁
income ₂	age ₂
income ₃	age ₃

3x2 matrix

Θ^A_1	Θ^B_1
Θ^A_2	Θ^B_2

Weights matrix 2x2
(2 inputs and 2 neurons)

Batch input (4 inputs)

Expanded to fit
number of input

Θ^A_0	Θ^B_0

+

Bias vector

Θ^A_0	Θ^B_0
--------------	--------------



=

4x2 matrix

$$\begin{array}{ll}
 \theta^A_0 + \theta^A_1 \times \text{income}_0 + \theta^A_2 \times \text{age}_0 & \theta^B_0 + \theta^B_1 \times \text{income}_0 + \theta^B_2 \times \text{age}_0 \\
 \theta^A_0 + \theta^A_1 \times \text{income}_1 + \theta^A_2 \times \text{age}_1 & \theta^B_0 + \theta^B_1 \times \text{income}_1 + \theta^B_2 \times \text{age}_1 \\
 \theta^A_0 + \theta^A_1 \times \text{income}_2 + \theta^A_2 \times \text{age}_2 & \theta^B_0 + \theta^B_1 \times \text{income}_2 + \theta^B_2 \times \text{age}_2 \\
 \theta^A_0 + \theta^A_1 \times \text{income}_3 + \theta^A_2 \times \text{age}_3 & \theta^B_0 + \theta^B_1 \times \text{income}_3 + \theta^B_2 \times \text{age}_3
 \end{array}$$

Exercise 2

Bias for
2 neurons

1	0
---	---

2x2 matrix

100	50
10	20

2x2 matrix

2	1
1	-1

Batch of 2 inputs

Weights for
2 neurons

2x2 matrix

=

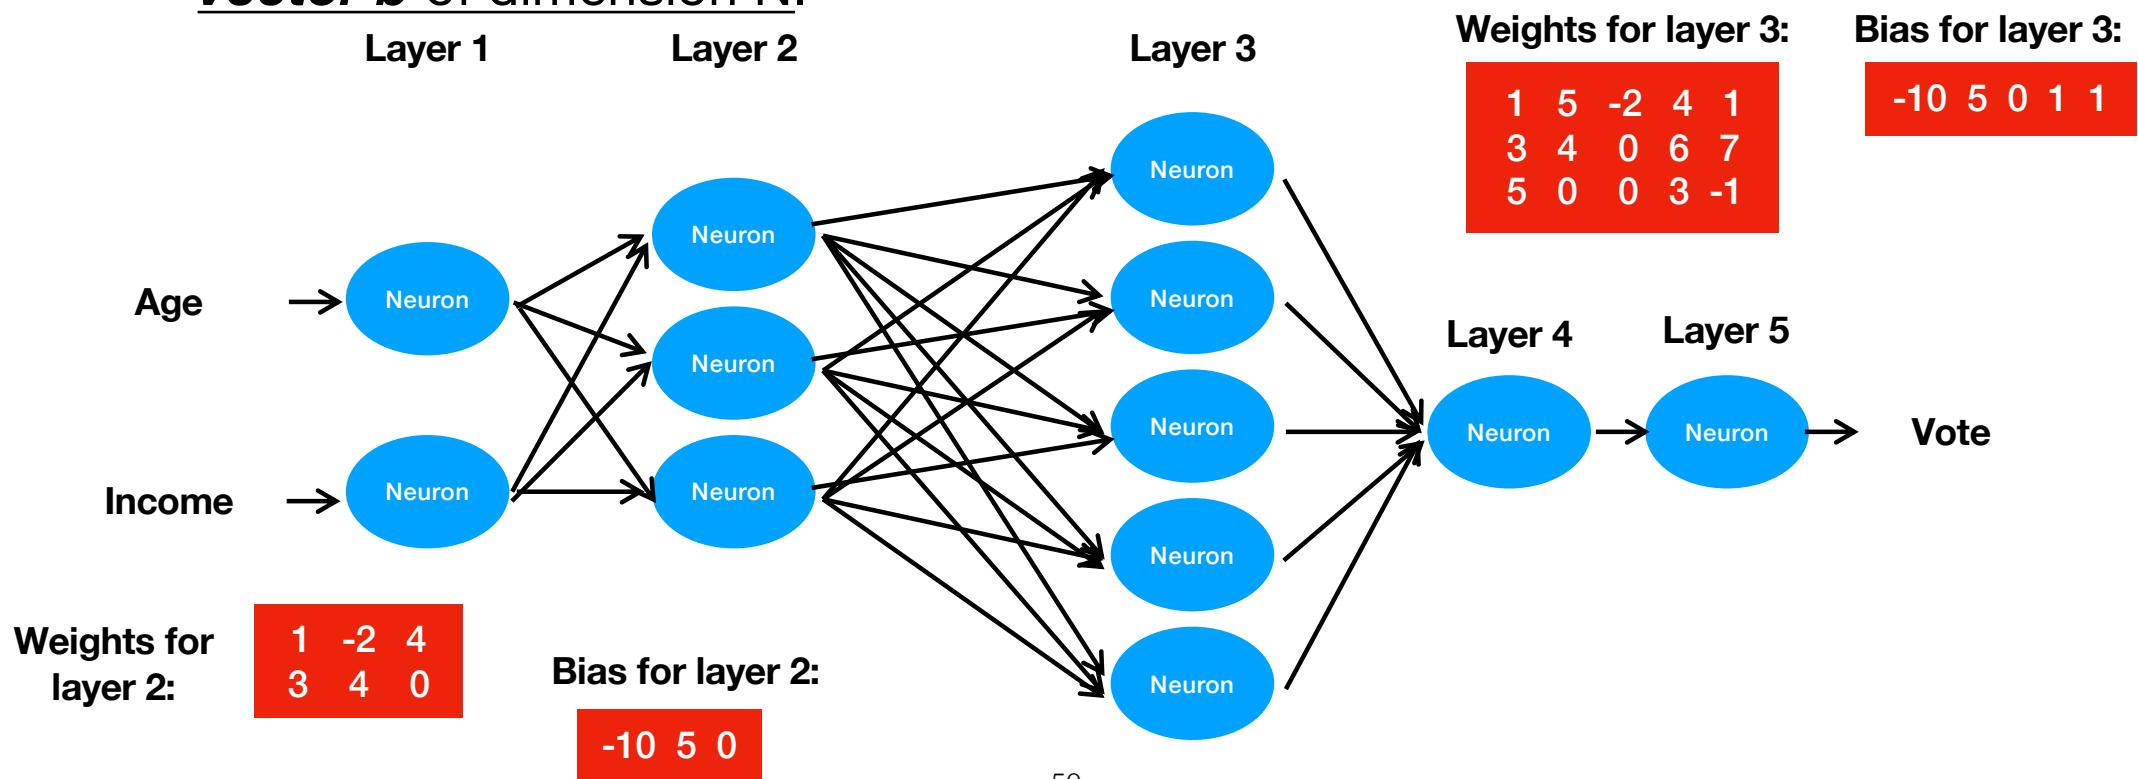
Score outputs of the
2 neurons for each of the 2 inputs

Matrix Multiplication for FNN

- Important: using these matrix representations only work if all the neurons have the same input and are fully connected
- Then, in practice, a **Fully Connected layer of N neurons with K input** can be represented by a **matrix of weights W** of shape $K \times N$, and a **bias vector b** of dimension N .

Feed-Forward Networks with Fully Connected Layers

- Then, in practice, a Fully Connected layer of N neurons with K input can be represented by a matrix of weights W of shape $K \times N$, and a bias vector b of dimension N.



Visualization for FNN

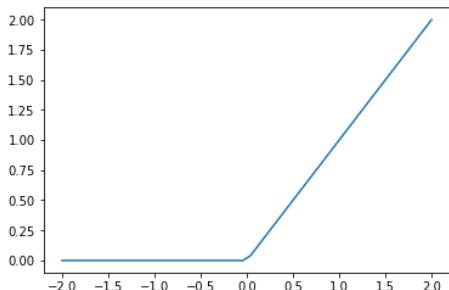
- An interactive Neural Network with Fully Connected Layers
 - From some nice people at Google
 - <https://playground.tensorflow.org>

Activations

- Some possible activation functions:

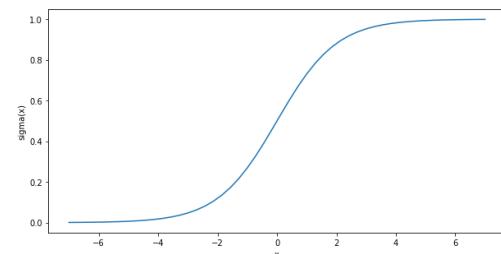
Rectified Linear Unit

$$ReLU(x) = \max(x, 0)$$



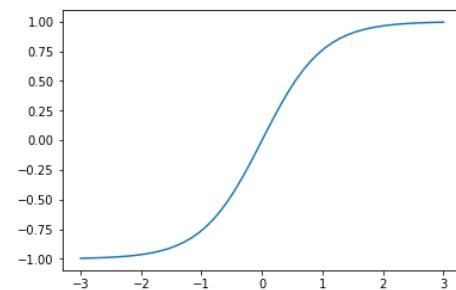
**Sigmoid
(a.k.a logistic function)**

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



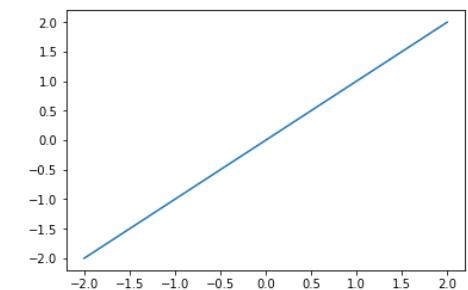
**Hyperbolic Tangent
(= a linear transformation
of the sigmoid)**

$$\tanh(x) = 2\sigma(2x) - 1$$



**Linear function
(useless as
activation function)**

$$f(x) = x$$



L1 and L2 Regularization

- We have discussed that L2 Regularization can be used for reducing overfitting
 - *-> Change Regularization and Regularization Rate to observe the effect of Regularization*

L1 Regularization

- We note $|\Theta|^1$ the sum of the absolute value of all parameters Θ_k (this is called the “L1 Norm”)
- Then we add this quantity to the loss we want to minimize:

$$\text{Loss} = \text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (\text{model}(\text{cig}_i, \text{bmi}_i, \text{ismale}_i) - \text{age}_i)^2$$



$$\text{Loss} = \frac{1}{N} \cdot \sum_i (\text{model}(\text{cig}_i, \text{bmi}_i, \text{ismale}_i) - \text{age}_i)^2 + \lambda \overrightarrow{|\theta|}$$

- Then we apply Gradient Descent to this new loss

L2 Regularization

- We note $|\Theta|^2$ the sum of the square of all parameters Θ_k (this is called the “L2 Norm”)
- Then we add this quantity to the loss we want to minimize:

$$\text{Loss} = \text{MeanSquaredError} = \frac{1}{N} \cdot \sum_i (\text{model}(\text{cig}_i, \text{bmi}_i, \text{ismale}_i) - \text{age}_i)^2$$



$$\text{Loss} = \frac{1}{N} \cdot \sum_i (\text{model}(\text{cig}_i, \text{bmi}_i, \text{ismale}_i) - \text{age}_i)^2 + \lambda \|\theta\|^2$$

- Then we apply Gradient Descent to this new loss

Random Restarts

- Because the weights of a neural network are initialized randomly, the result of a training will change every time we reinitialize the network
- Therefore, to obtain best performances, it is common to train a networks several times with different random initializations, and keep the best result
- -> *Press the reload button to restart a training with new random weight initialization*

Regularization Methods for Neural Network

- When we train a network with many neurons, the danger of overfitting is large
- There are a few technics that are very efficient at preventing this:
 - Early Stopping
 - Dropout
 - Weight Decay
 - Stochastic Gradient Descent

Early Stopping

- We keep a validation set separate from the training data
- We fix a **patience** number (typically patience = 10 or 20)
- During training, if we see no improvement on a validation set after **patience** measures, we stop the training
 - *Check the evolution of validation loss to detect when training should be stopped*

Dropout

- During training, we add random noise to disturb the network
- In practice, we randomly “cut” a certain proportion of connections (typically 10% to 50%)
 - *Check the effect of adding noise*

Stochastic Gradient Descent

- Instead of computing the gradient of the loss for all the training data, we compute it for **a subsampled part** of the training data
- This is actually done anyway to get faster training
- But it is also beneficial to prevent overfitting even if you could afford to compute the gradient for all the examples at once
 - -> *change batch size to observe what happens when we compute gradient on more or less examples*

Input Features

- We see that even for a simple model (like linear regression), adding new functions of the input can increase the capacity of the model
 - *-> Add input features and see the effect when there are few neurons*

Report

- Submit the report of **exercises 1 and 2 in pdf** via PandA
 - Submission due: **next lecture**
 - Name the pdf file as **student id_name**.

Google Colab Notebook

- Let us train FNN with Google Colab notebook

<https://shorturl.at/RfCjU>